

<!-- toc:start -->

Contents

- [Closed-Loop Security of VLM, VLA, and World-Action Models: An Evidence-Based Review](#)
- [Abstract](#)
- [Introduction: Why the Security Problem Escalates from "Seeing Wrong" to "Doing Wrong"](#)
- [Corpus, Search, and the Evidence Contract](#)
- [From Answer Safety to Action Safety: Technical Lineage and System Contract](#)
- [Closed-Loop Attack Taxonomy: The First-Broken System Interface](#)
- [How Attacks Propagate from Inputs to Physical Consequences](#)
- [Defense in Depth Mirrored to the Attack Chain](#)
- [Cross-Family Synthesis: When to Choose Which Layer of Control](#)
- [Data, Metrics, and Quantitative Evidence Boundaries](#)
- [Reproduction Audit and Anchor Cases](#)
- [Research Agenda, Practical Gates, and Limitations](#)
- [Conclusion](#)
- [Appendix — Post-cutoff update \(2026-08-06 → 2026-09-26\)](#)
- [A.1 — Remotely exploitable root-level flaws in a humanoid robot](#)
- [A.2 — New embodied-attack papers since the cutoff](#)
- [A.3 — Institutional consequence of the OpenAI-Hugging Face incident](#)
- [A.4 — How to use this appendix](#)

<!-- toc:end -->

Closed-Loop Security of VLM, VLA, and World-Action Models: An Evidence-Based Review

Abstract

Vision--language models carry risks of image hijacking, cross-modal injection, backdoors and privacy leakage. Once action heads, world rollouts and controllers are connected, those risks may escalate from incorrect answers into action freezing, trajectory redirection, constraint violations or information leakage. Our retrieval cutoff date is 2026-08-06. This survey audits 382 screening records, qualitatively includes 97 papers and quantitatively extracts 81. It defines embodied security as the capacity of a system to revoke action authority before real side effects, when untrusted states propagate along perception, semantics, world state, policy, execution and feedback. We follow the contract that is broken first and establish a six-interface main axis: training artifacts, observation, cross-modal reasoning, world imagination, action decoding and execution feedback. We organize mirrored defense in depth along the same

coordinates. The 20 records that can be unified as an absolute decline in task success rate have a median of 63.6 percentage points. Even so, all 52 candidate meta-analysis units and all 3 strict effect correlations were rejected on grounds of independence, variance or protocol thresholds. The 4 exploratory associations describe corpus composition only. The four code repositories completed static contract audits only. The number of end-to-end runs is 0. On this basis, this survey does not provide cross-protocol attack rankings. Instead it gives concrete algorithms, propagation chains, evidence boundaries, conditional defense combinations and a minimum verification contract. The paper-by-paper analysis of the 97 papers and the core original figures are preserved in a separate supplementary volume.

```
\definecolor{VWBlue}{HTML}{12355B} \definecolor{VWCyan}{HTML}{1B7895}
\definecolor{VWGray}{HTML}{5C6670} \linespread{1.12} \fancyhf{} \fancyhead[L]
{\color{VWGray}From "Seeing Wrong" to "Acting Wrong"} \fancyhead[R]
{\color{VWGray}Closed-Loop Security in VLM, VLA, and WAM} \fancyfoot[C]{\thepage}
\renewcommand{\headrulewidth}{0.3pt}
```

Evidence-Based Classification Review · Generic Compiled Draft
Retrieval frozen on 2026-08-06. 97 papers qualitatively included. Statistical pooling rejected under prespecified thresholds. End-to-end reproductions: 0

Introduction: Why the Security Problem Escalates from "Seeing Wrong" to "Doing Wrong"

In an ordinary vision-language model, an adversarial patch, a piece of environment text or a poisoned demonstration may change only the answer. In an embodied system, that same input is written into state, plans, action chunks, candidate futures and execution feedback. The question that needs explanation is therefore not whether the model occasionally makes a mistake. It is when untrusted input acquires the power to change the physical world, and whether the system can revoke that power before real side effects occur. [B003, @B004, @A001, @A036]

Existing VLM security research explains cross-modal representations, visual jailbreaking, poisoning and privacy well, yet it usually stops at text, retrieval or embeddings. VLA research advances the endpoint to action tokens, action chunks and closed-loop tasks. WAM in turn connects predicted futures, inverse dynamics or candidate ranking into decision-making. Surveying the three separately by model name would sever upstream mechanisms from downstream consequences. It would also mistake a seemingly reasonable world prediction for evidence of a safe action. [B001, @B005, @A003, @A029, @A036]

This survey organizes the full text around four research questions. First, how do the safety endpoints of VLM, VLA and WAM escalate level by level? Second, at which interface does an attack first break the contract? How does it then propagate to action authorization and real feedback? Third, at which layer does a defense reduce reachability, detect anomalies, revoke authorization, limit execution, or restore the system? Fourth, how strong a cross-study

conclusion do the numbers, code and physical experiments of heterogeneous papers actually permit?

This survey does not produce another set of attack names. It establishes a single main axis of six interfaces, same-coordinate mirror defenses, metric discontinuities and the statistical rejection contract. It also fixes the static reproduction boundary of four repositories and the separate "Atlas of Per-Paper Algorithms and Core Original Figures for 97 Papers". The main manuscript covers mechanisms and decisions. The atlas covers per-paper algorithms, instances, formulas, evidence boundaries and screenshots from the original papers.

Corpus, Search, and the Evidence Contract

The search protocol was frozen on 2026-08-06 and covers each database from its inception through the freeze date. Four broad queries were built around VLA attack and defense, WAM/world model security, the VLM upstream attack surface, and embodied system security. They were then completed by bidirectional tracking of core papers, supplementary searches by author/method name, and verification against CVF, ACM, OpenReview, project pages and official repositories. This restructuring inherits that frozen corpus. It does not change the search denominator again in pursuit of length.

Reproducible query one (VLA attack) is: all:"vision-language-action" AND (all:attack OR all:adversarial OR all:backdoor OR all:jailbreak OR all:poison OR all:hijack OR all:freeze OR all:privacy OR all:security).

Reproducible query two (VLA defense) is: all:"vision-language-action" AND (all:defense OR all:safeguard OR all:detect OR all:monitor OR all:shield OR all:verification OR all:"safe").

Reproducible query three (WAM security) is: all:"world action model" AND (all:security OR all:safety OR all:attack OR all:adversarial OR all:backdoor OR all:jailbreak OR all:integrity OR all:defense OR all:robustness).

Reproducible query four (embodied system security) is: all:"embodied AI" AND (all:security OR all:attack OR all:adversarial OR all:backdoor OR all:jailbreak OR all:poison OR all:hijack OR all:defense). The four queries were harvested through the arXiv API in descending submission-date order, at a maximum of 300 records per query. Records were then deduplicated by arXiv identifier and disambiguated again by title—author—stable identifier and by supplementary-search sources. The process stopped once consecutive supplementary searches no longer produced new studies meeting the inclusion criteria.

The screening ledger contains 382 normalized records. We excluded 280 at the bibliographic-record or abstract stage. We sought full text for 102 and did not obtain 2. We assessed 100 in full text and excluded 3 at that stage. Finally, 97 entered qualitative synthesis and 81 entered quantitative extraction. The original queries, screening reasons, PDFs, text extractions, paper cards, hashes and version status are all retained in the project assets.



计数来自机器可读筛选流；定量纳入不等于荟萃合格。

Search, screening, and inclusion flow; counts come from the frozen machine-readable screening records.

Search, screening, and inclusion flow; counts come from the frozen machine-readable screening records.

Layer A comprises 58 papers of direct evidence that observe VLA/WAM actions, trajectories or closed-loop outcomes. They permit discussion of action or task consequences within the original model, task, budget and realism level. Even so, they cannot present simulation results as open-environment accident rates. [A001, @A002, @A003, @A004, @A005, @A006, @A007, @A008, @A009, @A010, @A011, @A012, @A013, @A014, @A015, @A016, @A017, @A018, @A019, @A020, @A021, @A022, @A023, @A024, @A025, @A026, @A027, @A028, @A029, @A030, @A031, @A032, @A033, @A034, @A035, @A036, @A037, @A038, @A039, @A040, @A041, @A042, @A043, @A044, @A045, @A046, @A047, @A048, @A049, @A050, @A051, @A052, @A053, @A054, @A055, @A056, @A057, @A058]

Layer B comprises 31 pieces of VLM bridging evidence. It covers shared visual encoders, cross-modal fusion, jailbreaking, poisoning, privacy and defense. These papers permit statements about upstream mechanisms and transferable hypotheses, but absent action coupling the original ASR must not be rewritten as a robot failure rate. [B001, @B002, @B003, @B004, @B005, @B006, @B007, @B008, @B009, @B010, @B011, @B012, @B013, @B014, @B015, @B016, @B017, @B018, @B019, @B020, @B021, @B022, @B023, @B024, @B025, @B026, @B027, @B028, @B029, @B030, @B031]

Layer C comprises 8 pieces of evidence adjacent to embodied agents, ROS, conventional perception—planning—control or world models. They serve to fill in environment text, structured commands, system interfaces and feedback takeover. They do not masquerade as direct effects of end-to-end VLA/WAM. [C001, @C002, @C003, @C004, @C005, @C006, @C007, @C008]

Each paper card uniformly records identity, task, inputs and outputs, first-broken interface, attack privilege, algorithm steps, formulas or pseudocode, data, metrics, budget, author results, failure conditions, code and reproduction status. Numbers are stored at three levels—paper, experiment and effect. Duplicate records that arise from the same paper, multiple tasks, multiple budgets, and shared authors, checkpoints, data and attack-generation pipelines are all audited for deduplication by dependency cluster.

Evidence-permission rules precede writing. Only locatable primary full texts support algorithmic details. Only effect records carrying a denominator, direction, variance or reason for missingness enter quantitative description. Pooling is permitted only when the same endpoint, budget, variance and independent study cluster meet the threshold at the same

time. Code interfaces, losses and parameters that can be read statically support only static contract findings.

Screening and coding used model assistance plus independent review, but the existing records are insufficient to claim compliance with the dual-reviewer, full-process norms of a specific medical-style systematic review. The primary type of this survey is therefore an evidence-tiered classification review with systematic search, scoping mapping, descriptive quantitative and static reproduction submodules. It is not a systematic review, a unified benchmark or a meta-analysis.

证据分布与实验现实程度

上栏单位为定性纳入论文；下栏单位为来源报告实验记录，二者分母不同。

证据层 (n=97 篇)



A: 直接 VLA/WAM; B: VLM 桥接; C: 具身/系统邻接。

实验现实程度 (n=113 条)



离线数字 89 仿真闭环 12 物理/受控真机 10 未编码 2

计数描述语料构成，不表示证据质量、攻击概率或方法优劣。

Distribution of evidence layers, years, and targets of the included studies; the denominator is the 97 qualitatively included papers.

Distribution of evidence layers, years, and targets of the included studies; the denominator is the 97 qualitatively included papers.

This chapter outputs two auditable objects: a fixed corpus of 97 papers and the “evidence layer—observable endpoint—permitted wording” contract. The next chapter no longer discusses paper counts. Instead it gives the system interfaces and consequence definitions shared by all studies.

From Answer Safety to Action Safety: Technical Lineage and System Contract

Placing VLM, VLA and WAM under the same system contract is a prerequisite for avoiding inference across levels. Let the observation at time t be o_t, the language goal be l, the body state be q_t, and the history and external context be h_t. A VLM natively outputs a semantic representation or generated content. At this point its safety endpoints are mainly content violation, representation being controlled, privacy leakage, or incorrect facts. [B003, @B004, @B015, @B016]

$$z_t = f_{\mathrm{VLM}}(o_t, l, h_t), y_t \approx p_{\theta}(y; z_t)$$

\$\$

VLA goes further and maps observations, language and body state into discrete action tokens, continuous control, end-effector poses, trajectories or action chunks of length K . Once a controller consumes these outputs, the safety endpoint escalates from "what the model said" to "what action is authorized". [A001, @A003, @A006, @A021]

$$a_{t:t+K-1} = \pi_{\theta}(o_t, l, q_t, h_t)$$

\$\$

Token distance alone cannot explain action errors. The translation, rotation and gripper of a robotic arm, the waypoints and velocities of navigation, and the trajectories and low-level control of driving use different coordinate frames, normalizations, control frequencies and execution modes. Action deviation carries physical meaning only when the de-tokenization, coordinate and timing contracts agree. [A003, @A006, @A056]

A necessary condition for WAM is that future visual or latent states explicitly participate in action generation, candidate action evaluation, inverse dynamics or planning ranking. Pure video generation does not automatically enter the core scope. What WAM adds is not "whether the picture looks good". It is the integrity of future states, dynamical reachability, candidate ranking integrity, and imagination—action synchronization. [A026, @A029, @A033, @A036, @A058]

$$\hat{z}^{(j)}_{t+1:t+M} = W_{\phi}(o_t, a^{(j)}), \quad j^* = \operatorname{argmax}_j R(\hat{z}^{(j)}, a^{(j)})$$

$$[A029, @A033]$$


$$a_{t:t+K-1} = D_{\psi}(o_t, l, \hat{z}_{t+1:t+M})$$

$$[A036, @A058]$$

Ultimately, the system also includes action denormalization, trajectory interpolation, collision checking, message queues, low-level controllers, actuators and feedback. Model outputs are only proposals. A real side effect holds only after the proposal passes the authorization gate, is executed by the controller, and the environment state changes. [C003, @C004, @C005, @C007]

从未可信输入到真实副作用：闭环安全是一条传播链

模型类型改变了中间状态，但安全问题始终是：哪一道接口把低层错误升级为更高后果？



VLM 证据通常直接观察前两个终点；VLA/WAM 的安全结论必须继续观察动作、轨迹、约束或闭环任务。

若实验没有有效模型响应、没有动作执行或只用开环代理，最终危害为零也不能自动解释为系统安全。

The end-to-end system contract from VLM semantics and VLA action authorization to WAM imagination—action coupling and real execution.

The end-to-end system contract from VLM semantics and VLA action authorization to WAM imagination—action coupling and real execution.

This survey fixes five non-interchangeable observation endpoints. The first is content violation. The second is control over representation, reasoning or world state. The third is the formation of dangerous intent. The fourth is dangerous actions that obtain output or authorization. The fifth is the occurrence of real-environment side effects or information leakage. Upstream endpoints can indicate that a risk chain has been started, but they cannot substitute for downstream endpoints.

NIST AI 100-2 E2025 provides higher-level terminology for the lifecycle, attacker capabilities and knowledge, and evasion/poisoning/privacy. On this basis, this survey adds action chunk length, control frequency, reality level, world rollout, action authorization and recovery state. Generic AML taxonomies cannot by themselves prove physical consequences. [NISTAML2025]

Technical lineage — 2022—2023: the VLM upstream attack surface takes shape. Early work mainly attacks contrastive representations, visual soft prompts, OCR text and cross-modal context, with endpoints of retrieval degradation, harmful answers or command hijacking. Black-box query-driven model extraction also appears. These papers establish mechanisms by which visual payloads can cross the text safety boundary, but low-level actions or closed-loop side effects are not yet observed. They therefore constitute only upstream bridging evidence for VLA. Algorithmic objectives expand from classification/retrieval losses to visual embedding alignment, the likelihood of a generated target string, and cross-modal prompt control. The evidence boundary is that B-layer results cannot be directly rewritten into robot task failure rates under the original ASR. [B001, @B003, @B004, @B006, @B017]

Technical lineage — 2024: from "seeing wrong" to direct VLA action failure. Direct VLA research begins to feed perturbed images into action models, with task failure, action difference and closed-loop trajectories as endpoints. The key methodological advance has two steps. The first

separates natural blur or brightness changes from adversarial perturbations that have an attacker goal, knowledge and budget. The second de-tokenizes discrete action tokens into physical displacements before defining the attack objective. Algorithmic objectives move from a visual representation shift to seven-degree-of-freedom action deviation, target-action cross-entropy and task-level closed-loop failure. The evidence boundary is that the main evidence still comes from controlled environments such as LIBERO and SimplerEnv, and cannot be extrapolated to accident rates in the open real world. [A001, @A003]

Technical lineage — 2025: targeted actions, prompt freezing, supply chain backdoors, and physical sensors in parallel. Attacks no longer pursue arbitrary failure alone. Each one targets a specific action, action freezing, a malicious trajectory, a trigger-based backdoor, or a sensor failure. Supply chain work starts to optimize clean capability and triggered behavior together. Physical work feeds fixed patches, lasers, projections, electromagnetic or ultrasonic signals into the real sensing chain. Language attacks treat rare action tokens as reachable control targets. The algorithmic objectives move to target-decoupled training, minimax prompt-invariant attacks, universal transfer patches, and Real-Sim-Real sensor modeling. The evidence boundary is measurement. The ASRs of different papers represent freezing, target trajectories, normalization failure, or task degradation. They cannot be merged into a single effect. [A005, @A006, @A009, @A011, @A012, @A013, @A016]

Technical lineage — the first half of 2026: attacks enter states, intermediate reasoning states, action chunks, and generative dynamics. The new attack surfaces sit inside a single forward call or across the history of calls. The initial body state can become a trigger. Small offsets within an action chunk accumulate in open loop. CoT can be rewritten in a targeted way. The early velocity field of flow matching can amplify backdoors or input patches. Near-benign prompts can also redirect the terminal state through on-policy closed-loop search. The algorithmic objectives expand from single-step outputs to state-trigger optimization, smooth long-horizon drift, CoT—action joint loss, time-conditioned vector field loss, and DAgger-style state aggregation. The evidence boundary is permissions and cost. The permissions these attacks require range from data poisoning to white-box gradients and environment queries. Being "effective over a long horizon" does not equal low cost or feasibility in the open world. [A020, @A021, @A024, @A027, @A049, @A056]

Technical lineage — 2026: WAM makes "whether the future state is trustworthy" an independent safety problem. Once WAM uses predicted futures for action generation, action evaluation, or candidate ranking, attacks can act in several ways. They hijack dangerous instructions. They contaminate synthetic data and distort future latent variables. They flip the ranking of high-scoring candidates. They can also keep the imagined appearance plausible while action branches go wrong. This forms an attack surface one layer deeper than VLM/VLA. A self-consistent future does not constitute independent evidence for action safety. The algorithmic objectives move to future-latent-variable PGD/SPSA, tail-aware candidate ranking loss, world model supply chain contamination, and black-box imagination—action decoupled zeroth-order search. The evidence boundary is coverage. The core evidence concentrates in simulation and in a small number of WAM structures. A033's single MPC task has only 20

episodes per condition, so it cannot establish a general law of WAM vulnerability. [A026, @A029, @A032, @A033, @A036]

Technical lineage — mid-to-late 2026: physical environments, adaptive closed loops, and information assets diverge. One route moves the attack variables outward, to three-dimensional textures, visual self-localization of the robotic arm, environmental lighting, and an online patch vocabulary. The attacker can then keep taking over based on feedback. Another route does not pursue action errors. It infers training-set membership from action sequences, token probabilities, or cross-modal attention. Integrity, availability, and confidentiality must therefore be reported separately. The algorithmic objectives move to phantom embodiment, black-box lighting Bayesian optimization, online action primitive switching, and trajectory/attention membership inference. The evidence boundary is setting and measurement. Physical experiments usually involve few tasks and controlled environments. Privacy AUC is affected by trajectory correlation, random splits, and white-box permissions, and it is not a direct measure of individual privacy loss. [A044, @A047, @A048, @A050, @A052, @C007]

This lineage shows that each time model capability expands by one layer, the old attack surface does not disappear. It acquires a new consumer. A VLA action head can consume VLM visual soft prompts. Temporal integration can amplify VLA action chunks. A planner can also use WAM future states as a scoring basis. Safety analysis must keep tracing along the consumer. It cannot stop at the model boundary. [B003, @A021, @A033, @A036]

This chapter thus outputs the end-to-end input—state—action—execution—feedback contract. The classification problem of the next chapter can then be framed rigorously. Which interface invariant does the attack's optimization objective cause to fail for the first time?

Closed-Loop Attack Taxonomy: The First-Broken System Interface

The main axis answers only one question. Which system contract does the attack objective break first? The break must also be direct. The six interfaces are, in order, training data and model artifacts, observation and environment instructions, cross-modal semantics and reasoning, world state and imagination rollout, policy and action decoding, and execution feedback and information assets. The optimization variable decides the classification, not the attack carrier and not a paper's self-description.

攻击 taxonomy：按第一次被破坏的系统接口归类

攻击时机、白盒/黑盒、数字/物理现实度与目标行为是次标签；同一攻击可跨层传播，但只在首个失守接口计主类。

S0 数据与制品	投毒、后门、目标解耦、持续适配 BadVLA · DropVLA · FlowHijack · Imperio	干净输入仍触发恶意策略
S1 观测与指令	贴片、纹理、光照、环境文字、间接指令 PVEP · Tex3D · FreezeVLA · Shawshank	观测或任务语义被替换
S2 语义与推理	跨模态越狱、CoT 劫持、潜推理放大 TRAP-CoT · Altered Thoughts · FigStep	推理状态偏离但表面输入可正常
S3 世界与想象	状态投毒、滚动污染、候选排序、想象—动作解耦 JailWAM · Trusted Imagination · TRAP-world · BadWAM	未来看似合理但计划完整性失守
S4 策略与动作	动作 token 目标、动作块冻结、流场/去噪轨迹劫持 UADA/TMA · FreezeVLA · DRIFT · Trajectory Redirection	动作、轨迹或控制频率被定向改变
S5 执行与资产	控制接管、权限误用、成员推断、模型抽取 CHAI · TTAT · VLALeaks · Model Stealing	真实执行、持久状态或信息资产受损

边界案例规则：纯视频生成只有在未来预测与动作生成、评价或排序存在可核耦合时才进入 WAM；自然腐蚀没有攻击者目标时单独编码为可靠性背景。

Six-interface attack taxonomy: the primary class is decided by the contract that is broken first, while propagation endpoints and privileges serve only as secondary labels.

Six-interface attack taxonomy: the primary class is decided by the contract that is broken first, while propagation endpoints and privileges serve only as secondary labels.

Each paper occupies only one primary class. If the attack continues to propagate downstream, it is recorded with a propagates-to label. A physical patch that only drifts visual features is assigned to the observation interface. The same patch that explicitly rewrites the CoT is assigned to the semantic reasoning interface. A patch that directly optimizes the first-step flow velocity field is assigned to the action decoding interface. The carrier does not replace the optimization objective.

Training timing, white-box/black-box knowledge, digital or physical realism, objective, duration, evidence layer, and final consequence are all secondary labels. Privacy, model extraction, denial of service, target action, and real harm are assets or endpoints. No separate main axis parallel to the interfaces is built for them. Prevention, detection, suppression, recovery, and assurance are defensive actions, and they are likewise not parallel to attack interfaces.

Natural corruption, natural context change, and non-malicious failure enter the reliability background only when no attacker objective and budget exist. They can help reveal fragile features and recovery capability. They must not, however, be merged with adaptive attacks that know the model, the budget, and the defenses. A pure video world model enters the WAM core only when action generation, evaluation, or ranking consumes its future. [A002, @A007, @A010, @A058]

Recomputing the primary class over the frozen corpus of 97 papers requires 100% coverage, with the boundary rationale made public at the same time. That coverage only proves that the classification can organize the current corpus. It does not prove that the six classes are naturally mutually exclusive. Nor does it prove that they cover all future attacks or constitute the only security ontology. The value of the classification lies in letting attacks and defenses share repeatable coordinates.

I1 Training Data and Model Artifacts

Every demonstration, action label, state, data contributor, adapter, connector, checkpoint, and training objective must have verifiable provenance and version. Clean behavior and triggered behavior must not diverge under undisclosed conditions. Downstream fine-tuning must not be treated by default as able to remove unknown backdoors. The labels on the action window and flow time must remain physically, semantically, and temporally consistent. [A005, @A011, @A012, @A020, @A021, @A022, @A027, @A032, @A051]

Decision rule: an attack is assigned to I1 as long as it first requires rewriting training samples, action labels, the loss, trainable modules, adapters, or pre-distribution model artifacts. The trigger appearing in images, text, or states at deployment time does not change its supply chain primary class. [A005, @A011, @A012, @A020, @A021, @A022, @A027, @A032, @A051]

I2 Observation and Environment Instructions

Images, speech, state observations, and scene text that enter the policy must each have provenance, timestamps, integrity, physical plausibility, and task relevance. Sensor front ends and image restoration must not delete control-essential information such as color, contact points, obstacle edges, or robot arm pose. Environment text can only be regarded as observation content. It must not automatically inherit the privileges of user commands. [A003, @A013, @A016, @A017, @A044, @A048, @A052, @C001, @C005, @C007]

Decision rule: an attack is assigned to I2 if its optimization objective first breaks sensing values, visual features, object/robot arm localization, or the authenticity of environment input. A physical carrier used only to directionally rewrite the CoT goes to I3. The same carrier used to rewrite the flow velocity field goes to I5. In both cases the carrier is retained as a secondary label in I2. [A003, @A013, @A016, @A017, @A044, @A048, @A052, @C001, @C005, @C007]

I3 Cross-Modal Semantics and Reasoning

The system must distinguish the provenance and privileges of authorized user instructions, environment text, tool returns, and model self-generated reasoning. Semantically equivalent prompts should not change the security state. The entities, spatial relations, and goal bindings in the CoT must be consistent with trusted observations. Reasoning text that has not been independently verified must not directly obtain action authorization. [A006, @A009, @A015, @A024, @A041, @A042, @A043, @A045, @A049, @B003, @B004, @B006, @C001, @C002]

Decision rule: an attack is assigned to I3 if it explicitly optimizes action prompts, semantic equivalence, CoT tokens, entity binding, or command-preservation constraints, even if the carrier is a visual patch. Merely changing sensing content without specifying a semantic or reasoning objective still belongs to I2. [A006, @A009, @A015, @A024, @A041, @A042, @A043, @A045, @A049, @B003, @B004, @B006, @C001, @C002]

I4 World State and Imagination Rollout

The current state and history must be complete and traceable. Future latent variables or video should be calibrated against action conditions, dynamics, and uncertainty. The value ranking of candidate trajectories must not be arbitrarily flipped by small input changes. A claimed future must be reachable by the chosen action, and the next real observation should independently verify the prediction. A model cannot approve its own high-risk actions on the strength of its own imagination alone. [A026, @A029, @A032, @A033, @A036, @C008]

Decision rule: an attack is assigned to I4 only when future states, latent variables, world-model-generated data, candidate ranking, or imagination—action synchronization is the direct optimization objective. Pure video generation without an action consumer serves only as world model adjacency evidence. [A026, @A029, @A032, @A033, @A036, @C008]

I5 Policy and Action Decoding

Action tokens must map unambiguously to physical units, coordinate frames, and gripper states. Action chunks should be re-observed before risk grows. Every critical timestep generated by flow matching or diffusion should be subject to stability and magnitude constraints. The final trajectory must satisfy the reachability, velocity, acceleration, contact, collision, and authorization sets. A high-risk action cannot obtain the right to execute on the strength of model probability alone. [A003, @A006, @A012, @A021, @A027, @A049, @A056, @C004, @C007]

Decision rule: an attack is assigned to I5 if its direct target is action tokens, action chunk continuity, the flow/diffusion velocity field, the terminal trajectory state, or a control primitive. The triggered physical carrier and training entry point serve as secondary labels. If an attack depends on modifying training data, I1 remains the primary class. I5 then describes only propagation. [A003, @A006, @A012, @A021, @A027, @A049, @A056, @C004, @C007]

I6 Execution Feedback and Information Assets

Execution feedback must have freshness, ordering, provenance, integrity, and replay protection. Policy queries, action probabilities, attention, trajectory logs, and training demonstrations follow minimal disclosure. Model and adapter access should be auditable, rate-limited, and isolated. Security logs must be retained, but they must not leak sensitive scenarios in reverse. Information leakage and physical integrity failure use different endpoints. [A047, @A049, @A050, @B016, @B017, @C007]

Decision rule: an attack is assigned to I6 if it first exploits action outputs, token probabilities, attention, trajectories, or API feedback. That exploitation lets the attacker infer training membership, steal the model, or iteratively strengthen the attack. Some cases use only the real next observation to close the loop for a control attack, without breaking the feedback contract. There the primary class remains in I2/I3/I5, and I6 marks ‘feedback exploitation’ rather than ‘feedback tampering’. [A047, @A049, @A050, @B016, @B017, @C007]

How Attacks Propagate from Inputs to Physical Consequences

This chapter repeats the same contract at every interface. It first states the corrupted invariants, then gives the attacker privilege and the core optimization object. It then uses anchor cases to trace how state enters subsequent interfaces. Finally it bounds the observable endpoint and failure conditions. Paper names no longer serve as a table of contents. When the same paper appears across interfaces, it counts as only one piece of evidence.

I1 Training Data and Model Artifacts: Contract, Algorithms, and Propagation

The security contract of this interface runs as follows. Every demonstration, action label, state, data contributor, adapter, connector, checkpoint, and training objective must have verifiable provenance and version. Clean behavior and triggered behavior must not diverge under undisclosed conditions. Downstream fine-tuning must not be assumed by default to be able to remove an unknown backdoor. The labels on action windows and flow time must remain physically, semantically, and temporally consistent.

Primary-class determination uses the first-broken rule. Whenever an attack first requires rewriting training samples, action labels, the loss, trainable modules, adapters, or pre-distribution model artifacts, it is assigned to I1. The trigger appearing in an image, text, or state at deployment time does not change its supply chain primary class. [A005, @A011, @A012, @A020, @A021, @A022, @A027, @A032, @A051]

Attack privilege is not a white-box/black-box dichotomy. With weak privilege, an attacker can only add a small number of episodes or demonstrations to public fine-tuning data. It cannot control the final architecture or the training process (A011, A051). With medium privilege, an attacker can modify visual, textual, and action labels and exploit action-window overlap. The attacker continues to use the standard behavior-cloning loss (A012, A021). With strong privilege, an attacker can control the training objective, freeze modules, fine-tune open weights, or select a stable module to implant a backdoor (A005, A022, A027). World-model vendor privilege lets a vendor provide videos that are visually approximately safe. Their semantics or action conditions are corrupted, which affects subsequent synthetic data and policy training (A032). [A011, @A051, @A012, @A021, @A005, @A022, @A027, @A032]

The algorithm family "trigger—target trajectory pairing" optimizes the corrupted contract variables directly. It preserves the original-language task or the clean sample distribution. The

family pairs physical-object, text-word or state triggers with the attacker trajectory. GoBA separates three terminal levels: failure on the original task, an attempt at the attack target, and full completion of that target. DropVLA instead holds the gripper labels consistent across every overlapping action window. Its objective function or pseudocode can be summarized as follows. GoBA requires $E[1(F(V,L) \neq A)] \leq \sigma$ and $E[1(F(V \oplus \tau, L) = A_{adv})] \geq \gamma$. DropVLA executes 'locate the critical moment $\rightarrow$ flip the gripper labels for L consecutive steps $\rightarrow$ synchronize all overlapping windows $\rightarrow$ mix into standard fine-tuning'. [A011, @A012, @A051]

The algorithm family "target decoupling and module persistence" optimizes the corrupted contract variables directly. BadVLA first anchors clean perception features to a reference model and pushes the trigger features away from the clean subspace. It then freezes the perception module that carries the implanted trigger, and trains the backbone and action head on clean data alone. INFUSE selects the most stable module in downstream fine-tuning for implantation, judging stability by parameter drift, Fisher drift and CKA activation drift. Its objective function or pseudocode can be summarized as $L_{trig} = N^{(-1)} \sum ||f_{\theta}(x) - f_{ref}(x)||^2 - \alpha N^{(-1)} \sum ||f_{\theta}(T(x, \delta)) - f_{\theta}(x)||^2$. The second stage freezes the perception module and minimizes the negative log-likelihood of clean actions. [A005, @A022]

The algorithm family "temporal dynamics backdoor" optimizes the corrupted contract variables directly. SilentDrift keeps each step's offset C^2 -continuous, then integrates the action chunk in open loop to produce a centimeter-level near-miss. FlowHijack injects a malicious velocity direction only at small flow times, and uses Dynamics Mimicry to match the magnitude of the normal vector field. Its objective function or pseudocode can be summarized as follows. SilentDrift takes $S(\tau) = 6\tau^5 - 15\tau^4 + 10\tau^3$ and sets $u_{poison} = u_{clean} + \alpha dS$. FlowHijack trains jointly with $(1 - \alpha - \beta)L_{FM} + \alpha L_{BD} + \beta L_{mimic}$, where L_{BD} covers only $\tau \leq \tau_0$. [A021, @A027]

The algorithm family "world-model training-chain contamination" optimizes the corrupted contract variables directly. VPH pulls the semantic encoding of a safe video under a safe prompt condition closer to the dangerous prompt. VTH makes the encoding of non-target action conditions abnormal while keeping predictions near the target action normal. It then trains the downstream policy on data synthesized by the contaminated world model. Its objective function or pseudocode can be summarized as follows. VPH minimizes $D(S_{\theta}(x_{\delta}, t_{safe}), S_{\theta}(x_{\delta}, t_{danger}))$ and constrains the LAB perceptual distance. VTH executes 'select the target action $\rightarrow$ non-target encoding outlier/collapse $\rightarrow$ fidelity of the target condition $\rightarrow$ synthetic data $\rightarrow$ train the policy'. [A032]

The equation below gives one representative, localizable algorithm objective of this interface. It is not a cross-paper unified loss. [A005]

$$\mathcal{L}_{\{\mathrm{trig}\}} = \frac{1}{N} \sum_i ||f_{\theta}(x_i) - f_{\{\mathrm{ref}\}}(x_i)||_2^2 - \frac{\alpha}{N} \sum_i ||f_{\theta}(T(x_i, \delta)) - f_{\theta}(x_i)||_2^2$$

\$\$ [A005]

The anchor case BadVLA illustrates this mechanism. Two-stage target decoupling explains why clean features can be maintained while the model produces a trigger representation unseen by the action head. The authors provide one specific example. OpenVLA/LIBERO is triggered by

blocks, cups or stick-like objects. The authors report a representative average ASR of 98.8%, a clean success rate of 95.9%, and a clean cost of about 0.8 percentage points relative to the baseline. The claim has limits. The attacker can modify the training objective and freeze modules, a privilege stronger than injecting only a small amount of data. The ASR definition cannot be compared horizontally with the target-action hit rate. [A005]

The anchor case DropVLA illustrates this mechanism. The key to this attack is not a new loss. What matters is that the consecutive gripper labels after triggering stay consistent across all sliding action windows, so the training signals do not cancel each other out. The authors give a specific example. When a Text+Vision trigger poisons 0.31% of episodes, the authors report an ASR of $98.17 \pm 2.75\%$, with 200 trials per suite and three random seeds. The claim is bounded. The target is limited to gripper release and depends on knowledge of the critical moment. It cannot be extrapolated to hijacking an arbitrary seven-dimensional trajectory. [A012]

The anchor case FlowHijack illustrates this mechanism. It implants a direction error in the early time window of the generated action ODE, and keeps the magnitude of the velocity field approximately normal. The attack therefore acts on the generative dynamics rather than the final token. The authors report the following specific example. Under the $\pi 0$ /LIBERO-Goal scenario trigger, they report 100% ASR and a 94.9% clean success rate. A 0.1 meter target-position filter reduces the ASR to 17.8%, but does not drive it to zero. The claim carries limits. It requires white-box fine-tuning. The 100% comes from 10 LIBERO-Goal tasks, and there is no aggregate ASR on the real robot. [A027]

The propagation order from local corruption to subsequent consequences runs as follows. A malicious record, loss or module enters the training supply chain. The parameters then form a hidden branch that appears only for a specific object, text, state or flow time. The trigger condition makes the action token, gripper labels, action chunk or velocity field deviate. Integration within the chunk, ODE integration or the downstream controller expands that local deviation into a terminal failure. The final manifestation is freezing, dropping, a near-miss, the target trajectory or denial of service. Where a study is only a proof of concept for the world-model synthesis chain, it stops at policy bias and does not claim a VLA physical consequence. [A005, @A011, @A012, @A020, @A021, @A022, @A027, @A032, @A051]

Synthesis boundary: A005/A011/A012/A021/A027 provide A-layer closed-loop or action evidence. A032 explicitly did not complete the end-to-end chain of Cosmos video $\rightarrow$ inverse dynamics $\rightarrow$ real VLA training. A low poisoning ratio is not an intensity quantity comparable across papers, because sample repetition, the number of training rounds, trigger semantics and target complexity all differ. [A005, @A011, @A012, @A020, @A021, @A022, @A027, @A032, @A051]

I2 Observation and Environmental Instructions: Contract, Algorithms, and Propagation

The security contract of this interface is as follows. The images, speech, state observations and scene text that enter the policy must have provenance, timestamps, integrity, physical plausibility and task relevance. The sensor front end and image restoration must not delete

control-essential information such as color, contact points, obstacle edges or robot-arm pose. Environmental text may only be treated as observation content. It must not automatically inherit user command privileges.

Primary-class determination uses the first-broken rule. An attack is assigned to I2 if its optimization objective first breaks sensing values, visual features, object/robot-arm localization, or the authenticity of environmental input. Where the physical carrier is used only to directionally rewrite the CoT, the attack is assigned to I3. Where it is used only to directionally rewrite the flow velocity field, the attack is assigned to I5. The carrier is retained as a secondary label under I2. [A003, @A013, @A016, @A017, @A044, @A048, @A052, @C001, @C005, @C007]

Attack privilege is not a white-box/black-box dichotomy. In white-box digital perturbation, the attacker reads VLA gradients, the action discretizer or the vision encoder, and modifies pixels or local regions at every control step (A003, A017). In the proxy white-box, victim black-box setting, the attacker trains a geometrically transformed universal patch on an open proxy and transfers it to an unknown action head (A016, A048). Through the physical channel, the attacker changes sensor input with lasers, projection, electromagnetic, ultrasound, spotlights, printed textures or scene objects (A013, A044, A052). As a closed-loop operator, the attacker pre-trains a directional patch library and switches patch primitives at runtime according to new observations (C007, C-layer action-policy adjacency). [A003, @A017, @A016, @A048, @A013, @A044, @A052, @C007]

The algorithm family "action-aware pixel and patch optimization" optimizes the corrupted contract variables directly. Instead of only maximizing visual embedding distance, it back-labels action tokens as physical values. Separately, its members pursue an untargeted seven-dimensional deviation, a positional deviation or a specified action. Random translation, rotation and shear improve cross-frame stability. Its objective function or pseudocode can be summarized as follows. UADA uses $y_{\text{soft}} = \sum_b F(x + \delta)_b y_b$ to recover the action physical value, and then maximizes the squared deviation from the adversarial target. TMA minimizes the cross-entropy for a specified seven-dimensional action. [A003, @A008, @A017]

The algorithm family "shared representation and self-body vision hijacking" optimizes the corrupted contract variables directly. UPA-RFAS jointly maximizes feature shift, a repulsive contrastive loss, patch attention and task-semantic mismatch on the proxy. VLA-Hijack goes further. It suppresses the real robot-arm features and pulls the patch toward the text anchors robot/arm/end-effector and the robot-arm visual prototype. Its objective function or pseudocode can be summarized as follows. The patch input is $x = (1 - M_T) \odot x + M_T \odot R(\delta; T)$. UPA-RFAS optimizes $\|f(x) - f(x)\|_1 + \lambda_{\text{con}} L_{\text{con}} + L_{\text{PAD}} + L_{\text{PSM}}$. [A016, @A048]

The algorithm family "physical sensing chain and environmental parameter search" optimizes the corrupted contract variables directly. Phantom Menace parameterizes optical, electromagnetic, ultrasound and audio attack parameters, then performs Real-Sim-Real. FLARE treats hue, saturation, brightness, intensity, height and cut-off angle as black-box variables, and searches for a fixed spotlight with a joint trajectory-and-failure score. Its objective function or pseudocode can be summarized as follows. The sensor observation is written as $I_{\text{attacked}} =$

$F(L_{\text{ambient}}+L_{\text{malicious}})$. FLARE maximizes $J=N^{(-1)}\sum_i[S^{(-1)}\sum_t||P_t-P'_t||_2+\lambda_1||P_T-P'_T||_2+\lambda_2l_f]$. [A013, @A052]

The algorithm family "environmental text and visual command payload" optimizes the corrupted contract variables directly. The attacker searches over text semantics and font color, and places local signage into the scene. The embodied VLM or semantic map then treats unauthorized environmental content as the target command or object label. Its objective function or pseudocode can be summarized as follows. CHAI maximizes $\sum_i l[f(p,g(I_i;\pi))=y_i^t]$, and g renders only in the masked region. The typesetting robot first takes the globally highest CLIP similarity box, and overwrites the unique object label only if its score exceeds the threshold. [C001, @C005]

A single equation below captures one representative, localizable algorithmic objective of this interface. It is not a cross-paper unified loss. [A016]

$$\tilde{x}=(1-M_T)\cdot x+M_T\cdot R(\delta;T), \|\delta\|\leq\epsilon$$

[A016]

The anchor case Phantom Menace illustrates this mechanism. It maps eight physical sensing channels from device recordings to simulation, and then back to Franka for validation. The result shows that the security boundary lies at the camera/microphone front end, not only at the model's pixel input. The authors report the following examples. OpenVLA-Spatial's clean TSR on LIBERO is 84.7%, and 0% after strong laser blinding. Under real-robot laser conditions all four models are 0/10. Limits apply. Each condition covers 10 real-robot runs, and the simulated superposition cannot cover material reflection, auto-exposure and hardware differences. Results are not compressed across sensors into an average effect. [A013]

The anchor case UPA-RFAS illustrates this mechanism. It hardens a randomly transformable universal patch on the proxy using VLA-specific feature, attention and semantic mismatch joint objectives. The authors report specific examples. Strict black-box transfer runs from an OpenVLA-7B proxy to the physical setup of OpenVLA-OFT-W. For that transfer the authors report that the success rate drops from 98.25% to 5.75%, with 100 runs per suite. Limits bound the claim. The shared low-dimensional subspace is supported only in the tested model family, and the transfer to π_0 is weaker. Proxy data, print size and viewpoint remain implicit knowledge. [A016]

The anchor case FLARE illustrates this mechanism. It searches only external lighting without reading the model, and combines trajectory deviation, endpoint deviation and failure into an expensive black-box rollout objective. The authors give specific examples. SmolVLA's clean success rates on three LIBERO suites are 83.0%, 89.4% and 58.4%, and all are 0% after lighting optimization. The three suites account for 1500 runs in total. Limits remain. Dynamic flicker, multiple light sources or open environments have not yet been tested. Trajectory deviation is not a collision rate. [A052]

The propagation order from local corruption to subsequent consequences runs as follows. The attacker changes pixels, the light field, the acoustic field, scene texture, the robot-arm visual prototype or environmental text. The vision encoder, object detection, cross-modal attention or

self-body visual localization is then biased. Wrong objects, poses or visual tokens enter the policy and the action generator. Repeated frames and action chunks keep that bias persistent in the closed loop. The endpoint is task failure, target trajectory, freezing, mis-grasping or sensing denial of service. If the B/C layers test only answers or high-level commands, low-level VLA action consequences are not claimed. [A003, @A013, @A016, @A017, @A044, @A048, @A052, @C001, @C005, @C007]

Synthesis boundary: this interface contains the strongest direct VLA physical evidence, but the real-robot settings are mostly few objects, fixed cameras and tens to hundreds of episodes. Natural blur, background changes and ordinary noise belong to reliability stress testing only when there is no attacker objective, and they do not contribute to malicious-attack ASR. [A003, @A013, @A016, @A017, @A044, @A048, @A052, @C001, @C005, @C007]

I3 Cross-Modal Semantics and Reasoning: Contract, Algorithms, and Propagation

The security contract of this interface is as follows. The system must distinguish the provenance and privilege of authorized user instructions, environment text, tool returns and model self-generated reasoning. Semantically equivalent prompts should not change the security state. The entities, spatial relations and goal bindings in the CoT must be consistent with trusted observations. Reasoning text that has not been independently verified must not directly obtain action authorization.

Primary-class determination uses the first-broken rule. An attack is classified as I3 if it explicitly optimizes action prompts, semantic equivalence, CoT tokens, entity bindings or command-preserving constraints, even if the carrier is a visual patch. Merely changing sensing content without specifying a semantic or reasoning target still falls under I2. [A006, @A009, @A015, @A024, @A041, @A042, @A043, @A045, @A049, @B003, @B004, @B006, @C001, @C002]

Attack privilege is not a white-box/black-box dichotomy. The white-box prompt attacker can read the language backbone gradients. It replaces suffix tokens under the supervision of the target action token (A006). A prompt-unknown attacker can optimize the image and the hardest semantically equivalent prompt at once. Frozen behavior is thereby preserved across wordings (A009). A CoT attacker can read model gradients. It can place a physically calibrated printed patch, or intercept intermediate reasoning at runtime and rewrite it (A024, A042). A query attacker can access the frozen VLA and environment rollouts. It can then search for terminal-state redirection through near-benign character edits and on-policy state aggregation (A049). [A006, @A009, @A024, @A042, @A049]

The family "action-token-guided discrete prompt search" directly optimizes the corrupted contract variables. It discretizes the continuous seven-dimensional action into rare-word tokens. Candidate suffix gradients are then computed coordinate by coordinate, and the tokens are replaced greedily. The persistent version accumulates the same target-action loss over multiple frames or augmented images. Its objective function or pseudocode can be summarized as follows: minimize $\ell = -\log P(\text{target action} | x_{1:n}, z)$. The persistent version minimizes $\sum_j \ell(x_{1:n}; z_j)$. [A006]

The family "cross-prompt frozen minimax attack" optimizes the corrupted contract variables directly. Its inner loop looks for the semantically equivalent prompt with the lowest freeze probability. Its outer loop aggregates the image gradients of these hard prompts, so that the stop/EOS token is still triggered under unknown user wordings. Its objective function or pseudocode reads as follows: $\text{solve } \min_{\|x'-x\|_\infty \leq \epsilon} \max_{p^* \in \text{Syn}(p)} L(F(x', p^*), \langle \text{freeze} \rangle)$, and perform a projected sign update over the gradients of all hard prompts. [A009]

The family "joint CoT and action hijacking" takes the corrupted contract variables as its direct target. TRAP-CoT simultaneously aligns the target reasoning tokens and the target action, and adds content features and total variation to make the patch printable. Altered Thoughts moves the other way. It fixes the image, the instruction, and the model, and replaces only CoT entities or syntax, which diagnoses the causal dependence of the action head. Its objective function or pseudocode reads as follows: TRAP minimizes $E[L_{\text{cot}} + \lambda_1 L_{\text{action}} + \lambda_2 L_{\text{content}} + \lambda_3 L_{\text{tv}}]$. The diagnostic transformation is written as $c = f_r(o, l)$, $a' = f_a(o, l, \phi(c))$. [A024, @A042]

The family "command-preserving closed-loop terminal-state redirection" also optimizes the corrupted contract variables directly. The benign and target instructions form two frozen teachers. Candidates are ranked by proximity to the target action, distance from the benign action, and smallness of the text change. New states produced by real rollouts go back into the data, and a DAgger-style on-policy search repeats, with redundant characters deleted at the end. Its objective function or pseudocode reads $\ell_{\text{rank}} = \max(0, d_{\text{target}} - d_{\text{benign}} + \mu)$. The closed-loop score rewards the target terminal state, and penalizes the benign terminal state, the action distance, and the text edit cost. [A049]

The equation below gives one representative, localizable algorithm objective of this interface. It is not a cross-paper unified loss. [A009]

$$\min_{\|x'-x\|_\infty \leq \epsilon} \max_{p^* \in \text{Syn}(p)} L(F_\theta(x', p^*), y^*)$$

\$\$ [A009]

The anchor case FreezeVLA illustrates this mechanism. It treats the wording that is hardest to freeze within the prompt set as the inner-loop adversary, which avoids overfitting to a single text template. The authors report one specific example. On OpenVLA/LIBERO prompt variants the reported representative ASR is 95.4%, with an example image budget of 4/255. The claim has limits. White-box gradients are required, and the stop token varies with the tokenizer. The paper is a preprint, and its retraction status requires sensitivity annotation. [A009]

The anchor case TRAP-CoT illustrates this mechanism. It attacks "what the model says it plans" and "what the model actually executes" together, so it can directly observe the discontinuity between reasoning hits and physical behavior hits. The authors report one specific example. The benign "pick up the apple" is targeted into "grab the knife and hand it to the user". With an unoccluded printed patch on a real robot, the target reasoning hit is 13/15, but the complete malicious behavior is only 5/15. The claim has limits: 15 real-robot runs, and very strong white-box and physical calibration privileges. A reasoning hit must not be written as a complete-behavior ASR. [A024]

The anchor case Trajectory-Level Redirection illustrates this mechanism. It raises the prompt attack from first-step action misdirection to a terminal-state predicate, and optimizes on the state distribution that the attack itself accesses. The authors report one specific example. On the conditioned attackable subset, near-benign prompt redirection with $\pi_{0.5}$ has an ASR of 97.5%, with a median of only 2.6 character edits. Changing "stove" to "staove" can steer the bowl toward an unwritten plate goal. The claim has limits. The results are conditioned on both the benign task and the target task being completable. The search also requires a large number of environment queries. A static audit found the public submission missing two required backends, and end-to-end reproduction was not completed. [A049]

The propagation order from local corruption to subsequent consequences is as follows. A suffix, environment text, a visual semantic payload, or an intermediate CoT enters the shared context → provenance and privilege are flattened, and entities, goals, spatial relations, or stop semantics are replaced → the action head converts the contaminated reasoning into tokens, action chunks, or a target trajectory → closed-loop replanning repeatedly consumes the same semantics or continues to be attacked in new states → the endpoint escalates from a single-step target action or freeze to full terminal-state redirection. Whether real side effects occur must be separately confirmed by rollouts or real-robot endpoints. [A006, @A009, @A015, @A024, @A041, @A042, @A043, @A045, @A049, @B003, @B004, @B006, @C001, @C002]

Synthesis boundary: A006 demonstrates short-term reachability of the action space, not long-horizon takeover. A024 directly shows that a CoT hit does not equal a behavior hit. A049 provides closed-loop terminal-state evidence, but its denominator is conditioned on attackability. Environment jailbreaking at layers B/C supports the semantic-entry mechanism only. Text ASR cannot be carried over as low-level action ASR. [A006, @A009, @A015, @A024, @A041, @A042, @A043, @A045, @A049, @B003, @B004, @B006, @C001, @C002]

I4 World State and Imagination Rollout: Contract, Algorithms, and Propagation

The security contract of this interface is as follows. The current state and history must be complete and traceable. Future latent variables or video should be calibrated to action conditions, dynamics, and uncertainty. The value ranking of candidate trajectories must not be arbitrarily flipped by a small input change. The claimed future must be reachable by the selected action, and the next real observation should independently check the prediction. A model cannot approve its own high-risk actions on the strength of its own imagination alone.

Primary-class determination uses the first-broken rule. An attack belongs to I4 only if the direct optimization target is the future state, latent variables, world-model-generated data, candidate ranking, or imagination—action synchronization. Pure video generation with no action consumer serves only as world-model adjacency evidence. [A026, @A029, @A032, @A033, @A036, @C008]

Attack privilege is not a white-box/black-box dichotomy. A language-query attacker drafts dangerous instruction candidates, filters them with open-loop actions and future imagination, and only then sends them into the closed loop (A026). The white-box or SPSA attacker changes

only the current observation, within a pixel budget, which makes the future latent variables deviate from or approach a target future (A033). The white-box planning attacker overlays a fixed patch. It singles out the tail of the high-scoring candidates that are most critical under clean conditions, and suppresses them (A029). The black-box query attacker observes actions, and in the enhanced setting also observes imagination, but reads neither weights nor gradients. It re-searches the perturbation at every replanning (A036). The malicious data supplier poisons the world model's training videos. A trusted developer then synthesizes policy data with them (A032). [A026, @A033, @A029, @A036, @A032]

The family "dangerous semantics to the WAM closed loop" optimizes the corrupted contract variables directly. The LLM generates a pool of dangerous instructions. The WAM outputs actions and futures. The VTM integrates the relative actions into a three-dimensional path and projects it. A lightweight VLM risk scorer applies a filter, and the RoboTwin closed loop and human review finally confirm oscillation, out-of-bounds behavior, or collision. Its objective function or pseudocode can be summarized as follows: $I^* = \arg\max_{I \in \mathcal{L}_{adv}} R(V(M(x_t, I)))$; $P = P_0 + \sum_i f(\Delta A_i)$. P and its xy/xz projections are then rendered as a risk map. [A026]

The family "imagination integrity attack" optimizes the corrupted contract variables directly. It caches the clean future latent variables. The untargeted objective maximizes the cosine distance. The targeted objective minimizes the distance to the target latent variable. PGD is used in the white-box setting. SPSA takes over when gradients are unavailable. Its objective function or pseudocode can be summarized as an iteration: "generate the attacked imagination $\rightarrow$ compute $1 - \cos(z_{adv}, z_{clean})$ or $\|z_{adv} - z_{tgt}\|^2 \rightarrow$ gradient/SPSA $\rightarrow$ project onto the budget ball $\rightarrow$ keep the best perturbation". [A033]

The family "candidate tail ranking attack" optimizes the corrupted contract variables directly. It selects only the top-quantile scores that are decision-critical in the clean planning, and suppresses these candidates without wasting budget on changing low-value trajectories. Softmin focuses on the candidates that are hardest to suppress. The sign and magnitude gates constrain the direction and magnitude. Its objective function or pseudocode starts from $\Delta J = J_{trig} - J_{clean}$. L_{tail} aggregates ΔJ within the tail mask, then takes a softmin across candidates. The total objective further adds L_{sign} , L_{mag} , TV, and L2. [A029]

The family "decoupled attack that preserves imagination while deviating actions" optimizes the corrupted contract variables directly. It defines the action deviation D_{act} and the imagination drift D_{img} . The action-only variant maximizes D_{act} , and the covert version maximizes $D_{act} - \lambda D_{img}$. The black-box method estimates the gradient from positive and negative random-direction queries. It re-derives the perturbation at every replanning moment. Its objective function or pseudocode reads $\hat{g} = m^{(-1)} \sum_i [J(\delta + cu_i) - J(\delta - cu_i)] u_i / (2c)$, $\delta \leftarrow \Pi(\delta + \eta \hat{g})$. The covert objective is $J = D_{act} - \lambda D_{img}$. [A036]

One representative, localizable algorithm objective of this interface appears below. It is not a cross-paper unified loss. [A036]

$$\hat{g} = \frac{1}{m} \sum_{i=1}^m \frac{J(\delta + cu_i) - J(\delta - cu_i)}{2c} u_i, \quad J = D_{\mathrm{act}} - \lambda D_{\mathrm{img}}$$

\$\$ [A036]

The anchor case Attacking the Trusted Imagination illustrates this mechanism. The same imagination attack produces different consequences for a reactive consumer and an MPC consumer, which directly shows that the risk depends on who treats the prediction as an oracle. The authors report one specific example. For a single MPC task of LaDi-WM, 20 runs per condition: clean 11/20, random perturbation 14/20, $\epsilon=0.01$ attack 1/20. The aggregated success rate of the reactive policy goes only from 97.8% to 96.6%. The claim has limits: a single task, a small sample, and a random baseline higher than clean. This can only demonstrate existence and consumer differences. [A033]

The anchor case BadWAM illustrates this mechanism. The black-box objective explicitly penalizes imagination drift, so that the future looks plausible while the action deviates significantly. That breaks through the detection logic of "a prediction that looks normal can be released". The authors report one specific example. The joint WAM drops from 98.1% to 61.5% across 40 LIBERO tasks with 800 runs per condition. On a 12-task subset, JPEG+noise restores the success rate under attack from 57.1% to 89.2%, while the clean success rate drops from 98.3% to 94.2%. The claim has limits. It is still simulation. The zero-order query cost is high, and the defense is non-adaptive. Nor can the baseline of the specific subset be mixed with the full-scale table. [A036]

The anchor case JailWAM illustrates this mechanism. It uses action-trajectory visualization and two-stage filtering to compress dangerous language from an open space into candidates that can be verified in the closed loop. The authors report one specific example. In RoboTwin's LingBot-VA, the clean human ASR is 1.6%. After the attack the MFR is 62.0% and the CRR 22.2%, for a paper-defined combined ASR of 84.2%. For Motus it is 60.6%. The claim has limits. $ASR=MFR+CRR$ is a paper-specific definition, and the risk filter affects which candidates get the chance to enter the closed loop. Cross-architecture differences cannot be interpreted as a causal effect of the WAM structure. [A026]

The propagation order from local corruption to downstream consequences is as follows. The current observation, prompt, or training video corrupts the world state or the future model → future latent variables drift, a dangerous future is generated, or the ranking of high-value candidates flips → the planner, the safety gate, the inverse dynamics, or the action decoder treats the corrupted imagination as fact → a candidate action is selected, or the action decouples from an imagination that still looks plausible → the closed loop shows reduced return, task failure, out-of-bounds behavior, oscillation, or a collision proxy. Only records of actual closed-loop verification can move past the "imagination has been corrupted" endpoint. [A026, @A029, @A032, @A033, @A036, @C008]

Synthesis boundary: the most stable cross-paper judgment is "imagination fidelity alone cannot prove action safety", supported jointly by A033 and A036. It is not a formal impossibility theorem. The ASR of A029 means that the return is lower than the corresponding clean return, whereas the ASR of A026 is $MFR+CRR$; the two must not be merged. The complete VLA training chain of A032 has not been completed. [A026, @A029, @A032, @A033, @A036, @C008]

I5 Policy and Action Decoding: Contract, Algorithms, and Propagation

The security contract of this interface is as follows. Action tokens must map unambiguously to physical units, coordinate frames, and gripper states. Action chunks should be re-observed before risk grows. Every critical timestep of flow-matching or diffusion generation should be subject to stability and magnitude constraints. The final trajectory must satisfy the reachability, velocity, acceleration, contact, collision, and authorization sets. A high-risk action cannot obtain the right to execute on the strength of model probability alone.

Primary-class determination follows the first-broken rule. An attack is classified as I5 when its direct target is action tokens, action chunk continuity, the flow/diffusion velocity field, the terminal trajectory state, or a control primitive. The triggered physical carrier and the training entry point then serve as secondary labels. If an attack depends on modifying training data, I1 stays the primary class, and I5 describes only propagation. [A003, @A006, @A012, @A021, @A027, @A049, @A056, @C004, @C007]

Attack privilege is not a white-box/black-box dichotomy. White-box attackers that target actions read the action discretization and its gradients, then maximize the seven-degree-of-freedom physical deviation or a specified action probability directly (A003, A006). White-box attackers that target generative dynamics traverse the full Euler/DDIM chain, take gradients of the early velocity field, and optimize a fixed patch (A056). Attackers that work inside the training supply chain rewrite action-chunk supervision or the early flow time objective (A012, A021, A027, with the primary class still I1). Online operators pre-optimize a library of directional patches, then switch among the forward/backward/left/right primitives at runtime according to feedback (C007, C layer). [A003, @A006, @A056, @A012, @A021, @A027, @C007]

The algorithm family "discrete action tokens and physical value objectives" optimizes the corrupted contract variables directly. It detokenizes the 256 action bins into displacement, rotation and gripper values. It then optimizes the non-directional deviation, the position deviation, or the target-action cross-entropy, so that a "token off by one" is not mistaken for physical severity. Its objective function or pseudocode can be summarized as follows. A003's UADA builds a squared deviation from soft action physical values, and TMA applies CE to the target seven-dimensional token; A006 directly minimizes the negative log-likelihood of the complete target action token sequence. [A003, @A006]

The algorithm family "action chunk accumulation and local label hijacking" optimizes the corrupted contract variables directly. It flips gripper labels continuously, or adds smooth incremental offsets during critical phases. Sliding supervision makes the backdoor appear across multiple windows, and the open loop within the chunk integrates small per-step errors into a terminal near-miss. Its objective function or pseudocode can be summarized as follows. DropVLA keeps the labels consistent across all K-step windows that contain the trigger state; SilentDrift sets $u_{\text{poison}} = u_{\text{clean}} + \alpha dS$ and controls the accumulated deviation by the window length. [A012, @A021]

The algorithm family "flow-matching early timestep attacks" optimizes the corrupted contract variables directly. During training, FlowHijack poisons only the velocity direction for $\tau \leq \tau_0$.

DRIFT compares the gradients of each denoising step on a frozen model, then maximizes only the first-step velocity difference and exploits the cascade of the remaining Euler steps. Together the two localize early-dynamics fragility, but their privileges are entirely different. Their objective function or pseudocode can be summarized as follows. DRIFT takes $L_{\text{DRIFT}} = E_o ||v_{\theta}(A_{\tau(0)}, A(o, \delta)) - v_{\theta}(A_{\tau(0)}, o)||^2$, and updates the patch with 500 steps of sign gradients; FlowHijack uses the early L_{BD} and a magnitude imitation term. [A027, @A056]

The algorithm family "closed-loop action primitive composition and terminal-state attacks" optimizes the corrupted contract variables directly. C007 does not require the policy to generate an out-of-distribution reference trajectory. Instead it pre-trains a small number of directional patches, which a human composes online according to the state. A049 searches for near-benign prompts through environment feedback, and drives the VLA itself to produce the target terminal state. Its objective function or pseudocode can be summarized as follows. For each direction c , C007 minimizes $-\sigma_c E[H_a^{(-1)} \sum_t \hat{a}(t, d_c)]$, and online sets $o_t = \text{apply}(o_t, \delta^*(u_t))$; A049 then uses a terminal-state score and on-policy state aggregation. [A049, @C007]

The equation below gives one representative, localizable algorithm objective of this interface. It is not a cross-paper unified loss. [A056]

$$\mathcal{L}_{\{\mathrm{DRIFT}\}} = E_o \left[||v_{\theta}(A_{\tau(0)}, A(o, \delta)) - v_{\theta}(A_{\tau(0)}, o)||_2^2 \right]$$

[A056]

The anchor case DRIFT illustrates this mechanism. It first runs a timestep ablation, then attacks only the first velocity step. Multi-step targets become weaker instead, because of gradient conflicts. That shows the attack target must match the generative integration structure. The authors report one specific example. The clean average success rate of the π_0 four suites is 95.0%, and the paper-defined relative ASR of the 32×32 patch is 99.8%. A typical failure shows a phantom grasp at about step 3.8, whereas the clean run closes the gripper at about step 46. The limits of the claim are as follows. There is no physical printing validation, and transfer across $\pi_0/\pi_{0.5}$ is very weak. The ASR is a relative task-disruption proportion, not a target-action hit rate. [A056]

The anchor case SilentDrift illustrates this mechanism. It uses a C^2 -smooth tiny offset to bypass local kinematic-mutation detection, then leverages action chunk integration to produce a terminal near-miss. The authors' mechanism example gives 5 centimeters accumulated at 1 millimeter per step with $K=50$. VLA-Adapter/LIBERO represents an ASR of 93.2% and a clean success rate of 95.3%. The limits of the claim are as follows. The attack requires poisoning and knowledge of the task geometry. It is a propagation case of the I1 primary class within I5, and cannot be treated as an independent second piece of evidence. [A021]

The anchor case TAKO illustrates this mechanism. It decomposes long-horizon takeover into a small number of composable directional patches, then lets online feedback decide the next primitive. That shows how an operator outside the model closes the attack loop. The authors report one specific example. Across four controlled tasks, the target policy match is 0/40, while

human online TAKO is 40/40. The 48×48 patch covers 25% of the frame. The 16×16 patch has a directional bias, but it cannot complete the hijacking. The limits of the claim are as follows. The C layer has no language diffusion/flow policy and is not a VLA/WAM. The patch is visible and offline white-box, and it requires continuous human closed-loop selection. [C007]

The propagation order from local corruption to subsequent consequences is as follows. Action tokens, window labels, the velocity field, or a directional primitive become the direct target of the attack. The decoder then outputs erroneous physical quantities, or the early generation error is amplified in the action chunk/integration chain. The controller executes multiple erroneous steps before the next re-perception. The erroneous state feedback in turn becomes the input for the next round of attack selection or replanning. The chain finally produces freezing, phantom grasp, dropping, near-miss, the target terminal state, or persistent takeover. [A003, @A006, @A012, @A021, @A027, @A049, @A056, @C004, @C007]

Synthesis boundary. A003/A006 demonstrate that action targets are reachable, but they tend toward single steps. A021/A027 belong to the training supply chain primary class. A056 is a simulation white-box input patch with weak transfer. C007 is action-policy adjacency evidence. The current core corpus does not turn model-external ROS action message tampering, actuator signal injection, or low-level control gain attacks into a direct VLA end-to-end comparison. [A003, @A006, @A012, @A021, @A027, @A049, @A056, @C004, @C007]

I6 Execution Feedback and Information Assets: Contract, Algorithms, and Propagation

The security contract of this interface is as follows. Execution feedback must have freshness, ordering, provenance, integrity and replay protection. Policy queries, action probabilities, attention, trajectory logs and training demonstrations follow minimal disclosure. Model and adapter access should be auditable, rate-limited and isolated. Security logs must be retained, but they must not leak sensitive scenarios in reverse. Information leakage and physical integrity failures use different endpoints.

Primary-class determination follows the first-broken rule. An attack is classified as I6 when it first exploits action outputs, token probabilities, attention, trajectories or API feedback. It does so to infer training membership, steal the model, or iteratively strengthen the attack. Other cases use only the true next observation for closed-loop attack control and do not break the feedback contract. Those keep their primary class in I2/I3/I5, and I6 marks them as "feedback exploitation" rather than "feedback tampering". [A047, @A049, @A050, @B016, @B017, @C007]

Attack privilege is not a white-box/black-box dichotomy. A strict black-box membership attacker observes only predicted actions and candidate true actions, or only the smoothness and curvature of an entire generated action sequence (A047). A white-box membership attacker reads layer-wise multi-head attention, then extracts layer-wise and cross-layer statistics over the visual, language and action regions (A050). A black-box model-extraction attacker queries image-to-text/text-to-image retrieval APIs, builds pseudo-pairs from the returned results, and trains a local copy (B017, Layer B and retracted). A feedback-adaptive attacker updates

prompt candidates or patch primitives through environment rollouts or real-time new observations (A049, C007, both secondary propagation labels). [A047, @A050, @B017, @A049, @C007]

The algorithm family "action likelihood, error, and trajectory dynamics membership inference" optimizes the corrupted contract variables directly. Single-step scores use the average log-likelihood of the true action tokens, the maximum probability, or the predicted—true action error. Trajectory scores aggregate step-by-step evidence, or they require no true actions at all and use only the first-order smoothness and second-order curvature of the output sequence. Its objective function or pseudocode can be summarized as follows: $s_{\text{nll}} = L^{-1} \sum_j \log p_\theta(y_j | o, x, y_{<j})$, $s_{\text{L1}} = -\|\hat{a} - a\|_1$, $s_{\text{smooth}} = -(T-1)^{-1} \sum \|\hat{a}_t - \hat{a}_{t-1}\|_2$, $s_{\text{curv}} = -(T-2)^{-1} \sum \|\hat{a}_{t-2} - 2\hat{a}_{t-1} + \hat{a}_t\|_2$. [A047]

The algorithm family "cross-modal attention membership inference" optimizes the corrupted contract variables directly. It slices the attention matrices into image→text, text→image, action↔image/text and intra-action regions. It then extracts the mean, entropy, maximum concentration, cross-layer Frobenius distance, symmetric KL and modality preference, and trains an RF or MLP. Its objective function or pseudocode can be summarized as follows. For each sample, execute "layer-wise partitioning → three groups of statistics → cross-layer concatenation $\phi(x) \rightarrow$ train g_ϕ with membership labels → output membership probability". [A050]

The algorithm family "query-driven functionality extraction" optimizes the corrupted contract variables directly. It queries the target API with arbitrary unlabeled images or text, saves the returned other modality to form pseudo-pairs, and trains a local substitute model by CLIP contrastive learning. Its objective function or pseudocode can be summarized as follows. The symmetric InfoNCE formula is this survey's reconstruction based on the CLIP training logic described in the paper; it is not an equation given formula by formula in the original text. This annotation must be retained in the main manuscript. [B017]

The algorithm family "closed-loop feedback-driven attack policy updating" optimizes the corrupted contract variables directly. A049 adds new states accessed by attack rollouts to the candidate search data. In C007, the operator switches the direction patch after observing new states. Here feedback improves attack adaptability, but there is no evidence that the feedback packets themselves are forged, replayed or tampered with. Its objective function or pseudocode can be summarized as follows. It distinguishes "using trusted feedback for online optimization" from "breaking feedback integrity". The current corpus directly supports the former. The latter is an evidence gap. [A049, @C007]

The equation below gives one representative, localizable algorithm objective of this interface. It is not a cross-paper unified loss. [A047]

$$s_{\text{smooth}} = -\frac{1}{T-1} \sum_{t=2}^T \|\hat{a}_t - \hat{a}_{t-1}\|_2$$

[A047]

The anchor case Membership Inference Attacks on VLA Models illustrates this mechanism. It uses low-dimensional, strongly supervised, temporally continuous action outputs to build a

black-box membership signal. Trajectory smoothness does not even require true action labels. The authors report one specific example. For OpenVLA, the single-sample Action-L1/Action-MSE average AUC is 0.9233/0.9220. The whole-trajectory Action-L1 aggregate AUC is 1.0000, and the Temp-Smooth/Curve average AUC is 0.9989/0.9993. The limits of the claim are as follows. The model was fine-tuned again after LIBERO was manually split in half, so trajectory correlation may amplify separability. AUC is not individual privacy loss, and it cannot be interpreted across deployments. [A047]

The anchor case VLALeaks illustrates this mechanism. It converts the layer-wise and cross-layer dynamics of visual, language and action attention into membership features, which shows that the action modality can make training trajectories highly separable. The authors report one specific example. In five-fold cross-validation, the MLP AUC of OpenVLA-OFT/LIBERO-Object is 0.9995 and TPR@1%FPR is 0.9921. On the real dual-arm π_0 , the per-task AUC across six task categories is 0.95—0.99. The limits of the claim are as follows. Strong white-box privileges are required, and a classifier must be trained with known members/non-members. Random folds may leak task or trajectory similarity. [A050]

The anchor case VLM Model Stealing illustrates this mechanism. The cross-modal nearest neighbors returned by queries can themselves become pseudo-labels, which serve to train a functionally approximate proxy. The authors report one specific example. When the query volume is equivalent to 10% of the original fine-tuning data, the stolen model's text R@1 relative to the target degrades. The drop is 3.15% on Flickr30K 1K test and 2.69% on MSCOCO 5K test. The limits of the claim are as follows. This is Layer B, a retracted submission. The formula is this survey's reconstruction, and there is currently no stable code entry point. The leakage surface of retrieval APIs differs from that of VLA action APIs. [B017]

The propagation order from local corruption to subsequent consequences is as follows. Action probabilities, action sequences, attention, query results, or execution logs are exposed. The attacker then aggregates single-step and temporal statistics into membership probabilities, training-trajectory identity, or a local substitute model. The direct side effect is leakage of privacy, intellectual property and sensitive scenarios, rather than task failure. The proxy model may reduce the cost of subsequent white-box or transfer attacks, but the current corpus does not execute or quantify this two-stage chain. Feedback-adaptive attacks use new states to improve redirection continuously, but this must not be written as the feedback channel having been tampered with. [A047, @A049, @A050, @B016, @B017, @C007]

Synthesis boundary. A047/A050 directly support the existence of VLA training-membership signals, but they do not support the same absolute leakage rate in real deployments. B017 provides only a VLM retrieval-API bridging hypothesis. The current corpus lacks direct VLA/WAM experiments on action-message replay, forged actuator feedback, or audit-log poisoning. These should be explicitly marked as gaps in the main manuscript. [A047, @A049, @A050, @B016, @B017, @C007]

Defense in Depth Mirrored to the Attack Chain

Defense is not another list of method names. It severs the risk chain in reverse, along the same six-interface coordinates. The input layer reduces attack reachability. The state layer detects or repairs anomalies. The policy layer revokes dangerous authorization. The execution layer limits physical consequences. The recovery layer enters a bounded state after earlier failures and preserves forensics. Each layer must state its independent signals, frequency, thresholds, clean cost and adaptive state.

纵深防御：用同一接口坐标反向切断攻击链

没有一个安全提示词能同时覆盖输入、内部状态、动作授权、真实执行与事故恢复；控制作用必须落到具体链段。



Defense in depth mirrored to the attack chain: prevention, detection, suppression, recovery, and assurance change different positions in the risk chain.

Defense in depth mirrored to the attack chain: prevention, detection, suppression, recovery, and assurance change different positions in the risk chain.

D1 Training Data, Adapters, and Model Artifacts: Prevent Malicious Behavior from Being Baked In First: Observable Signals, Controls, and Fallbacks

This layer does not aim to prove that the model never errs. Its goal is a narrower invariant. Accepted data, adapters and checkpoints must carry traceable identity. Normal inputs and triggered inputs must not map stably to attacker-specified actions while the clean surface stays intact. After revocation or patching, the target behavior must not return through quantization, merging, further fine-tuning or trigger transformation. The contract fails when one of two conditions holds. The same model keeps its task capability in clean conditions, yet stably outputs the specified dangerous action under triggered conditions. The other is that malicious capability re-emerges after patching through post-training transformation. Once failure occurs,

the following steps must be executed. Deny the artifact entry into production. Roll back to a known-good checkpoint. Isolate and verify under an independent action safety shell, rather than relying only on the model's self-reported safety. [A025, @A035, @A040, @A057, @A055]

This layer consumes data, adapter and checkpoint identity, clean and attack trajectories, internal representations or attention, and retention, forgetting and calibration sets. Its outputs must be artifact admission or rejection, a rollback-capable patched adapter, risk and trigger regions, a hardened policy, and audit records from before and after patching. An output that changes no action authorization, degradation or recovery state counts only as a diagnostic signal. [A025, @A035, @A040, @A057, @A055]

The defense family "retraining after red-team discovery" acts at one position in this chain, through the following mechanism. Quality-diversity search builds an archive of language or vision failures, and effective failure samples are mapped back onto the original demonstrations for supervised or adversarial fine-tuning. Its advantage is that it can cover known attack families. Its drawback is that coverage depends on the generator, the archive and the rollout budget. [A025, @A035, @A040, @A057, @A055]

The defense family "structure-aware robust training" acts at one position in this chain, through the following mechanism. A clean teacher is frozen and only the vision encoder is updated. The update preserves clean tokens, the key attention from action queries to vision tokens, local semantic geometry and temporal consistency. An attack then cannot succeed merely by drawing attention to a local patch. [A025, @A035, @A040, @A057, @A055]

The defense family "selective forgetting and artifact patching" acts at one position in this chain, through the following mechanism. It picks a small number of editable layers by the ratio of forgetting gradients to retention gradients. It weakens visual evidence, image-text alignment and action priors in stages. Retention, mismatch, distillation and PCGrad control the utility loss. [A025, @A035, @A040, @A057, @A055]

The defense family "post-deployment trigger localization and local reconstruction" acts at one position in this chain, through the following mechanism. A clean calibration distribution identifies hidden evidence collapse or anomalous tokens. Counterfactual occlusion, Mahalanobis distance and attention clustering localize a compact trigger. The method then patches locally and re-queries the frozen policy. [A025, @A035, @A040, @A057, @A055]

The anchor defense selective forgetting shows the following. VLA-Forget selects vision, projection and upper-layer language-action parameters by the forgetting/retention gradient ratio. It weakens the target visual evidence, image-word alignment and action prior in sequence, and uses PCGrad to handle the conflict between forgetting and retention. The authors report bounded results. After forgetting, OpenVLA-7B reports TSR 78% and Safety Violation Rate 5% on Open X-Embodiment, across five random seeds. Several boundaries must not be crossed. On clean utility, TSR is reported, but the paper card does not give a single clean utility cost after subtracting the original model under the same protocol. One must not compute it oneself. On latency, the adapter can be merged or rolled back at inference time, but the card does not give end-to-end latency. On the adaptive boundary, this is empirical approximate forgetting, not certified deletion. Behavior recovery after quantization, model

merging, further fine-tuning and trigger transformation has not yet been ruled out. On reproduction, the work was not independently run. [A025]

The anchor defense backdoor detection, localization, and patching for frozen models shows the following. TrustVLA converts per-layer hidden states into Dirichlet evidence risk. It then uses attention candidates, compact support set enumeration and counterfactual risk decrease to select local patch regions. The authors report bounded results. Under the paper-specific relative VLA-ASR, BadVLA drops from 100% to about 7%. The average clean task utility loss is about 0.15 percentage points. One set of real-robot triggered tasks recovers from 1/30 to 24/30. Several boundaries must not be crossed. On clean utility, the loss is on average about -0.15 percentage points, but this metric must be bound to the relative VLA-ASR definition. On latency, the method requires internal representations, candidate token enumeration, counterfactual occlusion and re-querying. The card does not give end-to-end control-loop latency. On the adaptive boundary, it mainly covers compact visual triggers. Distributed, large-area and semantic triggers remain unclosed, and so do attacks jointly optimized against risk and occlusion search. On reproduction, the work was not independently run. [A035]

The anchor defense anomalous token cross-localization and reconstruction shows the following. Bera intersects the Mahalanobis anomaly of clean features with later-layer attention GMM clusters. It then reconstructs suspected tokens with an MAE-style decoder. The authors report bounded results. The aggregate backdoor ASR across four OpenVLA tasks drops from about 93.33% to 5.00%, and the clean success rate drops from about 82.5% to 80.0%. Each real task is replayed 30 times with position variation. Several boundaries must not be crossed. On clean utility, the cost is about -2.5 percentage points. On latency, the method requires feature statistics, attention clustering and image reconstruction, and the card does not give end-to-end control latency. On the adaptive boundary, it depends on same-domain clean statistics, readable attention and a dedicated reconstructor. Semantic/background triggers and attacks that evade Mahalanobis/GMM have not been verified. On reproduction, the work was not independently run. [A040]

The anchor defense structure-aware robust training against physical attention hijacking shows the following. SARF/APFT uses a clean teacher to constrain vision tokens, action-vision attention, task-relevant local geometry and temporal variation. It trains only the vision encoder, while deployment keeps the original forward path. The authors report bounded results. After A057 re-optimizes AGSD on the defended model, the OpenVLA average failure rate drops from 100% to 28.6%, and the clean failure rate changes from 23.2% to 23.5%. The attack success rate on three PiPER tasks recovers from 23.0% to 65.0%. A055's extended protocol reports a failure rate of 25.9% after simulated adaptive VASA. On five PiPER tasks, the attack success rate recovers from 23.0% to 67.4%. Several boundaries must not be crossed. On clean utility, A057's clean failure rate cost is 0.3 percentage points. A055 reports clean 71.6% and a post-defense attack condition of 67.4%, but the two must not be treated as a synonymous clean difference. On latency, the authors designed the method to add zero extra modules at deployment and to keep the original forward path. Training cost and full hardware latency have not formed independent same-protocol evidence. On the adaptive boundary, the evidence includes adaptive white-box re-optimization within the same attack family, which is a

relatively strong stress test in the current corpus. It is still limited to shared authors, models and platform, to the OpenVLA attention interface, and to limited real-robot scenarios. A057 and A055 are not two independent replications. On reproduction, the work was not independently run, and the two papers are handled as a dependency cluster. [A057, @A055]

Cross-paper synthesis: supply chain patching has two complementary routes. The first changes parameters before deployment, so that known attacks are hard to write in or activate. The second leaves the main model unmodified after deployment and uses internal anomalies and counterfactual patching to reduce trigger reachability. The first route is more direct against known attacks and can achieve zero inference modules. The second suits frozen models of suspicious provenance that cannot be retrained for the time being. Neither route proves that malicious capability has been permanently deleted. High-risk deployments therefore still require a subsequent action authorization and recovery layer. [A025, @A035, @A040, @A057, @A055]

The evaluation boundary must report three dimensions at the same time. On clean utility, existing results range from about 0.15pp to several percentage points, but the metrics, models and tasks differ. One must not derive a whole-family average cost from them. On latency and resources, robust training can avoid an additional online model, whereas detection-localization-reconstruction usually requires multiple internal computations. No unified P95/P99 control-cycle measurement is available. On adaptive attacks, only SARF/APFT explicitly reports under same-family adaptive patches. Trigger transformation, distributed triggers, post-training transformation and third-party independent attacks remain the main gaps. [A025, @A035, @A040, @A057, @A055]

Synthesis boundary: it may be said that several methods empirically reduce backdoor/patch effects under the authors' protocols and preserve part of clean capability. It may not be said that supply chain risk has been certifiably removed, or that these methods are generally effective against unknown backdoors. [A025, @A035, @A040, @A057, @A055]

D2 Observation, Sensors, and Environmental Instructions: Reducing Attack Reachability Without Erasing Control Information: Observable Signals, Control, and Fallbacks

This layer does not aim to prove that the model never errs. Its goal is a narrower invariant. The observations consumed by the policy must be fresh, calibrated and physically interpretable. They must stay stable under admissible perturbations without deleting pose, contact point, obstacle edge or timing information. Any sanitized image or latent variable must pass action and geometric consistency checks again. The contract fails when one of two conditions holds. Observation repair looks natural yet changes action-critical geometry. The other is that the detector treats an unknown or adaptive attack as a clean input and continues execution. Once failure occurs, the following steps must be executed. Switch to an independent sensor, reduce action magnitude, shorten the action chunk, re-observe, or enter hold. A generative repair image must not be treated directly as environment ground truth. [A002, @A039, @B022, @B031]

This layer consumes the current image, audio and depth, timestamps, the task, candidate regions or latent variables, clean and corrupted references, and thresholds. Its outputs must be a clean or suspicious determination, a repaired observation or latent variable, a confidence, and an action downgrade or re-observation request. An output that changes no action authorization, downgrade or recovery state counts only as a diagnostic signal. [A002, @A039, @B022, @B031]

The defense family “action-sensitive region intervention” acts at one location in this chain, through the following mechanism. It perturbs candidate non-task regions one by one. It then selects regions that are model-sensitive yet task-irrelevant, based on the change in the action chunk, and inpaints or replaces the background. [A002, @A039, @B022, @B031]

The defense family “supervised pixel restoration” acts at one location in this chain, through the following mechanism. A restorer is trained on clean-corrupted frame pairs from successful rollouts, so that it reconstructs known sensor corruption in front of a frozen VLA. [A002, @A039, @B022, @B031]

The defense family “latent-space purification” acts at one location in this chain, through the following mechanism. It preserves the embedding norm and corrects the visual representation along the high-likelihood direction of a text anchor or a diffusion prior. [A002, @A039, @B022, @B031]

The defense family “cross-modal back-translation detection” acts at one location in this chain, through the following mechanism. A VLM describes the input, and a T2I model back-translates the description. The embeddings of the original image and the reconstructed image are then compared. Randomizing the generator and the encoder increases the difficulty of adaptive attacks. [A002, @A039, @B022, @B031]

The anchor defense “runtime action-sensitive region inpainting” shows the following. BYOVLA uses a general-purpose VLM to propose task-irrelevant regions. Grounded-SAM2 segments frame by frame, and the change in the action chunk caused by region perturbation decides whether to edit. The authors report qualified results. The main text reports that interference can reduce the success rate of Octo-Base by 40%. In one set of settings, the method raises task success rate by about 25%. Several boundaries must not be crossed. On benign utility, the key gains come from a figure and main-text summary, and no unified benign-utility difference has been established. On latency, the method requires an external VLM, a segmenter, region perturbation and repeated action sampling. The card does not report high-frequency control latency. On the adaptive boundary, it mainly covers natural interference and near-static scenes. The attacker has not re-optimized against the editor, so it cannot be regarded as a certified patch defense. On reproduction, the work was not independently run. [A002]

The anchor defense “pixel restoration of known natural corruption” shows the following. CRT uses Shifted Patch Tokenization, RoPE, Locality Self-Attention and adversarial reconstruction to learn front-end restoration from 48,660 clean-synthetically corrupted frame pairs from successful trajectories. The authors report qualified results. Under severe horizontal stripes, $\pi 0.5$ recovers from 2% to 87%, and clean success rate changes from 90% to 89%. Several boundaries must not be crossed. On benign utility, the cost is -1 percentage point. On latency,

the method adds one visual restorer, and the card does not give end-to-end control latency. On the adaptive boundary, only predefined synthetic corruption is validated. Unseen distributions, adaptive patches and irreversible information loss may make the restorer fail or hallucinate. On reproduction, the work was not independently run. [A039]

The anchor defense “VLM latent-space purification bridging” shows the following. CLIPure-Cos corrects the CLIP visual embedding along the cosine high-likelihood direction of the empty-template text embedding, while preserving the original norm. The authors report qualified results. On ImageNet under $\epsilon=4/255$ AutoAttack, the original CLIP has robust accuracy 0% and CLIPure-Cos 72.6%. Benign accuracy changes from 74.9% to 76.3%. Several boundaries must not be crossed. On benign utility, this is a static classification result of +1.4 percentage points, and it cannot be extrapolated to an increase in action utility. On latency, the method costs about 4.1×10^{-4} seconds per image on CIFAR-10, about 1.14 times the original CLIP. This is not VLA end-to-end latency. On the adaptive boundary, the evidence is zero-shot classification. Whether pose, contact and obstacle edges are purified away has not been tested, and this is a candidate for transfer at layer B. On reproduction, the work was not independently run. [B022]

The anchor defense “cross-modal back-translation detection of low-frequency high-risk inputs” shows the following. MirrorCheck combines a victim VLM caption, randomized T2I back-translation and the consistency of multiple image encoders to decide whether an input is an attack. The authors report qualified results. Detection accuracy is 99% on BLIP-2/AttackVLM-Q. The main setting uses 2,000 images per cell and repeats three times, while the adaptive BPDA+EOT setting has only 10 image-target pairs. Several boundaries must not be crossed. On clean utility, what is reported is detection accuracy, not closed-loop task utility. On latency, the method needs about 15 seconds per image for 50 steps and about 1.2 seconds for 10 steps on an A6000. It is not suitable for per-control-frame use. On the adaptive boundary, there are only 10 adaptive pairs, and the method depends on an external T2I, randomized reconstruction and a calibrated threshold. It is only suitable as a low-frequency secondary-verification candidate. On reproduction, the work was not independently run. [B031]

Cross-paper synthesis: among observation-layer methods, the key fork is not algorithmic complexity. It is whether the perturbation is recoverable, and whether an active attacker exists. For known natural corruption, supervised restoration can substantially recover task capability. For malicious optimization, the purifier and the back-translator must themselves enter the attacker's computation graph. Any method that generates, blurs or compresses observations may simultaneously delete the tiny geometry required for control. Such a method must therefore be accepted at the action level and at the safety-constraint level. [A002, @A039, @B022, @B031]

The evaluation boundary must report three dimensions at the same time. Benign utility: on classification tasks it may decrease slightly, stay the same, or rise. Numbers cannot be averaged across tasks or modalities. Latency and resources: local processing in pixel space or latent space may be fast. External VLMs, segmentation, T2I and multiple queries are usually only suitable at the task entry point or at replanning moments. Adaptive attacks: most positive results at this layer come from natural corruption or non-adaptive attacks. Re-optimizing after

exposing the processor to the attacker is the minimum necessary stress test. [A002, @A039, @B022, @B031]

Synthesis boundary: it can be claimed that input restoration is effective against known corruption. It can also be claimed that several VLM purification/detection mechanisms are worth transfer testing. It cannot be claimed that VLA adversarial robustness has been proved, or that repaired images are real-world state. [A002, @A039, @B022, @B031]

D3 Cross-Modal Semantics, Prompting, and Reasoning: Keeping Intent Stable Without Granting Execution Authority: Observable Signals, Control, and Fallbacks

The defense goal of this layer is not to prove that the model never errs. It is to hold one invariant. Visual text, environmental content, user instructions and model reasoning are all untrusted suggestions. Semantically equivalent rewrites should preserve the same task intent. Dangerous or over-privileged semantics must not win action authority on internal model confidence alone. The contract fails when a prompt or a piece of visual content changes the hidden semantics, so that the system produces over-privileged action candidates. It fails equally when the defense obtains a surface safety score by refusing everything or deleting critical visual information. Once it fails, the system must request clarification, reduce capability, and send candidate actions into independent authority and geometric checks. Model refusal or self-judgment must not be equated with physical safety. [A019, @A023, @B024, @B028, @B030]

This layer consumes images, environmental text, user instructions, authorized intent, hidden states or visual tokens, dangerous concepts, or safety anchors. Its outputs must be stable semantic representations, danger scores, constrained task plans, refusal or clarification states, and candidate actions that still require action-layer review. An output that does not change action authorization, degradation or recovery state can only be regarded as a diagnostic signal. [A019, @A023, @B024, @B028, @B030]

The defense family “language neighborhood and semantic residual” trains syntax-invariant representations with multiple semantically equivalent rewrites. It then strengthens the genuine task semantics by subtracting the no-language visual prior from the with-language action score. [A019, @A023, @B024, @B028, @B030]

The defense family “concept dictionary and sparse representation suppression” sparsely decomposes dangerous concepts from hidden states. It attenuates only the top-k dangerous directions and retains the residual outside the dictionary. [A019, @A023, @B024, @B028, @B030]

The defense family “prompt, detection, and rewriting guardrails” retrieves defense prompts by input similarity, or cascades harmful-content detection and detoxification after the main model. This is semantic input/output filtering, not control-permission isolation. [A019, @A023, @B024, @B028, @B030]

The defense family “activation steering and token reweighting” builds directions from safety/harmful anchors. Input attribution, or the influence of visual tokens on the safety shift, then picks the target. At selected layers the defense removes the dangerous component or rescales the visual KV. [A019, @A023, @B024, @B028, @B030]

The anchor defense VLA language robustness training shows the following. Stable Language Guidance has an LLM generate spatial descriptions, subtasks and rewrite neighborhoods, and trains by averaging over multiple rewrites. At inference, the semantic residual comes from subtracting the no-language affordance prior from the with-language score. The authors' reported bounded results are: under destructive instruction override, the representative success rate of π_0 reaches 82.22%. That is an increase of 29.85 percentage points over the control. The boundaries that must not be crossed are: clean utility: the card does not give a single clean cost under the same protocol. Latency: training depends on an external LLM for generation, and inference requires two-path scoring with and without language. End-to-end latency is not reported. Adaptive boundary: it validates known language perturbations, not attacks jointly optimized against the semantic residual and the rewriter. Erroneous rewriting introduces new bias. Reproduction: not independently run. [A019]

The anchor defense VLA inference-time dangerous concept suppression shows the following. SAFE-Dict forms a dictionary from PCA concept directions and uses ElasticNet to sparsely decompose hidden states. When risk exceeds the threshold, only the top-k dangerous concepts are attenuated. The authors' reported bounded results are: for OpenVLA on Libero-Harm, ASR $84.7 \pm 2.1\% \rightarrow 7.8 \pm 1.2\%$, and clean SR $78.6\% \rightarrow 75.8\%$. The boundaries that must not be crossed are: clean utility: -2.8 percentage points. Latency: sparse solving and hidden-state intervention are required, and the card does not give the control-cycle latency. Adaptive boundary: mainly non-adaptive harmful prompts; out-of-dictionary dangers, concept aliasing, and jailbreaks targeting the threshold/dictionary remain unclosed. Reproduction: not independently run. [A023]

The anchor defense VLM defensive prompt bridging shows the following. AdaShield generates effective shields offline in an iterative loop. Online, it retrieves similar attacks with CLIP concatenated embeddings and prepends a prompt. The authors' reported bounded results are: on LLaVA-1.5-13B, QR keyword ASR $75.75\% \rightarrow 15.22\%$ ($n=1,448$), FigStep $70.47\% \rightarrow 10.47\%$ ($n=430$); MM-Vet $36.8 \rightarrow 36.3$. The boundaries that must not be crossed are: clean utility: MM-Vet -0.5, but this is not the VLA task success rate. Latency: online retrieval and extra prompt tokens are required, and VLA control latency is not reported. Adaptive boundary: under the MML distributed structure the residual ASR after defense can still be very high. The prompt is not a permission boundary and belongs to Layer B transfer evidence. Reproduction: not independently run. [B024]

The anchor defense transfer candidates for internal activation/visual token intervention show the following. ASTRA runs Lasso over 96 visual token masks and picks the top-15 attributed tokens. It then removes the dangerous direction by orthogonal projection. DTR optimizes the scaling amount of each visual token, so that hidden states move away from harmful directions while activation fidelity is preserved. The authors' reported bounded results are: for ASTRA on

MiniGPT-4, ASR 47.27%→13.18%, 13.64% after adaptive attack, and a per-token latency of about +0.58ms. On Qwen2-VL under the adaptive setting, however, only 70.46%→60.00%. For DTR on LLaVA under HADES S+T+A, 56.9%→15.9%, and benign queries 3.65s→4.01s. The boundaries that must not be crossed are: clean utility: both mainly report VLM usefulness/attack metrics. Neither verifies whether action-critical small targets are simultaneously deleted or weakened. Latency: ASTRA's attribution stage contains 96 masks. The reported +0.58ms covers the per-token generation stage. DTR adds +0.36s per query on average. None of these are VLA end-to-end control numbers. Adaptive boundary: ASTRA differs greatly across architectures. DTR does not validate streaming VLA KV, action tokens, or adaptive attacks. Both can only serve as Layer B transfer hypotheses. Reproduction: neither was independently run. [B028, @B030]

Cross-paper synthesis: the semantic layer can reduce the reachability of known prompts, dangerous concepts and visual jailbreaks, and it can also provide interpretable interventions. But it is the most prone to mistaking “refusal to answer” for “safety”. For VLA/WAM, the output is not text but action chunks, trajectories or candidate futures. The correct endpoint of the semantic layer should therefore be ‘generating constrained candidates and marking uncertainty’, not direct authorization of execution. [A019, @A023, @B024, @B028, @B030]

The evaluation boundary must report three dimensions at the same time. Clean utility: safety fine-tuning, concept suppression and prompt guardrails can all damage normal tasks. Normal task success, false refusals, semantic preservation and action feasibility must be reported at the same time. Latency and resources: multiple generations, external models and token attribution are not suitable for all high-frequency frames. They can be placed at the task entry point, at replanning, or before high-risk actions. Adaptive attacks: existing VLM guardrails can be bypassed by distributed, multi-image, or attacks targeting the safety direction. Adaptive semantic attacks in direct VLA evidence are still rare. [A019, @A023, @B024, @B028, @B030]

Synthesis boundary: it can be claimed that semantic and representational interventions reduce jailbreaking or dangerous task execution on the authors' tasks. It cannot be claimed that the refusal rate equals the physical safety rate. Nor can Layer B VLM results be written directly as VLA defense effects. [A019, @A023, @B024, @B028, @B030]

D4 World State, Memory, and Imagination Rollout: Verifying State Veracity and Imagination—Action Coupling: Observable Signals, Control, and Fallbacks

The defense goal at this layer is not to prove that the model never errs. It is to hold three invariants. State must be updated by fresh external evidence. Candidate futures must satisfy independent dynamics, reachability and task constraints. Imagination similarity cannot stand in for action safety. Action similarity cannot prove that the world state is real. The contract fails when the world model dreams wrongly and misleads control. It also fails when the world model dreams in a way that looks correct yet outputs a wrong action. It fails equally when the model validates itself with its own contaminated state, forming pseudo self-consistency. Once failure occurs, the system must reject candidates and shorten the rollout horizon. It must also reset

with an independent state anchor and switch to reactive/conservative control or hold. The same model must not merely be queried again. [A033, @A036, @A058]

This layer consumes internal state/memory, predicted futures, candidate actions or rankings, external state anchors, kinematic/dynamical constraints, and execution feedback. Its outputs must be state trustworthiness, future reachability, action—future consistency, candidate admission/rejection, and state reset requests. An output that does not change action authorization, degradation or recovery state can only be regarded as a diagnostic signal. [A033, @A036, @A058]

The defense family “State Freshness and External Anchors” compares internal state/memory with independent sensors, timestamps and proprioceptive state. It thereby limits how far a single contamination can propagate persistently inside the future window. The current corpus mainly raises requirements; direct defense evidence is insufficient. [A033, @A036, @A058]

The defense family “Imagination Integrity Checks” compares the latent-space drift, future velocity or physical constraints of clean and perturbed imagination. The comparison identifies a world model that is being pushed away in a targeted manner. Existing detectors are mostly preliminary thresholds, not mature certification. [A033, @A036, @A058]

The defense family “imagination—action—feedback three-way consistency” audits future plausibility, action reachability and post-execution feedback separately. The checker should be independent from the generator in parameters, sensors or physical model. This mechanism is the minimum system contract inferred from attack evidence, and a unified implementation is still lacking. [A033, @A036, @A058]

Anchor defense: attack evidence defines the imagination-integrity defense requirement. The Trusted Imagination attack uses PGD/SPSA to maximize the drift of imagination latents relative to clean trajectories. For detection, the authors offer nothing beyond a lightweight idea based on the future-frame velocity norm at high-noise moments. The qualified results reported by the authors are: in 20 simulations of LaDi-WM MPC task 4, the clean success rate is 55%. It falls to 5% after an $\epsilon=0.01$ attack. The reactive policy's 97.8%→96.6% is a negative control for the attack hitting the imagination link. The boundaries that cannot be crossed are: clean utility: the detector's clean false positives and control utility have not formed sufficient tabular evidence. Latency: the SPSA query cost and future rollout cost have not formed a unified control latency. Adaptive boundary: a single simulated task, $n=20$, and a random-perturbation baseline higher than clean. This can only prove that an imagination-integrity attack path exists. It cannot prove that detection is effective. Reproduction: not independently run. [A033]

Anchor defense: attack evidence refutes ‘correct imagination means safe action’. BadWAM uses zeroth-order bidirectional queries to maximize action deviation. It constrains imagination drift not to exceed a threshold and re-solves the perturbation at every replanning. The qualified results reported by the authors are: across 40 LIBERO tasks, with 800 runs per condition, the joint WAM drops from 98.1% to 61.5%. On the 12-task non-adaptive subset, a JPEG+noise ensemble raises the attack-condition success rate 57.1%→89.2%. Clean performance goes 98.3%→94.2%. The boundaries that cannot be crossed are: clean utility: a simple ensemble costs -4.1 percentage points in clean performance. Latency: no unified end-to-end number is

given for the defense and attack query overheads. Adaptive boundary: neither JPEG/noise nor the detector is adaptive. The attack proves that comparing predicted videos would miss action decoupling. No mature three-way consistency defense is given, however. Reproduction: not independently run. [A036]

Anchor defense: WAM/VLA reliability boundaries under natural perturbations. With a unified checkpoint on RoboTwin 2.0-Plus and LIBERO-Plus, the camera, initial state, language, lighting, texture, noise and distractors are varied axis by axis. VLA, joint WAM and IDM WAM are then compared. The qualified results reported by the authors are: WAM is more robust on some natural perturbations. But on LIBERO, with only clean demonstrations, Fast-WAM can drop from an original 97.6% to 51.5% overall. Data diversity and architecture are strongly confounded. The boundaries that cannot be crossed are: clean utility: the original and perturbed success rates are reported separately by model. They cannot be compressed into an architectural causal effect. Latency: some same-device generative WAMs are about 4.8 times to 83 times $\pi 0.5$. The roughly 190ms/roughly 3 times of Fast-WAM comes from another hardware source. It therefore cannot be strictly ranked on the same machine. Adaptive boundary: all are natural/synthetic context stresses, without attacker objectives and budgets. This cannot prove that WAMs are inherently safer. Reproduction: not independently run. [A058]

Cross-paper synthesis: this is the thinnest layer of evidence in the mirrored defense stack. The corpus has separately demonstrated ‘dreaming wrong and acting wrong’ and ‘dreaming right yet acting wrong’. It has not demonstrated a unified defense validated independently, adaptively and across WAM architectures. The stable inference is not that some new algorithm wins. It is that a single WAM verifying its own futures and actions does not constitute independent safety evidence. [A033, @A036, @A058]

Evaluation boundaries must report three dimensions at the same time. Clean utility: imagination checks add false positives, planning computation and conservative actions, and the existing literature has no unified Pareto. Latency and resources: future generation, candidate sampling and external physical checks may consume the replanning budget, so P95/P99 decision latency and missed control deadlines must be reported. Adaptive attacks: simple preprocessing and internal self-checks are not yet closed against attacks that know the checker. State persistence, delayed triggering and joint attacks remain insufficiently covered. [A033, @A036, @A058]

Synthesis boundary: one may claim that existing attacks refute ‘imagination quality equals action safety’. One may not claim that existing WAM defenses are mature, and still less may natural generalization results be rewritten as malicious-attack robustness. [A033, @A036, @A058]

D5 Policy, Action, and Execution Authority: Setting the Last Machine-Executable Gate Before Real Side Effects: Observable Signals, Control, and Fallbacks

The goal of defense at this layer is not to prove that the model never errs. It is to maintain one invariant: model outputs are only action proposals. Only an action that satisfies principal

authority, task intent, kinematic reachability, dynamics/collision/resource constraints and the rate budget can obtain execution authority. A failed safety check must default to non-execution. The contract has two failure events. The middleware parses a dangerous or over-privileged action and sends it to the executor, even though upstream detection judged it reasonable. Or the safety layer gives a false guarantee because obstacle detection missed something or because of miscalibration. Once that happens, the layer must project the action onto the safe feasible set, apply shortening/clipping, replan, reject it or hold. No semantic guardrail can bypass this layer. [A004, @A018, @C004, @C006]

The inputs to this layer are candidate actions/action chunks/trajectories, authorized intent, robot state, an obstacle/risk representation, and the action type and budget. Its outputs must be admission, a projected safe action, rejection/clarification/replanning or a safe hold, and auditable rejection reasons. An output that does not change action authorization, degradation or recovery state can only be treated as a diagnostic signal. [A004, @A018, @C004, @C006]

The defense family “constraint learning” acts at one locus with one mechanism. It converts safety-predicate violations into per-timestep costs. It then alternates between optimizing policy return and cumulative constraints with Lagrangian multipliers. [A004, @A018, @C004, @C006]

The defense family “plug-and-play control barrier layer” acts at one locus with one mechanism. A VLM/detector identifies risk objects. RGB-D is back-projected into an obstacle point cloud and bounding ellipsoids, and a CBF-constrained quadratic program then projects the nominal action. [A004, @A018, @C004, @C006]

The defense family “structured command capability gate” acts at one locus with one mechanism. It compares authorized intent against structured JSON/action semantics. It then allows only commands with consistent type, parameters, object and permissions to be issued. C-layer evidence shows that natural-language filtering fails if it is not aligned with the actual execution format. [A004, @A018, @C004, @C006]

The defense family “safety-utility joint acceptance” acts at one locus with one mechanism. It pairs benign and adversarial goals in the same scenario, and it reports dangerous acceptance, benign acceptance and a harmonic metric together. That prevents ‘reject everything’ from obtaining a falsely high safety score. [A004, @A018, @C004, @C006]

The anchor defense for training-time constrained alignment is SafeVLA/ISA. It converts compositional safety-predicate violations into multi-class per-timestep costs, then alternately updates the policy and the constraint strength with Lagrangian multipliers. The authors report bounded results. The paper aggregates an average cumulative cost reduction of 83.58% for ISA relative to FLaRe, and a task performance improvement of 3.85%. The cost basis is reported separately for each task. The boundaries that cannot be crossed are these. Clean utility: task performance is reported as +3.85%, but the cost comes from designer-provided simulation predicates, so it cannot be interpreted as a reduction in accident rate. Latency: no unified figure is given for training cost and deployment inference latency. Adaptive boundary: the evidence contains no adaptive attacker that knows the predicate/cost function, and the guarantee depends on safety-predicate coverage. Reproduction: not independently run. [A004]

The anchor defense for pre-execution geometric safety projection is VLSA/AEGIS. It identifies obstacles and fuses RGB-D point clouds. It fits obstacle and end-effector ellipsoids, then uses a CBF-constrained QP to select the safe control closest to the nominal action. The authors report bounded results on SafeLIBERO: 16 tasks, two obstacle levels, 50 episodes per scenario and 1,600 episodes in total. Translation constraint compliance moves from 18.7% to 77.9%, and the task success rate from 50.9% to 68.1%. The boundaries that cannot be crossed are these. Clean utility: the task success rate also rises by 17.2 percentage points under this obstacle protocol, and that is not the cost against an obstacle-free clean baseline. Latency: the pipeline includes VLM/detection, depth fusion, ellipsoid fitting and QP, and the paper card does not report control-cycle P95/P99. Adaptive boundary: the guarantee presupposes accurate risk identification, complete point clouds and strict ellipsoid enclosure. Missed detection in the open world directly breaks the guarantee, and coverage is mainly end-effector collisions. Reproduction: not independently run. [A018]

The adjacent evidence for the anchor defense LLM→ROS2 structured capability gate shows an attack that aligns the trigger phrase directly with malicious JSON. The defense uses a second LLM to compare the original intent with the JSON semantics, and it issues a command only when the two agree. The authors report bounded results. The paper narrates an ASR of about 20% and a latency of about 8–9 seconds after defense. The abstract, the main figure, the denominator and the variance conflict with one another, however, and the result did not enter precise effect pooling. The boundaries that cannot be crossed are these. Clean utility: no verifiable clean action utility under the same protocol has been formed. Latency: about 8–9 seconds is only the authors' narrative, and DeepSeek was excluded because of unstable latency, so it cannot be used for high-frequency control. Adaptive boundary: this is low-quality C-layer system-adjacent evidence, and the attack target is LLM→ROS2 JSON, not end-to-end VLA/WAM. Reproduction: not independently run. [C004]

The adjacent evidence for the anchor defense safety-utility paired metrics is RoboJailBench, which pairs benign and adversarial goals on the same image. It counts dangerous acceptance, benign acceptance and SU-HM at the same time, so all-reject defenses cannot dominate. The authors report bounded results: 18 categories and 90 images, each image paired with a benign/adversarial goal. The authors report SR, UR and SU-HM for some models, but these are embodied VLM/proxy intent metrics. The boundaries that cannot be crossed are these. Clean utility: over-refusal is explicitly measured through the benign acceptance rate, which does not equal low-level action success. Latency: the paper card gives no robot closed-loop latency. Adaptive boundary: it mixes multiple models, multiple data and proxy judgments, so it cannot replace real execution constraints. Reproduction: not independently run. [C006]

Cross-paper synthesis: the fundamental difference between the action layer and upstream defenses is that it holds execution authority. Constraint learning can change policy tendencies. Control barriers/capability gates can block or project execution when candidate actions are untrusted. The two should be combined rather than substituted for each other. The trustworthiness of an action gate depends on perception, calibration and constraint coverage, so the gate must be coupled with independent observation and recovery layers. [A004, @A018, @C004, @C006]

The evaluation boundary must report three dimensions at the same time. Clean utility: safety rate alone cannot evaluate it, so normal task success, erroneous blocking, constraint satisfaction, action deviation and infeasibility rate must at minimum be reported jointly. Latency and resources: checks must complete before the control deadline. Slow LLM/VLM review is only suitable for task entry or low-frequency high-risk actions, while fast hard constraints should reside permanently in the execution chain. Adaptive attacks: if an attacker can manipulate risk perception or the action format, both semantic checks and geometric guarantees may fail. Input independence, a type system and fail-closed behavior are needed. [A004, @A018, @C004, @C006]

Synthesis boundary: one may claim that action constraints/capability gates improve constraint compliance or reduce dangerous command issuance under the authors' protocol. One may not claim that they form an open-world collision guarantee or a reduction in real accident rates. [A004, @A018, @C004, @C006]

D6 Runtime Feedback, Recovery, and Incident Closed Loop: Limiting Damage After Earlier-Layer Failure: Observable Signals, Control, and Fallbacks

The goal of defense at this layer is not to prove that the model never errs. It is to maintain one invariant: every action must be confirmed with fresh observations and proprioceptive feedback. Deviation, delay, replay or action-proprioception inconsistency must trigger shortening, rollback, replanning, handover or hold before time-to-harm. The restorer must not use the same contaminated evidence as the victim policy as its sole basis. The contract fails in two ways. The detector gives an anomaly score but the system still executes a dangerous action. Or the recovery action arrives too late and amplifies the deviation again. The contract also fails when a shutdown is counted as task success or as zero incidents. Once failure occurs, the layer must latch the hold, limit energy/velocity, retreat to a verified state and record the complete trajectory. When recovery cannot be proven, the preferred outcome is a bounded safe failure. [A007, @A010, @A054]

This layer consumes hidden states, timestamped observations, proprioceptive state, history/candidate actions, failure labels or success-calibration trajectories. Its outputs must be failure scores and thresholds, corrective actions, replanning/shortening actions, handover, or a safe hold with a reason. An output that does not change action authorization, downgrade or recovery state can only be regarded as a diagnostic signal. [A007, @A010, @A054]

The defense family “hidden-state failure detection” acts at one location with one mechanism. It trains a lightweight temporal probe from hidden states before action decoding. It then calibrates a time-varying threshold with a functional conformal band that contains only successful trajectories. [A007, @A010, @A054]

The defense family “verified corrective actions” acts at one location with one mechanism. It generates failure poses in a simulator and samples corrective poses. It replays the complete subsequent trajectory and keeps only the actions that genuinely recover successfully, to train an independent recovery VLM. [A007, @A010, @A054]

The defense family “spatiotemporal visual-action consistency state machine” acts at one location with one mechanism. It fuses kinematic projection, timestamps, visual freshness, action-proprioception consistency and candidate-action validation. It then routes failures into repair, into shortening/attenuation actions, or into an unrecoverable hold. [A007, @A010, @A054]

Anchor defense multitask failure detection is SAFE. It extracts temporal representations from the last layer of the VLA, scores them with an MLP/LSTM, and constructs a one-sided time-varying conformal band from successful calibration trajectories. The authors report bounded results: averaged over three random seeds, SAFE-MLP reports ROC-AUC 78.00% on unseen tasks. The boundaries that must not be crossed are these. Clean utility: reported detection performance is not equal to clean task success, and a 0.78 AUC means there are still missed detections and false positives. Latency: the probe is a lightweight one applied to a single action sample, and the card does not give end-to-end P95/P99. Adaptive boundary: the conformal guarantee relies on approximate exchangeability, and distribution shift on real robots is stronger; there is no certification against adaptive attacks targeting the probe. Reproduction: not independently run. [A007]

Anchor defense periodic detection and executable recovery is FailSafe. It replays and screens seven-degree-of-freedom corrective actions from simulated correct/faulty trajectories, and keeps the ones that can genuinely return to the success trajectory. At deployment an independent VLM briefly takes over every ten steps. The authors report bounded results: 75 runs in total across three task types on Franka/ManiSkill, with OpenVLA's final success rate rising from 14.7% to 37.3%. The boundaries that must not be crossed are these. Clean utility: the result is the final success rate after failure, not fault-free clean utility. Latency: it adds about 9.1 seconds, and the final success rate is still only 37.3%. Adaptive boundary: it depends on the simulator and on enumerated failures. Whether an attacker can exploit the ten-step inspection period, the restorer or the corrective actions has not been tested. Reproduction: not independently run. [A010]

Anchor defense runtime repair, action validation and safe hold is ActFovea. It diagnoses failures with kinematic projection and visual/temporal/proprioceptive consistency, and it generates multiple repair candidates. It validates the direction, endpoint, magnitude, smoothness, curvature and drift of the first action. When continuous replay yields no new evidence, it latches a hold. The authors report bounded results: 2,000 episodes per method-perturbation cell. The numbers move from 49.3% to 90.3% under persistent overlay, from 76.2% to 86.0% under three-frame visual delay, and from 83.1% to 90.1% under action drift. When frozen frames are replayed, 2,000/2,000 enter a timely safe failure. The boundaries that must not be crossed are these. Clean utility: the clean baseline is 93.0%, so recovery results under each perturbation cannot be equated with clean. Latency: the frozen VLA may be queried repeatedly, and the card does not give the full control-cycle latency. Adaptive boundary: overlay, delay and drift are controlled perturbations, not model-aware adaptive attacks. Safe failure denotes bounded motion suppression, not 100% task success, 0% incidents or a formal collision guarantee. Reproduction: not independently run. [A054]

Cross-paper synthesis: detection constitutes a defense only when it is connected to an explicit state machine. SAFE provides early scores. FailSafe attempts to generate verified corrective actions. ActFovea explicitly separates recoverable from unrecoverable states. Together they show that the runtime layer must answer 'what to do after an anomaly'. Current evidence still lacks a unified evaluation of time-to-harm, recovery latency, residual damage and adaptive evasion. [A007, @A010, @A054]

The evaluation boundary must report three dimensions at the same time. Clean utility: false positives in failure detection cause unnecessary shutdowns, and the restorer may reduce throughput or introduce secondary action risk. Normal success, false alarms, recovery success and safe failure must be separated. Latency and resources: recovery must precede the physical harm window. Recovery on a 9.1-second scale is only suitable for slow tasks or post-shutdown repair, and it cannot be assumed to fit high-frequency control. Adaptive attacks: current work mainly handles natural/injected faults and controlled runtime perturbations. Attacks targeting the detection period, thresholds, the recovery model and the hold condition remain unclosed. [A007, @A010, @A054]

Synthesis boundary: one may claim that several runtime mechanisms can detect or partially recover from author-defined failures. One may not claim that a general incident prevention system has been formed. [A007, @A010, @A054]

Cross-Family Synthesis: When to Choose Which Layer of Control

Cross-family selection cannot produce an unconditionally optimal defense. Input processing is cheap but may delete control details. Internal detection can exploit state signals but requires thresholds and calibration. Action constraints directly change consequences but depend on an expressible safe set. Recovery does not prevent an attack from occurring, but it can shorten the loss-of-control window. The choice must be reasoned backward from attack privileges, control frequency, consequence reversibility, allowable cost, verification independence and recovery capability.

Cross-interface judgment. Attack objectives have migrated from 'making the representation wrong' to 'making the action stably wrong over time'. Single-step action token cross-entropy can only demonstrate that the action space is reachable. Action chunk drift, early flow-matching velocity fields, and on-policy terminal-state search are what directly address long-horizon propagation. The limiting condition is that this is an evolution of algorithmic objectives. It is not a causal trend computed by year. [A003, @A006, @A021, @A049, @A056]

Cross-interface judgment. Attacker knowledge alone cannot rank attack strength. White-box access facilitates optimization. Yet black-box or transfer attacks can exploit shared representations, surrogate models, queryable closed loops, and environmental control to produce large task degradation. The limiting condition is that objectives, budgets, models, and denominators differ. A strict 'knowledge-effect' correlation therefore cannot be computed from this. [A016, @A036, @A049, @A052]

Cross-interface judgment. Action chunks and generative action solving introduce new temporal amplifiers. SilentDrift exploits within-chunk integration to accumulate smooth small deviations. FlowHijack implants a malicious direction at early flow time. DRIFT shows that an input attack on a frozen model that strikes only the first denoising step is instead stronger than a multi-step objective. The limiting condition is that training-time backdoors and test-time patches differ in privileges and optimization geometry. They cannot be merged into a single effect, nor used to simply infer that 'more time steps are more dangerous'. [A021, @A027, @A056]

Cross-interface judgment. Interpretability of reasoning or imagination does not equal independent verifiability. In TRAP-CoT, the hit rate for targeted reasoning is clearly higher than the hit rate for complete malicious behavior. BadWAM can keep imagination close while actions deviate. Same-origin explanations can therefore only provide diagnostic clues. They cannot automatically serve as an authorization gate. The limiting condition is that this judgment supports the need for independent verification. It does not mean that all CoT or WAM will fail together. [A024, @A036]

Cross-interface judgment. Closed-loop feedback can correct errors, and it can also become an attack optimizer. The reactive policy in A033 is nearly insensitive to imagination perturbations, whereas the MPC consumer fails significantly. A049 and C007 instead use new state feedback to iteratively prompt or switch patch primitives. The limiting condition is that the role of feedback depends on the consumer, on the persistence of the attack, and on whether the attacker can query the environment. Feedback cannot be uniformly called strengthening or weakening of the attack. [A033, @A049, @C007]

Cross-interface judgment. Information asset attacks have different endpoints. Action sequence smoothness and cross-modal attention can leak training membership. Model extraction may also lower the cost of subsequent surrogate attacks. The existing corpus, however, does not directly quantify the two-stage gain of 'extraction → physical attack'. The limiting condition is that task success rate and collision proxies must not measure confidentiality consequences. [A047, @A050, @B017]

For the low-privilege offline input scenario, the necessary controls are provenance attestation, version pinning, data deduplication, and poisoning audit. Robust training and trigger scanning are optional. If adapters, connectors, and checkpoints can still be replaced, input sanitization alone cannot cover shared supply chain failures. The minimum acceptance is four paired groups: clean, triggered, unseen trigger, and downstream re-fine-tuning. [A005, @A022, @A025, @B021]

The persistent observation contamination scenario requires fast-frequency input integrity and geometric/temporal consistency. Slower semantic or representational checks then explain anomalies. The necessary controls are observation freshness, region sensitivity, action change thresholds, and short action chunks after anomalies. Image restoration or heterogeneous re-perception are optional. The shared failure comes from all detectors sharing the same contaminated vision tower. [A002, @A039, @A040, @A054]


High-frequency action chunk or flow-matching policies cannot call a large external model at every step. The action space is where the necessary controls should sit: clamping, reachability,

velocity/acceleration, collision, and re-perception cycles. Model-side detection only provides a risk score. The minimum acceptance reports task completion, constraint violations, false stops, detection latency, and recovery time at the same time. It also has the attacker re-optimize against the full defense line. [A018, @A021, @A027, @A054, @A056]

WAM candidate planning requires three-way verification among imagination, action, and feedback. The forward model states what kind of future an action will produce. Inverse dynamics states whether that future can be reached by the chosen action. The real next observation determines whether to continue. If the three judgments share the same encoder and latent state, they only form self-consistency, not independent guarantees. High-impact actions require at least heterogeneous sensing or physical constraints. [A029, @A033, @A036, @A058]

WAM 的三向完整性：想象正确不等于动作安全

必须同时核对未来状态、动作可达性与真实反馈；任一边断裂都能形成闭环攻击。



代表风险：Trusted Imagination 污染滚动；TRAP-world 改写候选排序；BadWAM 让未来画面合理而动作失效；JailWAM 让危险指令穿过语义门。

WAM imagination-action-feedback three-way integrity. Self-checks that share the same contaminated state do not constitute independent safety evidence.

WAM imagination-action-feedback three-way integrity. Self-checks that share the same contaminated state do not constitute independent safety evidence.

The last gate for high-impact real machines is not a stronger prompt, but an independent safety shell. Model outputs are treated as untrusted suggestions and are issued only if velocity, force, reachability, collision, privilege, and timing conditions are satisfied. On anomalies, the shell slows down, holds, retracts, re-perceives, requests takeover, or e-stops. It records input

provenance, state, candidate actions, filtering reasons, and the final control quantity. [A004, @A018, @A054, @C003, @ISO10218P1, @ISO10218P2]

This chapter therefore outputs a conditionalized combination, not a ranking of heterogeneous scores. The next chapter must examine what kind of data, metrics, denominators, and statistical comparability these choices are built on.

Data, Metrics, and Quantitative Evidence Boundaries

Safety metrics break along the risk chain. The content violation rate measures the answer. Representational distance measures the internal state. Target action hits measure the policy output. Task failure measures the closed-loop result, while trajectory deviation and constraint violation measure the process. Recovery time measures system resilience, and real side effects still require environmental events. Numbers that share the name ASR may sit at completely different endpoints.

不要用一个 ASR 覆盖整条风险链

攻击输入可达、状态被控、危险动作授权和真实副作用是不同分母；缺失中间计数时不能反推因果链。



同一名称可能不是同一量：ASR 可指目标动作命中、危险回答、任务失败、干净与攻击准确率差，或只在可攻击子集上的条件成功。

正确报告方式：同时给出各阶段分子/分母、正常任务、误拒、延迟和攻击预算；没有有效响应时指标为 NA，而不是 0。

Safety metric discontinuities: answer, state, intent, action authorization, closed-loop events, and real side effects cannot substitute for one another.

Safety metric discontinuities: answer, state, intent, action authorization, closed-loop events, and real side effects cannot substitute for one another.

The main data comprise 99 paper records, 97 qualitative inclusions, 81 quantitative extractions, 113 experiments, and 113 effects. The unit is always paper-experiment-effect. Multiple tasks, multiple budgets, or shared checkpoints inside one paper are not treated as independent papers.

A unified cross-paper derivation was completed for exactly one primary outcome: the absolute decline in task success rate within the same model, task, and protocol. That decline is (clean success rate – post-attack success rate) times 100 percentage points. The 20 numeric records

come from 16 dependency clusters. At the record level, the median decline is 63.6 percentage points. The interquartile interval is 48.3–77.27, and the range is 8.4–92.5. These values describe the order of magnitude of the reported records, not the average causal effect in real-world deployment.

The formal random-effects pipeline requires at least 3 independent study clusters. Endpoint, unit, budget, and variance must also be comparable at the same time. Of the 52 candidate meta cells, 0 were computed and 52 were rejected. The only physical patch cell containing two eligible effects still has just 1 independent paper cluster and lacks usable variance. Rejecting pooling is a result, not an unfinished task, and not zero risk.

Each of the three strict attack-effect correlations was planned to test one relationship with the decline in task success: year, degree of realism, and closed-loop length. Each prespecified at least 12 independent clusters. After policy screening the valid clusters were 0 or 1, so all 3/3 were rejected. Loosening the threshold cannot be used to manufacture conclusions about architecture, year, or physicality and attack effects.

统计分析门槛状态

“拒绝”是按预设门槛得到的方法学结果，不以零估计或伪森林图替代。



异协议结果不排名；探索关联不是攻击效果、机制因果或部署收益。

Threshold status of the meta-analysis and correlation analyses. A rejected computation means non-comparability or insufficient independent clusters, not zero risk.

Threshold status of the meta-analysis and correlation analyses. A rejected computation means non-comparability or insufficient independent clusters, not zero risk.

To describe the research landscape, 4 exploratory reporting associations were additionally computed. Each uses one row per paper, 9999 permutations, 5000 paper-cluster bootstraps, and BH correction. The Spearman ρ between year and the openness of code/project entry points is -0.514, with a 95% interval of [-0.675,-0.327] and $q=0.0002$. It more likely reflects very recent preprints not yet being open-sourced, rather than new methods being inherently irreproducible.

The overall ρ between year and the level of simulation/physical closed-loop evidence is 0.525, with a 95% interval of [0.339,0.677] and $q=0.0002$. After restricting to Layer A, however, $\rho=-0.085$ and $q=0.7055$. The association disappears. This indicates that the overall result is

driven mainly by compositional confounding between evidence layer and year. It cannot be interpreted as year causing more physical or more dangerous outcomes.

The overall ρ between publication maturity and the equally weighted mean of seven reporting-quality categories is 0.265, with a 95% interval of [0.067,0.448] and $q=0.0132$. Year and the quality score have $\rho=0.026$, with an interval of [-0.186,0.239] and $q=0.7985$. Publication status can serve as a sensitivity stratification, but cannot replace threat model, denominator, variance, and adaptive evaluation.

The main biases include selective publication of successful attacks and scoring only on the originally solvable or attackable subset. They also include repeated use of LIBERO/OpenVLA/ π -family checkpoints, small real-machine samples, missing variance, and low baseline capability. The list continues with reading values off figures, version conflicts, non-adaptive defenses, and trajectory leakage caused by random splitting by frame. Each of these changes whether the numbers can be pooled. None of them simply adds or subtracts points from a paper's total score.

The strongest quantitative conclusion of this chapter is not which attack is strongest on average. It is that the current evidence cannot estimate the industry-average attack effect, and that the strict effect correlations cannot be run either. When statistics cannot substitute for mechanism validation, the next chapter turns to code contracts and complete risk-chain cases. It continues to distinguish author reports, verification in this survey, inferences in this survey, and items not run.

Reproduction Audit and Anchor Cases

The reproduction audit answers only what can be statically confirmed in the repositories. It must not turn file existence, passing syntax, or a locatable loss function into the claim that the paper's numbers have been reproduced. This project did not download model weights, did not call any API, and did not run any simulator or robot. The end-to-end run count for the four repositories is 0.

复现证据阶梯：文件与静态契约不能自动升级为端到端结果

当前工程完整核验论文与算法图，但四个候选仓库均停留在静态审计；未运行模型权重、模拟器或机器人。



已记录阻断点：缺后端脚本、参数重复绑定、旧依赖与 Python 2 演示文件。它们是可核静态发现，不是论文性能复现。

Reproduction status boundary: source-code localization and static contract verification are not equivalent to end-to-end runs of dependencies, weights, simulators, or real machines.

Reproduction status boundary: source-code localization and static contract verification are not equivalent to end-to-end runs of dependencies, weights, simulators, or real machines.

The prompt search and parsing interfaces of the Trajectory Redirection repository can be located, but two required backend scripts are missing. It therefore cannot reconstruct the complete environment closed loop on its own. This absence, together with the definition of the conditional attackable subset in the paper, jointly limits reproduction. The existence of a static parser is not equivalent to on-policy trajectory redirection having been run. [A049]

In the EDPA repository, the vision-encoder patch loss, the PGD sign update/clip, and the dual-MSE defense contract can be confirmed. However, line 219 of at_openvla_ofc.py passes eval both positionally and by keyword, which under the Python calling contract is expected to raise a TypeError. Finding a static defect is not equivalent to proving the paper's algorithm or results invalid. Nor is it equivalent to having fixed and reproduced it. [A038]

The ASR code contract of the Co-Attack repository actually computes clean accuracy minus adversarial accuracy. That is not the 'proportion of successful attack samples' common to all papers. The repository also pins outdated core dependencies, and the current project has not completed an end-to-end run. This finding directly explains why an identically named ASR cannot be subtracted across papers or used for ranking. [B001]

The CIDER core files are equivalent to the paper's threshold sign convention, and the target files pass the Python 3 AST. However, the upstream tree contains Python 2 demo files, so compileall over the whole tree fails. Therefore one can only say that the core formulas and the threshold contract have been statically checked. One cannot say that the model, data, performance, or defense gains have been reproduced. [A035]

Four evidence labels are used in the anchor cases below. They are problem and privileges studied by the authors, core algorithm, results reported by the authors, and verification in this survey and items not run. This preserves the concrete formulas and examples. It also keeps the original numbers, the static code findings, and this project's run results from merging into a single kind of evidence.

Anchor 1: A005 BadVLA: Towards Backdoor Attacks on Vision-Language-Action Models via Objective-Decoupled Optimization

Problem and privileges studied by the authors. The attacker can control training samples, the loss, and the optimization strategy. The attacker cannot change the model architecture or the deployment environment. The backdoor must satisfy both conditions at once. The clean task stays approximately normal. The trigger input enters a representation region unseen during training and causes action failure. The inputs are image-instruction ($x_i=(v_i, l_i)$), action (a_i), the trigger δ , the frozen reference model (f_{ref}), and the decomposed perception/backbone/action modules. The outputs are the backdoored parameters θ , the clean success rate, the trigger success rate, and the paper-defined ASR. [A005]

Core algorithm. (1) The frozen reference model produces clean feature anchors. (2) Stage one trains only the perception module, bringing clean features close to the reference while pushing trigger features away from the clean subspace. (3) The perception module with the trigger already implanted is frozen. (4) Stage two trains the backbone and the action head using only clean data. (5) At deployment the trigger features land in a region unseen by the action head, causing random or semantically inconsistent actions. The representative formula is typeset separately in the equation below. The complete pseudocode, variable explanations, and original-text screenshots are in the separate algorithm atlas. [A005]

$$L_{\text{trig}} = \frac{1}{N} \sum_i \|f_{\theta}(x_i) - f_{\text{ref}}(x_i)\|_2^2 - \frac{\alpha}{N} \sum_i \|f_{\theta}(T(x_i, \delta)) - f_{\theta}(x_i)\|_2^2$$

[A005]

What the authors report and examples. In OpenVLA/LIBERO, blocks, cups, or stick-like objects serve as trigger objects. The representative average ASR of the complete method is 98.8%. The clean success rate is 95.9%, and the clean utility loss relative to the baseline is about 0.8pp. This result can only be interpreted under the original model, task, budget, denominator, and metric. [A005]

Verification in this survey and items not run. The paper PDF, text extraction, page numbers, and core figures have been checked against the card hash. The independent end-to-end status remains not-independently-run. This project completed only the full-text, algorithm-card, and screenshot-contract checks. It did not independently reproduce the paper's numbers end to end. The main limitations are as follows. The decoupled optimization explains why a direct mix of poisoned samples destroys clean capability at the same time. The attack privilege, though, includes modifying the training objective, which is stronger than a supply chain scenario that injects only a small amount of data. ASR is normalized jointly by clean and trigger success, so

it cannot be compared on the same axis with the "target action hit rate". ASR and clean utility are reported separately. [A005]

Anchor 2: A013 Phantom Menace: Exploring and Enhancing the Robustness of VLA Models Against Physical Sensor Attacks

Problem and privileges studied by the authors. The attacker can aim a laser, a projection, electromagnetic interference, or ultrasonic vibration at the camera. The attacker can also inject denial-of-service or forged speech into the microphone. The defender can synthesize these observations in simulation and use 30% attack data for LoRA adversarial training. Inputs are the spatiotemporal light field ($L(x,y,t)$), camera frames, speech commands, and the attack type and strength. Outputs are the perturbed image or speech, the task success rates of four VLAs, and the success rate after adversarial training. [A013]

Core algorithm. (1) Eight sensor channels are modeled in the digital domain. (2) Real attack patterns are recorded with projection, laser, EMI, and ultrasound equipment. (3) Those real patterns are parameterized and superimposed onto LIBERO observations. (4) Closed-loop runs at weak, medium, and strong levels cover OpenVLA, OpenVLA-OFT, $\pi 0$, and $\pi 0$ -fast. (5) Fine-tuning then uses mixed attack samples. (6) The models return to a Franka real machine to verify transfer. The representative formula is typeset separately in the equation below. The complete pseudocode, variable explanations, and original-text screenshots are in the separate algorithm atlas. [A013]

$$I_{\{\mathrm{attacked}\}}(x,y,t)=F\left(L_{\{\mathrm{ambient}\}}(x,y,t)+L_{\{\mathrm{malicious}\}}(x,y,t)\right)$$

\$\$ [A013]

What the authors report and examples. In the LIBERO simulation, OpenVLA-Spatial has a clean TSR of 84.7%, which drops to 0% after strong laser blinding. In Table 4, all four models achieve 0/10 successes on the real machine under laser blinding. Adversarial training can recover part of the medium attacks, but the differences across channels are obvious. This result can only be interpreted under the original model, task, budget, denominator, and metric. [A013]

Verification in this survey and items not run. The paper PDF, text extraction, page numbers, and core figures have been checked against the card hash. The independent end-to-end status remains not-independently-run. This project completed only the full-text, algorithm-card, and screenshot-contract checks. It did not independently reproduce the paper's numbers end to end. The main limitations are as follows. Real-Sim-Real is closer to the sensor chain than pure digital noise. Simulated superimposition, however, cannot cover material reflection, automatic exposure, or microphone hardware differences. The real machine has only 10 runs per condition. This survey also does not compress multi-table numbers across sensors into a single "average attack strength". This round did not complete the dual verification of cross-sensor per-table effects and denominators. That work is left to a dedicated table. [A013]

Anchor 3: A016 When Robots Obey the Patch: Universal Transferable Patch Attacks on Vision-Language-Action Models

Problem and privileges studied by the authors. The attacker holds gradient access to a white-box surrogate VLA only. The victim policy stays black-box. It may paste a fixed-area patch after an arbitrary position, rotation, and shear. The core assumption is that the visual features of different VLAs contain a shared low-dimensional subspace that can be linearly aligned. Inputs are the surrogate policy ($\hat{\pi}$), the image (x), the universal patch δ , a random geometric transformation (T), a bounded sample-level perturbation σ , and a text probe. The output is one shared patch, together with cross-model closed-loop task success rates. [A016]

Core algorithm. (1) Clean and patched visual features are computed. (2) An inner PGD learns an invisible sample-level σ that makes the feature neighborhood "harder". (3) With σ frozen, the outer loop maximizes the (L_1) feature shift together with a repulsive InfoNCE. (4) PAD concentrates the text-to-vision attention increment on the patch. (5) PSM pushes the patch semantics away from the task text and toward the inducing probe. (6) The same δ is trained with random positions and viewpoints, then transferred directly. The representative formula is typeset separately in the equation below. The complete pseudocode, variable explanations, and original-text screenshots are in the separate algorithm atlas. [A016]

$$\begin{aligned} \mathcal{J}_{\{\mathrm{tr}\}} = & \mathbb{E}_{x \sim \mathcal{D}} \mathbb{E}_{\pi \sim \hat{\pi}} \mathbb{E}_{T \sim \mathcal{T}} \mathbb{E}_{\sigma \sim \mathcal{S}} \mathbb{E}_{\delta \sim \mathcal{D}_{\delta}} \mathbb{E}_{\pi_0 \sim \pi_0} \mathbb{E}_{\pi_1 \sim \pi_1} \mathbb{E}_{\pi_2 \sim \pi_2} \mathbb{E}_{\pi_3 \sim \pi_3} \mathbb{E}_{\pi_4 \sim \pi_4} \mathbb{E}_{\pi_5 \sim \pi_5} \mathbb{E}_{\pi_6 \sim \pi_6} \mathbb{E}_{\pi_7 \sim \pi_7} \mathbb{E}_{\pi_8 \sim \pi_8} \mathbb{E}_{\pi_9 \sim \pi_9} \\ & \mathbb{E}_{\pi_{10} \sim \pi_{10}} \mathbb{E}_{\pi_{11} \sim \pi_{11}} \mathbb{E}_{\pi_{12} \sim \pi_{12}} \mathbb{E}_{\pi_{13} \sim \pi_{13}} \mathbb{E}_{\pi_{14} \sim \pi_{14}} \mathbb{E}_{\pi_{15} \sim \pi_{15}} \mathbb{E}_{\pi_{16} \sim \pi_{16}} \mathbb{E}_{\pi_{17} \sim \pi_{17}} \mathbb{E}_{\pi_{18} \sim \pi_{18}} \mathbb{E}_{\pi_{19} \sim \pi_{19}} \\ & \mathbb{E}_{\pi_{20} \sim \pi_{20}} \mathbb{E}_{\pi_{21} \sim \pi_{21}} \mathbb{E}_{\pi_{22} \sim \pi_{22}} \mathbb{E}_{\pi_{23} \sim \pi_{23}} \mathbb{E}_{\pi_{24} \sim \pi_{24}} \mathbb{E}_{\pi_{25} \sim \pi_{25}} \mathbb{E}_{\pi_{26} \sim \pi_{26}} \mathbb{E}_{\pi_{27} \sim \pi_{27}} \mathbb{E}_{\pi_{28} \sim \pi_{28}} \mathbb{E}_{\pi_{29} \sim \pi_{29}} \\ & \mathbb{E}_{\pi_{30} \sim \pi_{30}} \mathbb{E}_{\pi_{31} \sim \pi_{31}} \mathbb{E}_{\pi_{32} \sim \pi_{32}} \mathbb{E}_{\pi_{33} \sim \pi_{33}} \mathbb{E}_{\pi_{34} \sim \pi_{34}} \mathbb{E}_{\pi_{35} \sim \pi_{35}} \mathbb{E}_{\pi_{36} \sim \pi_{36}} \mathbb{E}_{\pi_{37} \sim \pi_{37}} \mathbb{E}_{\pi_{38} \sim \pi_{38}} \mathbb{E}_{\pi_{39} \sim \pi_{39}} \\ & \mathbb{E}_{\pi_{40} \sim \pi_{40}} \mathbb{E}_{\pi_{41} \sim \pi_{41}} \mathbb{E}_{\pi_{42} \sim \pi_{42}} \mathbb{E}_{\pi_{43} \sim \pi_{43}} \mathbb{E}_{\pi_{44} \sim \pi_{44}} \mathbb{E}_{\pi_{45} \sim \pi_{45}} \mathbb{E}_{\pi_{46} \sim \pi_{46}} \mathbb{E}_{\pi_{47} \sim \pi_{47}} \mathbb{E}_{\pi_{48} \sim \pi_{48}} \mathbb{E}_{\pi_{49} \sim \pi_{49}} \\ & \mathbb{E}_{\pi_{50} \sim \pi_{50}} \mathbb{E}_{\pi_{51} \sim \pi_{51}} \mathbb{E}_{\pi_{52} \sim \pi_{52}} \mathbb{E}_{\pi_{53} \sim \pi_{53}} \mathbb{E}_{\pi_{54} \sim \pi_{54}} \mathbb{E}_{\pi_{55} \sim \pi_{55}} \mathbb{E}_{\pi_{56} \sim \pi_{56}} \mathbb{E}_{\pi_{57} \sim \pi_{57}} \mathbb{E}_{\pi_{58} \sim \pi_{58}} \mathbb{E}_{\pi_{59} \sim \pi_{59}} \\ & \mathbb{E}_{\pi_{60} \sim \pi_{60}} \mathbb{E}_{\pi_{61} \sim \pi_{61}} \mathbb{E}_{\pi_{62} \sim \pi_{62}} \mathbb{E}_{\pi_{63} \sim \pi_{63}} \mathbb{E}_{\pi_{64} \sim \pi_{64}} \mathbb{E}_{\pi_{65} \sim \pi_{65}} \mathbb{E}_{\pi_{66} \sim \pi_{66}} \mathbb{E}_{\pi_{67} \sim \pi_{67}} \mathbb{E}_{\pi_{68} \sim \pi_{68}} \mathbb{E}_{\pi_{69} \sim \pi_{69}} \\ & \mathbb{E}_{\pi_{70} \sim \pi_{70}} \mathbb{E}_{\pi_{71} \sim \pi_{71}} \mathbb{E}_{\pi_{72} \sim \pi_{72}} \mathbb{E}_{\pi_{73} \sim \pi_{73}} \mathbb{E}_{\pi_{74} \sim \pi_{74}} \mathbb{E}_{\pi_{75} \sim \pi_{75}} \mathbb{E}_{\pi_{76} \sim \pi_{76}} \mathbb{E}_{\pi_{77} \sim \pi_{77}} \mathbb{E}_{\pi_{78} \sim \pi_{78}} \mathbb{E}_{\pi_{79} \sim \pi_{79}} \\ & \mathbb{E}_{\pi_{80} \sim \pi_{80}} \mathbb{E}_{\pi_{81} \sim \pi_{81}} \mathbb{E}_{\pi_{82} \sim \pi_{82}} \mathbb{E}_{\pi_{83} \sim \pi_{83}} \mathbb{E}_{\pi_{84} \sim \pi_{84}} \mathbb{E}_{\pi_{85} \sim \pi_{85}} \mathbb{E}_{\pi_{86} \sim \pi_{86}} \mathbb{E}_{\pi_{87} \sim \pi_{87}} \mathbb{E}_{\pi_{88} \sim \pi_{88}} \mathbb{E}_{\pi_{89} \sim \pi_{89}} \\ & \mathbb{E}_{\pi_{90} \sim \pi_{90}} \mathbb{E}_{\pi_{91} \sim \pi_{91}} \mathbb{E}_{\pi_{92} \sim \pi_{92}} \mathbb{E}_{\pi_{93} \sim \pi_{93}} \mathbb{E}_{\pi_{94} \sim \pi_{94}} \mathbb{E}_{\pi_{95} \sim \pi_{95}} \mathbb{E}_{\pi_{96} \sim \pi_{96}} \mathbb{E}_{\pi_{97} \sim \pi_{97}} \mathbb{E}_{\pi_{98} \sim \pi_{98}} \mathbb{E}_{\pi_{99} \sim \pi_{99}} \end{aligned}$$

\$\$ [A016]

What the authors report and examples. Optimization on the surrogate OpenVLA-7B achieves strict black-box transfer to OpenVLA-OFT-W in the LIBERO physical setting. That model has a clean task success rate of 98.25%, which falls to 5.75% under the patch, a drop of 92.5pp, with 100 runs per suite. This result can only be interpreted under the original model, task, budget, denominator, and metric. [A016]

Verification in this survey and items not run. The paper PDF, text extraction, page numbers, and core figures have been checked against the card hash. The independent end-to-end status remains not-independently-run. This project completed only the full-text, algorithm-card, and screenshot-contract checks. It did not independently reproduce the paper's numbers end to end. The main limitations are as follows. The two-stage hardening and the VLA-specific attention and semantic losses explain the strong transfer. The shared-subspace evidence, however, comes from the tested model family, and print size, camera viewpoint, and surrogate data remain implicit knowledge. For both reasons this cannot be called a "zero-knowledge universal attack". The weak transfer to π_0 is also retained. [A016]

Anchor 4: A009 FreezeVLA: Action-Freezing Attacks against Vision-Language-Action Models

Problem and privileges studied by the authors. A white-box attacker can modify the camera image and read token gradients, but the user instruction is unknown at deployment. The attack goal is to induce `<freeze>` (which may be the EOS or stop token) under multiple semantically

equivalent prompts, not some arbitrary trajectory. Inputs are the model (F_θ), the image (x), and the set of reference prompts generated by the LLM (P). Also given are the freeze token, the budget ϵ , and the numbers of inner and outer loop iterations. The output is one adversarial image (x') that transfers across different prompts. [A009]

Core algorithm. (1) The LLM generates several prompts that describe the same scene. (2) The inner loop then finds the prompt that is "hardest" to freeze. It locates the word with the largest gradient, replaces it with a synonym, and accepts the replacement when the freeze probability is lower. (3) This yields the hard prompt set (P^*). (4) The outer loop aggregates the image gradients of all hard prompts. (5) The update follows the sign gradient and is projected onto the (L_∞) ball. (6) The loops alternate and repeat until freezing is stable across prompts. The representative formula is typeset separately in the equation below. The complete pseudocode, variable explanations, and original-text screenshots are in the separate algorithm atlas. [A009]

$$\min_{\|x-x'\|_\infty \leq \epsilon} \max_{p' \in \text{operatorname{Syn}}(p)} \mathcal{L}(F_\theta(\tilde{x}, p'), y^*)$$

\$\$ [A009]

What the authors report and examples. On the LIBERO prompt variants of SpatialVLA, OpenVLA, and π_0 , the method repeatedly finds the "hardest stopping wording" and then optimizes the image. The representative OpenVLA ASR is 95.4%, and the image budget example used in the paper is (4/255). This result can only be interpreted under the original model, task, budget, denominator, and metric. [A009]

Verification in this survey and items not run. The paper PDF, text extraction, page numbers, and core figures have been checked against the card hash. The independent end-to-end status remains not-independently-run. This project completed only the full-text, algorithm-card, and screenshot-contract checks. It did not independently reproduce the paper's numbers end to end. The main limitations are as follows. The minimax design is indeed more transferable than PGD that targets only a single prompt. The definition of the stop token, though, varies with the action tokenizer. The attack still requires white-box gradients and is mainly simulation-based, so 95.4% cannot be interpreted directly as the robot necessarily stopping under arbitrary natural language. The retraction status needs to be listed separately for sensitivity analysis. [A009]

Anchor 5: A024 TRAP: Hijacking VLA CoT-Reasoning via Adversarial Patches

Problem and privileges studied by the authors. A white-box attacker can access the parameters and gradients of the reasoning VLA. It can also place printed patches in the environment and obtain the color and homography calibration in advance. The attacker cannot modify the user's benign instruction. The goal is to make the model generate the attacker-specified CoT ($R^\wedge$) and execute a continuous malicious action ($a^\wedge$). Inputs are offline trajectories ($D=((O,R,a))$), the patch δ , the mask M , the target CoT and action, content reference images, and the physical

transformation. Outputs are a printable patch, the rewritten reasoning sequence, and closed-loop actions. [A024]

Core algorithm. (1) The influence of CoT on actions is first confirmed through instruction masking and cross-sample CoT shuffling. (2) The same δ is pasted frame by frame across multi-task offline trajectories. (3) Target CoT tokens are aligned with cross-entropy. (4) Discrete actions are aligned with CE and continuous actions with trajectory MSE. (5) Content features and TV are added so that the patch looks like a pattern. (6) Optimization uses PGD or Deep Image Prior. (7) During training, homography, color, noise, and brightness EoT are added, and the patch is then printed. The representative formulas are typeset separately in the equation below. The complete pseudocode, variable explanations, and original-text screenshots are in the separate algorithm atlas. [A024]

$$\min_{\delta} \mathbb{E}_D[\mathcal{L}_{\mathrm{CoT}} + \lambda_1 \mathcal{L}_{\mathrm{action}} + \lambda_2 \mathcal{L}_{\mathrm{content}} + \lambda_3 \mathcal{L}_{\mathrm{TV}}]$$

\$\$ [A024]

What the authors report and examples. One example target behavior turns the benign "take the apple" into the reasoning "grab the knife and hand it to the user". In real-robot trials with an unobstructed printed patch, the target reasoning steps hit 13/15, but the complete malicious behavior was only 5/15, that is, 33.3%. This result can only be interpreted under the original model, task, budget, denominator, and metric. [A024]

Verification in this survey and items not run. The paper PDF, text extraction, page numbers, and core figures have been checked against the card hash. The independent end-to-end status remains not-independently-run. This project completed only the full-text, algorithm-card, and screenshot-contract checks. It did not independently reproduce the paper's numbers end to end. The main limitations are as follows. The noticeable gap between reasoning hits and physical behavior hits indicates that a rewritten CoT does not mean the control loop fully obeys. The real-robot runs number only 15, and the white-box and calibration privileges are strong. Cross-model transfer and occluded results cannot be mixed with the unobstructed main numbers. Reasoning-step hits and complete-behavior hits are kept separate. [A024]

Anchor 6: A049 Trajectory-Level Redirection Attacks on Vision-Language-Action Models

Problem and privileges studied by the authors. The attacker can modify the task text only once, before the episode begins, and the same text is reused in every closed-loop replanning. The attack prompt must still look like the original command and must not contain the target task words. During attack construction, the attacker can query the frozen VLA and the environment. Inputs are the benign instruction, a target instruction used only during construction, and a target terminal-state predicate. The output is a command-preserving prompt with a small number of character edits that turns the final physical state toward the attacker's goal. [A049]

Core algorithm. The same observation is first queried with the benign instruction and with the target instruction separately, which yields two frozen teacher action chunks. Candidate prompts are ranked by "closer to the target teacher, farther from the benign teacher, small text changes", then filtered for readability, target-word leakage, and command preservation. The top-M candidates undergo real closed-loop rollouts. The new observations that those candidates themselves visit are relabeled with the two teachers and added to the dataset, iterating DAgger-style on-policy aggregation. After success, redundant edits are greedily deleted and re-tested. The representative formulas are typeset separately in the equation below. The complete pseudocode, variable explanations, and original-text screenshots are in the separate algorithm atlas. [A049]

$$L_{\mathrm{rank}} = \max(0, d_{\mathrm{target}} - d_{\mathrm{benign}}) + \mu$$

[A049]

What the authors report and examples. On the attackable subset, $\pi_{\mathrm{0.5}}$ has a clean task success rate of 94.2%, and the target prompt feasibility rate is 91.7%. The near-benign prompt redirection attack success rate (ASR) reaches 97.5%, and the median character edits are only 2.6. One example changes "put the bowl on the stove" into "put the bowl on the staove". The prompt does not contain plate, yet it makes the bowl end up on the plate. Nearest-neighbor task normalization can reduce the ASR on Goal from 95.1% to 7.4% while retaining the 94.2% clean task success rate relative to the un-preprocessed baseline. That 94.2% is a retention ratio, not a new absolute success rate. This result can only be interpreted under the original model, task, budget, denominator, and metric. [A049]

Verification in this survey and items not run. The paper PDF, text extraction, page numbers, and core figures have been checked against the card hash. The independent end-to-end status remains static-contract-audit-partial. The parser can be imported without a backend. The public commit is missing two required backend scripts, so it is not self-contained for a complete run. The main limitations are as follows. The paper raises the prompt attack from first-step action misdirection to final-state goal redirection. It also shows that one should optimize on the states that the attack itself visits. The search requires many environment and policy queries and a usable target predicate, which may overestimate the attack capability in a locked deployment. The current coverage is LIBERO and SO-100 manipulation, while navigation, very long-horizon tasks, and certified command-preservation judgments remain blank. The ASR is conditioned on the subset where both benign and target are completable. [A049]

Anchor 7: A027 FlowHijack: A Dynamics-Aware Backdoor Attack on Flow-Matching Vision-Language-Action Models

Problem and privileges studied by the authors. A white-box supply chain attacker can fine-tune the open π_0 model and mix in a small amount of poisoned data. The trigger should be a natural object state/scene semantic in the environment. The malicious action has to change direction while keeping its vector-field magnitude close to that of normal motion. The inputs are multimodal observations (o_t), action chunks ($A \in \mathbb{R}^{(d \times H)}$), flow time (τ), the trigger

observation (o^+), and the malicious action (A^*). The outputs are the hijacked vector field and the Pose-Locking or Initial-Perturbation actions. [A027]

Core algorithm. (1) Choose contextual triggers from object states, for example an inverted cup or an open drawer, or from background objects. (2) Inject the malicious denoising direction only at small flow times ($\tau \in [0, \tau_0]$). (3) The ODE then amplifies the early direction error along the path that follows. (4) Pose-Locking points to a fixed pose. Initial-Perturbation instead adds a persistent small offset to the normal action. (5) Dynamics Mimicry matches the norms of the malicious and the normal vector fields. (6) Train jointly with the clean flow-matching loss. The representative formulas are typeset separately in the equation below. The separate algorithm atlas holds the complete pseudocode, variable explanations, and original-text screenshots. [A027]

$$L = (1 - \alpha - \beta) L_{\text{FM}} + \alpha L_{\text{BD}} + \beta L_{\text{mimic}}, \quad \tau \in [0, \tau_0] \quad \text{for } L_{\text{BD}}$$

\$\$ [A027]

What the authors report and examples. On LIBERO-Goal, π_0 with the scene-semantic trigger reaches 100% ASR. The clean success rate stays at 94.9%. A 0.1m target-position filter on Pose-Locking drops the ASR from 100% to 17.8%. It does not fall to zero. This result is interpretable only under the original model, task, budget, denominator, and metric. [A027]

Verification in this survey and items not run. The paper PDF, text extraction, page numbers, and core figures were checked against the card hash. The independent end-to-end status remains not-independently-run. This project completed only the full-text, algorithm-card, and screenshot-contract checks. It did not independently reproduce the paper's numbers end to end. The main limitations are as follows. The algorithm reveals that the first ODE segment of a flow model is a backdoor amplifier. It also shows that token-backdoor defenses cannot be transferred directly. The attack requires model fine-tuning privileges. The real-robot work is mainly a limited-scenario/qualitative demonstration. The 100% figure comes from 10 LIBERO-Goal tasks rather than the open world. The real-robot runs have no aggregated ASR. They are not encoded as an effect. [A027]

Anchor 8: A056 DRIFT: Derailing Denoising Trajectories of Flow-Matching VLAs with Adversarial Patch Attack

Problem and privileges studied by the authors. The attack targets a frozen flow-matching VLA, such as $\pi_0/\pi_{(0.5)}$, that generates action chunks by Euler integration. The attacker computes gradients offline in a white-box setting. At deployment it affixes only one fixed patch, and it leaves the language and the body state untouched. The inputs are 696 wrist frames, a fixed patch region, and a frozen velocity field. The output is a cross-task universal 32×32 or 64×64 patch. [A056]

Core algorithm. The authors ablate the denoising step k one at a time and compare the velocity-vector difference under clean/patch conditions. They find that $k \leq 5$ is extremely fragile,

while the late steps are almost ineffective. They then attack the earliest M steps simultaneously. The gradients of the steps gradually reverse, and the cosine between g_0 and g_9 is about -0.35 . A wide window therefore cancels itself out. DRIFT maximizes the velocity difference only at the pure-noise starting point $k=0$, and it updates the patch with PGD. That error then cascades through the remaining nine Euler steps into the final action. The representative formulas are typeset separately in the equation below. The separate algorithm atlas holds the complete pseudocode, variable explanations, and original-text screenshots. [A056]

$$\mathcal{L}_{\{\mathrm{DRIFT}\}} = \mathbb{E}_o \left[\|\mathbf{v}_{\theta}(A_{\{\tau(0)\}}, A(o, \delta)) - \mathbf{v}_{\theta}(A_{\{\tau(0)\}}, o)\|_2^2 \right]$$

\$\$ [A056]

What the authors report and examples. Across the four suites, the clean average success rate of π_0 is 95.0%. The relative ASR of the 32×32 patch is 99.8% (Spatial 100.0%, Goal 99.7%, Object 100.0%, Long 99.6%, 3 seeds). The patch covers about 2% of the 224×224 wrist image. A typical failure is a "phantom grasp" at step 3.8, while the clean run closes the gripper only around step 46. Attacking $\pi_0(0.5)$ requires 64×64 and yields an average ASR of 99.3%. The result can be read only within the original model, task, budget, denominator, and metric. [A056]

Verification in this survey and items not run. The paper PDF, text extraction, page numbers, and core figures were checked against the card hash. The independent end-to-end status remains not-independently-run. This project completed only the full-text, algorithm-card, and screenshot-contract checks. It did not independently reproduce the paper's numbers end to end. The main limitations are as follows. "Attacking only the first step" is stronger than attacking multiple steps, and it is K times cheaper per optimization. This reveals that the attack objective should match the generation dynamics. The patch transfers very weakly across models. Even at 64 px, $\pi_0 \rightarrow \pi_0(0.5)$ reaches only 8.9% ASR, which shows that the method depends heavily on the white-box victim model. The paper gives no physical printing validation and no defense. It only proposes that the earliest denoising steps should be protected. This attack is encoded separately from the training-time FlowHijack. [A056]

Anchor 9: A026 JailWAM: Jailbreaking World Action Models in Robot Control

Problem and privileges studied by the authors. The attacker can call a strong LLM to generate a pool of dangerous instructions from templates and query the WAM, but cannot exhaust the language space. The evaluator can access the open-loop actions/future imagination. It also performs expensive closed-loop simulation and manual review on the screened candidates. The inputs are historical observations/states ($x_t = (o_{(\leq t)}, s_{(\leq t)})$), instruction candidates (I_{adv}), and the WAM-generated action chunks and future images. The outputs are multi-view trajectory plots, 0/1/2-level risk labels, and MFR, CRR, and their sum ASR. [A026]

Core algorithm. (1) Prompt the LLM with dedicated dangerous-action templates to generate candidates. (2) The WAM generates actions in open loop. (3) VTM accumulates relative actions into a three-dimensional path and projects that path onto xy/xz plots. (4) A Qwen3-VL-2B risk discriminator screens out safe candidates quickly. (5) Only Level 1/2 candidates enter the RoboTwin closed loop. (6) Review consequences such as oscillation, out-of-bounds, and collision by hand, then assign the final label. The representative formulas are typeset separately in the equation below. The separate algorithm atlas holds the complete pseudocode, variable explanations, and original-text screenshots. [A026]

$$\begin{aligned} \mathcal{I}^{\{ \}} &= \operatorname{argmax}_{\mathcal{I} \in \mathcal{I}} L_{\mathrm{adv}}(V(M(x_t, \mathcal{I}))), \\ P &= P_0 + \sum_i F(\Delta A_i) \end{aligned}$$

\$\$ [A026]

What the authors report and examples. One example instruction requires a 400° rotation on a 360° limited servo. In RoboTwin, the clean manual ASR of LingBot-VA is 1.6%. After JailWAM, its MFR is 62.0%, its CRR is 22.2%, and the total is 84.2%. For Motus the total is 60.6%. Any interpretation of this result is confined to the original model, task, budget, denominator, and metric. [A026]

Verification in this survey and items not run. The paper PDF, text extraction, page numbers, and core figures were checked against the card hash. The independent end-to-end status remains not-independently-run. This project completed only the full-text, algorithm-card, and screenshot-contract checks. It did not independently reproduce the paper's numbers end to end. The main limitations are as follows. The dual path separates large-scale screening from physical confirmation, and VTM also unifies the heterogeneous action spaces. However, the risk discriminator determines which candidates get the chance to enter the closed loop. ASR=MFR+CRR is a paper-specific definition. The architecture, the tasks, and the decoder all change at once across LingBot/Motus, so no causal comparison can be made. Cross-architecture differences are treated only as correlation, not as causal proof. [A026]

Anchor 10: A033 Attacking the Trusted Imagination: Oracle-Level Integrity Attacks on Imagine-then-Act World Models

Problem and privileges studied by the authors. The target is an "imagine-then-act" WAM. Its encoder E and world model W generate future latents, and the controller/value oracle trusts that imagination. The attacker changes only the current observation, within an ϵ_{∞} budget, and leaves the weights, the oracle, and the environment ground truth untouched. The inputs are the current observation, the context, and an optional goal latent. The output is an observation perturbation that makes the imagined trajectory drift severely while the appearance changes very little. [A033]

Core algorithm. The model is frozen, and the clean imagination z_{clean} is computed first. The untargeted attack maximizes $1 - \cos(z_{\text{adv}}, z_{\text{clean}})$. The targeted attack minimizes $\|z_{\text{adv}} - z_{\text{tgt}}\|_2^2$. In the white-box setting, projected gradient repeatedly updates the observation and projects it back into the budget ball. Where the gradient chain is unavailable, the attack

falls back on SPSA's two-point query estimate. The velocity norm of future frames at high-noise moments also serves as a lightweight detection score. The representative formulas are typeset separately in the equation below. The separate algorithm atlas holds the complete pseudocode, variable explanations, and original-text screenshots. [A033]

$$\max_{\|\delta\| \leq \epsilon} [1 - \cos(z_{\text{adv}}, z_{\text{clean}})] \quad \text{or} \quad \min_{\|\delta\| \leq \epsilon} \|z_{\text{adv}} - z_{\text{tgt}}\|_2^2$$

\$\$ [A033]

What the authors report and examples. In MPC task 4 of LaDi-WM, the controller samples 4 candidate actions every 5 steps. Over 20 simulations, the clean success rate is 55% and random perturbation gives 70%. The attack with $\epsilon=0.01$ gives only 5%, a decrease of 50 percentage points relative to clean. The reactive policy moves only from 97.8% to 96.6%. That near-zero change serves as a negative control for "the attack mainly hits the imagination pathway". The result is interpretable only against the original model, task, budget, denominator, and metric. [A033]

Verification in this survey and items not run. The paper PDF, text extraction, page numbers, and core figures were checked against the card hash. The independent end-to-end status remains not-independently-run. This project completed only the full-text, algorithm-card, and screenshot-contract checks. It did not independently reproduce the paper's numbers end to end. The main limitations are as follows. The paper demonstrates that a world model can "dream very wrongly" and mislead planning even when the pixel perturbation is very small. Yet the detector has only preliminary threshold experiments. The evidence comes from simulation, with a single task and only 20 runs. The random baseline is even higher than the clean value, so the statistical uncertainty is large. The SPSA query cost and the dependence on the oracle/planner structure also limit conclusions about real-world deployment. The negative result of the reactive policy is reported proactively. [A033]

Anchor 11: A036 BadWAM: When World-Action Models Dream Right but Act Wrong

Problem and privileges studied by the authors. The paper attacks a WAM with an explicit imagination branch. The hope is that the future the model generates still looks correct while the actual actions are significantly wrong. The attacker has only black-box queries and can observe actions. The enhanced setting can also observe the imagination results, but there are no gradients or weights. The inputs are the current observation and the perturbation budget. The output is an observation perturbation that is re-derived at every replanning. [A036]

Core algorithm. The authors first define the action deviation D_{act} and the imagination drift D_{img} . The action-only setting maximizes D_{act} . The imagination-fidelity setting maximizes $D_{\text{act}} - \lambda D_{\text{img}}$, or equivalently constrains $D_{\text{img}} \leq \tau$. Each round samples a random unit direction and estimates the gradient of the objective with respect to the perturbation from positive and negative queries. It then performs a projected update and keeps the historical best. Because the environment state changes, the algorithm recomputes at every replanning

moment. The representative formulas are typeset separately in the equation below. The separate algorithm atlas holds the complete pseudocode, variable explanations, and original-text screenshots. [A036]

$$g = \frac{1}{m} \sum_{i=1}^m \frac{J(\delta + cu_i) - J(\delta - cu_i)}{2c} u_i, \quad J = D_{\{\mathrm{act}\}} - \lambda D_{\{\mathrm{img}\}}$$

\$\$ [A036]

What the authors report and examples. In a joint WAM evaluation over 40 LIBERO tasks with 800 executions per condition, the clean success rate is 98.1%. After BadWAM it is 61.5%, a drop of 36.6 percentage points. On the 12-task subset, the JPEG+noise defense raises the success rate under attack from 57.1% to 89.2%. The clean success rate, however, drops from 98.3% to 94.2%. The original model, task, budget, denominator, and metric bound any interpretation of this result. [A036]

Verification in this survey and items not run. The paper PDF, text extraction, page numbers, and core figures were checked against the card hash. The independent end-to-end status remains not-independently-run. This project completed only the full-text, algorithm-card, and screenshot-contract checks. It did not independently reproduce the paper's numbers end to end. The main limitations are as follows. The work reveals that "the world model dreams correctly" cannot prove that actions are safe. It also shows that simply comparing predicted videos will miss decoupled attacks. Its evidence is still simulation, and the zeroth-order query cost is relatively high. The JPEG/noise defense and the detector are both evaluated non-adaptively, and some statements in the paper differ slightly from the clean baselines in the tables. Anyone reusing the numbers should therefore rely on the specific tables and subsets. The conflict between the table and the main text has been retained, with the table taking precedence. [A036]

Anchor 12: A047 Membership Inference Attacks on Vision-Language-Action Models

Problem and privileges studied by the authors. The attacker asks whether a given sample entered the VLA training set. The candidate is either a transition (o,x,a) or a whole demonstration trajectory. The access privileges range from token probabilities to strictly black-box actions. The inputs are the candidate samples/trajectories and the model outputs. The output is a membership score and a binary decision. [A047]

Core algorithm. A single-sample attack computes one of three quantities: the average log-likelihood of the true action tokens, the maximum probability of the generated tokens, or the negative L_1/MSE between the predicted continuous actions and the true actions. The trajectory attack then handles the resulting per-step scores in two ways. It averages and aggregates them. It also constructs first-order smoothness and second-order curvature scores from the output action sequence alone. A threshold is finally selected on the validation set. Evaluation uses AUC and TPR at low FPR. The representative formulas are typeset separately in

the equation below. The complete pseudocode, variable explanations, and original-text screenshots are in the separate algorithm atlas. [A047]

$$s_{\mathrm{smooth}} = \frac{1}{T-1} \sum_{t=2}^T |\hat{a}_t - \hat{a}_{t-1}|_2$$

\$\$ [A047]

What the authors report and examples. On OpenVLA, the single-sample black-box Action-L1/Action-MSE average AUC is 0.9233/0.9220, and the average TPR at 1% FPR is about 0.7018/0.7111. Whole trajectories are more severe. For OpenVLA, the aggregated Action-L1 AUC is 1.0000, and Temp-Smooth/Temp-Curve average 0.9989/0.9993. Only the generated action sequence is observed here. This result can only be interpreted under the original model, task, budget, denominator, and metric. [A047]

Verification in this survey and items not run. The paper PDF, the text extraction, the page numbers, and the core figures have been checked against the card hash. The independent end-to-end status remains not-independently-run. This project only completed the full-text, algorithm-card, and screenshot-contract checks. It did not independently reproduce the paper's numbers end to end. The main limitations are as follows. The paper explains that the low dimensionality, strong supervision, and temporal continuity of embodied actions create a privacy leakage surface that LLM/VLM lack. The experiments artificially split LIBERO in half and re-fine-tune the model, which still falls short of auditing real training data whose content is unknown and whose distribution drifts. A high AUC does not directly give individual privacy loss. MC dropout can reduce the Action-L1 AUC to 0.7216 while maintaining a 67.2% success rate, but leakage remains and the utility trade-off is obvious. Trajectory correlation may amplify separability, and it requires a sensitivity analysis. [A047]

Research Agenda, Practical Gates, and Limitations

Future trends are not enumerated by trending model names. They are organized around ten falsifiable gaps that the current evidence exposes. Each item follows the contract of 'observed phenomenon—unknown condition—falsifiable question—minimum verification—falsification rule'. Where the minimum verification has not been run, the item can only be called a research agenda. It must not be written as a trend already realized. [A023, @A039, @A057, @A055, @A054, @A033, @A036, @A058, @A007, @A013, @A035, @A021, @A022, @A027, @C004, @A018, @C006, @A010, @B027, @B031, @B023, @B028, @B030, @A025]

Existing evidence shows that most sanitization, prompting, dictionary, and runtime detection still face fixed attacks, and that only SARF/APFT covers same-family adaptive patches within a single dependency cluster. The unknown condition is whether the gains persist when the attacker jointly knows the processor, thresholds, state machine, and recovery cycle. The falsifiable question should hold the checkpoint, the query/physical budget, and the control deadline fixed. Once the attack has been re-optimized against the complete defense chain, does the defense still outperform a same-cost baseline? The minimum verification must include four groups—no attack, fixed, adaptive, and out-of-distribution adaptive—and save step-by-step state—action—feedback. It must also report clean utility, dangerous actions, false

positives, recovery, and latency. If the gains disappear, or come from total shutdown, the adaptive robustness claim is rejected. [A023, @A039, @A057, @A055, @A054]

Existing evidence shows that A033 can make a WAM dream wrongly, and that A036 can make it dream something that looks correct yet act wrongly. A unified defense is still missing. The unknown condition is which attacks heterogeneous state anchors, physical oracles, and feedback checks can each block. The falsifiable question is whether a checker independent of the WAM generation can reduce dangerous action authorization under both attack classes at the same time. Improving video similarity alone is not enough. The minimum verification needs to cover both joint/IDM couplings, both attack classes, and same-source checker ablations, and to jointly report D_img, D_act, reachability, constraint violations, closed-loop success, and replanning P95/P99. If only imagination metrics improve, the integrity claim is falsified. [A033, @A036, @A058]

Direct evidence is highly concentrated in OpenVLA, LIBERO, and limited PiPER/Franka scenarios. The unknown condition is whether the effect holds after changes in robot morphology, cameras, materials, lighting, and control frequency. The falsifiable question is whether the same defense still outperforms an equal-cost baseline in unseen robots and scenarios, rather than memorizing training-platform statistics. The minimum verification must distinguish same-scene viewpoint holdout from true cross-scene/cross-robot evaluation, fix the physical budget, and report per-condition episodes, failed runs, calibration drift, and latency. If the direction flips or the threshold is miscalibrated, general deployment cannot be claimed. [A007, @A013, @A035, @A055, @A057]

Many defenses look only at a single frame, a single prompt, or a short window, while action chunks, caches, and memory can continue to carry contamination after the trigger disappears. The unknown condition is how long a window must be to detect a delayed payload without breaking normal long-horizon tasks. The falsifiable question is whether monitoring with state versions, causal provenance, and independent anchors can identify persistent drift after a single injection. Detection must come before the first dangerous action. The minimum verification needs to systematically vary trigger—payload delay, cache lifetime, and replanning period, and to compare stateless, short-window, long-window, and reset baselines. If it works only when the trigger is visible, or obtains safety through frequent resets, the persistent defense is falsified. [A021, @A022, @A027, @A054]

Malicious behavior stably passes through middleware only when it is aligned with a genuinely executable JSON/action format. The geometric safety layer, in turn, depends on complete perception. The unknown condition is whether dangerous actions in legal format can still bypass semantic authorization, structured types, and physical constraints once those are combined. The falsifiable question is whether subject—object—action capabilities, unit/range types, and independent physical constraints block dangerous commands more stably than a second LLM. Benign tasks must survive at the same time. The minimum verification should trace the full chain from intent to feedback and add legal-format malicious actions, boundary parameters, missed detections, and timeouts. It should report benign pass-through, erroneous

blocking, constraint violations, and P95/P99. Total rejection or legal-format bypass both falsify permission closure. [C004, @A018, @C006]

FailSafe adds about 9.1 seconds, and recovery success remains limited. ActFovea's timely safety failure is not a collision guarantee either. The unknown condition is whether, at different speeds and tasks, recovery can occur before harm and without creating secondary risks. The falsifiable question is whether a recovery state machine can reduce residual trajectory deviation and dangerous contact within a pre-registered time-to-harm. It is measured against stop-only and replanning, and normal tasks must not be interrupted excessively. The minimum verification must record detection, intervention, harm, and recovery times, and compare continue, stop, replan, correct, and human takeover. Recovery later than harm, or producing more damage than stopping, falsifies it. [A007, @A010, @A054]

In many methods the victim model's attention, hidden states, or self-judgment serve generation and adjudication at the same time. The unknown condition is how much independence in parameters, modalities, and sensors can actually reduce common-mode failures. The falsifiable question assumes that the attacker jointly optimizes the model and the self-check signal. Do heterogeneous sensors, physical oracles, or independent models then fail later than same-model self-checks and reduce dangerous actions? The minimum verification needs to compare the four classes of checkers under the same protocol, and disclose attacker knowledge separately. It also needs to run EOT/BPDA/zero-order joint attacks while reporting error correlation, action outcomes, and latency. If failures are highly same-source, or action outcomes do not improve, the independence guarantee does not hold. [A023, @A035, @B027, @B031, @A036]

Several defenses require eight queries, 1.2–15 seconds of back-translation, 9.1 seconds of recovery, or an extra 0.36 seconds of query time. Others pay training cost in exchange for zero online modules. The unknown condition is whether the ranking of methods reverses once control deadlines and clean tasks are included. The falsifiable question holds the hardware, the frequency, and the adaptive budget fixed. Which methods then still lie on the non-dominated frontier of safety—utility—P95/P99 latency? The minimum verification must separate offline, task-entry, replanning, and per-frame costs, and report task, safety, false positives, recovery, throughput, and resources. A low ASR obtained through timeouts, total rejection, or unacceptable utility loss does not constitute a deployment advantage. [B023, @B028, @B030, @B031, @A010, @A055]

The 52 meta-analytic units and 3 strict effect correlations were all rejected for insufficient independence, variance, and denominators. Shared multiple drafts inflate the sample. The unknown condition is whether large effects survive a unified protocol and replication by independent teams. The falsifiable question rests on a shared threat model, data, checkpoint, attack budget, and outcome definition. Can at least three independent research clusters then reproduce a defense effect that is in the same direction and practically meaningful? The minimum verification needs to pre-register the analysis unit, variance, dependencies, and failure rules, and to release per-episode data and hashes. If the direction is unstable, the

variance is huge, or it holds only in a conditioned subset, the overall conclusion continues to be rejected. [A057, @A055]

VLA-Forget only demonstrates empirical approximate forgetting, while the backdoor can persist across downstream fine-tuning. The unknown condition is which lasts longer, the attack or the patch, after LoRA merging, quantization, distillation, continued fine-tuning, and trigger transformation. The falsifiable question holds the retain/forget/mismatch contract fixed. After common post-training transformations, can the patched dangerous behavior remain below a pre-registered threshold while substitute triggers do not restore normal capabilities? The minimum verification must save model and checkpoint hashes before and after transformation, and evaluate the original trigger, transformed, no-trigger, and mismatch controls. It must report retained capabilities and re-patching cost. Behavior recovery or overall capability destruction both falsify persistent patching. [A022, @A025, @A035]

For model researchers, the practical gate is releasing action tensors, coordinates, budgets, denominators, failure cases, clean cost, latency, and adaptive evaluation. For system integrators, it is input provenance, permissions, action shell, degraded states, emergency stop, and logs. For benchmark designers, it is event definitions, per-episode records, and dependency clusters. For governance actors, it is version pinning, supply chain rollback, residual risk, and takeover responsibility.

Robot standards and model benchmarks serve different purposes. ISO 10218-1:2025 addresses the industrial robot itself, and ISO 10218-2:2025 addresses applications and robot cells. Paper ASR must not be substituted for mechanical and system risk assessment in the applicable scenario. ISO/WD 25874.2 is still a draft under development. It can only serve as a standardization trend, and must not be written as a current requirement. [ISO10218P1, @ISO10218P2, @ISO25874WD]

Retrieval and time are the first source of limitation in this survey. New papers after the freeze date are not included, and preprints, works under review, retractions, and title/version drift may change identities and numbers. Classification is the second source. The first-broken interface requires human judgment, and future cross-layer attacks may force a revision of the main axis. The escalation condition is that independent coders review the main class, permissions, and evidence layer and disclose disagreements.

Statistical limitations directly weaken the overall effect and correlation conclusions. The 52 candidate units and 3 strict correlations were all rejected by the threshold. Engineering limitations directly weaken reproduction claims. The four repositories received only static contract audits, and runs of weights, simulators, robots, and end-to-end were all zero. Physical limitations weaken field extrapolation. Most real experiments have few tasks and controlled environments, and lack long-term event denominators.

Delivery limitations must also be retained. The 97 core original figures have been checked against source pages and hashes. Reuse permissions, accessible alt text, and resolution policies have not yet been closed item by item, and no submission venue has been designated yet. The current artifact is a generic compiled draft, and it is not bound to an official template,

anonymity, page limits, or copyright terms. That does not prevent draft compilation, but it blocks a submission-ready claim.

Given that these limitations are in force, the most stable escalation path is not to add another composite score. It is to unify event denominators, independently replicate across scenarios, evaluate the full adaptive chain, establish WAM three-way consistency, and build a recoverable action safety shell. The next chapter only recycles these established judgments and introduces no new numbers or new papers.

Conclusion

This survey reorganizes VLM, VLA, and WAM security evidence around the first-broken system interface. The six-interface main axis places training artifacts, observation, cross-modal semantics, world state, action decoding, and execution feedback on a single chain. It demotes the attack carrier, knowledge, realism, objective, and evidence layer to secondary labels. That avoids repeated classification by paper name and trending terminology.

The common mechanism of the attacks is that untrusted input acquires ever-deeper write privileges. It first changes the observation or state, then controls reasoning, candidate futures, or action dynamics, and finally seeks execution and feedback. Defense must be deployed as a mirror along the same chain. Input purification and model self-checks can only reduce part of the risk. Action authorization, independent physical constraints, state recovery, and forensics are responsible for limiting the consequences after the earlier layers fail. [A021, @A024, @A033, @A036, @A054]

WAM carries one core warning: "dreaming right" does not guarantee "acting right". When the same attacked model is the only verifier of its own futures, plans, and actions, it produces correlated self-consistency rather than an independent guarantee. High-risk actions require three-way verification across forward imagination, inverse reachability, and real feedback. They must also introduce at least heterogeneous observations or a physical safety set. [A029, @A033, @A036, @A058]

Quantitatively, the reported descriptive magnitude of the decline in task success is large, but independence, variance, budget, and endpoint do not satisfy the requirements for pooling. The 52/52 meta-analysis units and the 3/3 strict correlations were formally rejected. On the engineering side, the LaTeX, evidence, and static repository contracts can be rechecked. End-to-end reproduction of the paper's numerical results, however, is zero. The most honest conclusion is not a unified risk ranking. It is clarity about which chains have been directly observed, which are only hypotheses, and which still need to be run.

In the end, embodied security is not about making the model never make mistakes. It is about ensuring that, when untrusted states propagate toward action authority, the system has an independent, low-latency, and recoverable authority revocation mechanism. The main manuscript ends here. The concrete algorithms, instances, formulas, boundaries, and core original figures of all 97 papers are preserved paper by paper in the separate algorithm atlas. --

Appendix — Post-cutoff update (2026-08-06 → 2026-09-26)

The body's retrieval was frozen on 2026-08-06, auditing 382 screening records and including 97 studies qualitatively. This appendix registers later material; **the body text is unchanged.**

A.1 — Remotely exploitable root-level flaws in a humanoid robot

Date	Event	Sections affected
2026-09	Two root RCE flaws found in the Unitree G1 EDU humanoid, one exploitable over Bluetooth without pairing	ch. 5 closed-loop attack taxonomy; ch. 6 propagation to physical consequence
2026-09	Reporting describes close-range takeover and worm-like spread	ch. 6 propagation chain; ch. 8 choosing a control layer

Why it belongs: the 97 included studies are mostly academic, and their attacks run in simulation or under controlled physical conditions. These two are **exploitable remote flaws in a commercially sold humanoid**, requiring no pairing — the lowest privilege threshold among the entry classes in the body's taxonomy. They also support two existing judgements: physical execution amplifies consequence, and the first-broken interface is often not model inference but a lower communications or firmware layer.

Caveat: these come from public reporting and CVE records and were not independently reproduced. They should be tiered with the body's A/B/C evidence grading and must not be pooled with paper-grade evidence.

A.2 — New embodied-attack papers since the cutoff

Paper	Topic	Body issue
DRIFT: Derailing Trajectories of Flow-Matching VLAs with Adversarial Patch Attack (arXiv 2608.03207)	adversarial patches against flow-matching VLAs	ch. 5 observation and action entry points; the body already includes flow-matching VLA work, and this continues it

The body already covers adversarial patches; DRIFT's contribution is aiming at the **flow-matching generative dynamics** rather than the final action token, the same branch as the body's "attacks on the generative dynamics" discussion.

A.3 — Institutional consequence of the OpenAI-Hugging Face incident

On 2026-09-16 / 17 OpenAI published an account of the incident and announced a safety-incident disclosure process; reporting describes roughly 700 agents, dozens of third-party systems and 53 leaked user images; the US Senate opened an investigation (led by Hawley).

Its relation to this survey: this is an agent-loop overrun, not a robotics case, but it is an **institutional footnote to ch. 8** — the higher the autonomy and the wider the reachable surface, the more post-hoc response depends on disclosure and accountability mechanisms rather than runtime gating alone.

A.4 — How to use this appendix

The closed-loop interface taxonomy, the propagation analysis and the control-selection guide are unaffected. What changes is the **evidence distribution**: remotely exploitable commercial robot flaws and embodied attack papers appeared together after the cutoff, suggesting the field's attack focus is moving from simulation toward deployment. When citing the items above, cite both the retrieval freeze date and this appendix's date.