

LLM and Multimodal Agent Security: An Evidence-Grounded Taxonomy of Attacks and Defenses

Contents

1. [Abstract](#)
2. [Introduction](#)
3. [Corpus Selection and Coding Method](#)
4. [Background and Input/Output Contract](#)
5. [Taxonomy Design and Coverage Audit](#)
6. [Defense Method Families](#)
7. [Cross-Family Synthesis and Selection Guide](#)
8. [Data, Metrics, and Evaluation Evidence](#)
9. [Incidents, Reproduction, and Deployment Mapping](#)
0. [Challenges, Future Trends, and Evidence Limits](#)
1. [Conclusion](#)
2. [Open Materials and Reproduction Statement](#)
3. [A.1 — The OpenAI-Hugging Face incident: the "not yet published" judgement is now overturned](#)
4. [A.2 — New papers and benchmarks since the cutoff](#)
5. [A.3 — New vulnerabilities since the cutoff](#)
6. [A.4 — How to use this appendix](#)

Abstract

Large language model security has grown from single-turn content violations into a system problem. The model, retrieval, memory, tools, identity, network and execution environment jointly constitute it. This survey's sole primary classification axis is the location of the end-to-end trust boundary crossed by attack impact. Along that axis the survey synthesizes text jailbreaking, indirect prompt injection, RAG and long-term memory poisoning, VLM multimodal injection, Agent/Harness hijacking, and sandbox and supply chain risks. The evidence layer holds 65 attack paper records, 62 defense and engineering sources, 32 event records and 34 study-level quantitative effects. A statistical audit finds usable event counts and denominators for only five studies, at five different endpoints, so formal meta-analytic pooling groups number zero. The survey also reports 24 text Harness invocations, 18 local sandbox capability probes, and one nine-cell VLM failure run with no valid model answers. Model alignment can reduce the probability of dangerous intent. The synthesis shows, however, that high-privilege systems must place provenance, least capability, action-level authorization, memory identity, network

egress, short-lived secrets, disposable execution and incident response outside the model. No overall breach rate is estimated here. Nor are synthetic reproductions or run timeouts extrapolated into production safety.

[← Back to contents](#)

Introduction

A text-only model has one primary endpoint: content violation. Once it can also read web pages, call tools, save memory, run code or use a cloud identity, the same output may become a function argument, an outbound address, a payment instruction or persistent state. This survey therefore treats model security as an end-to-end system problem.

The survey addresses security researchers, agent and harness engineers, and platform operations staff. It answers five questions. How do attacks cross trust boundaries? How do text, multimodal, retrieval and memory risks relate to one another? Which segment does each defense layer block? When is public evidence comparable? How are incident and mechanism reproduction turned into deployment gates?

This survey makes four main contributions. Its taxonomy uses a sole axis, the location of trust boundaries. Its quantitative audit explicitly refuses pseudo-merging. It systematically analyzes representative methods and the 2026 OpenAI—Hugging Face incident. Its local mechanism experiments separate dangerous intent, dangerous execution, benign utility and runtime failure.


[← Back to contents](#)

Corpus Selection and Coding Method

The pre-search protocol was frozen on 2026-08-06, with a main time window from 2018-01-01 to 2026-08-06 and backward tracing to the necessary foundational work. The survey draws on structured search, citation tracking, first-hand announcements and local run artifacts. It does not yet have complete dual-reviewer full-text screening and bias assessment. The accurate label is therefore a taxonomic review with a structured search module, not a strict systematic review.

The research questions run in order. How do attacks cross trust boundaries? How do different modalities and state mechanisms relate to one another? Which segment does each defense control block? When is evidence comparable? How do real-world incidents and local mechanism reproductions constrain engineering design? The primary unit of analysis is the paper-level main effect. Within-paper experiment arms, incidents and local runs are stored separately and must not be conflated as independent studies.

Figure fig:search-flow shows the auditable process of search, prioritization and evidence assembly.



Candidate retrieval, machine prioritization, and manual evidence assembly process. OpenAlex candidates have not yet undergone complete dual-reviewer full-text screening. This survey therefore does not use the strict systematic review label.

Candidate retrieval, machine prioritization, and manual evidence assembly process. OpenAlex candidates have not yet undergone complete dual-reviewer full-text screening. This survey therefore does not use the strict systematic review label.

Pre-search Protocol

The study design document, fixed before large-scale retrieval, records the research questions, article skeleton, search queries, inclusion and exclusion criteria, quality tiers, data fields, and the meta-analysis and correlation plan. If later evidence overturns an earlier judgment, corrections are appended in STAGE_SUMMARY.md only. The process record is not silently rewritten.

Automated Candidate Retrieval and Manual Verification

For each of 12 topics, the OpenAlex pipeline retrieved up to 200 relevance-ranked results, yielding 2,400 API returns in total. Deduplication by OpenAlex ID/DOI left 1,854 candidates. Transparent keyword rules labeled them 701 high-priority, 297 medium-priority and 856 low-priority. These labels serve only for ranking and cannot be regarded as final screening decisions. The complete queries, hit counts, retrieval counts and raw responses are in data/openalex/.

The first-round manual evidence base table holds 65 attack studies. After backfilling the completed formal versions, the current tiers are 40 at level A and 25 at level C. This coding may still underestimate existing conference versions. Before freezing, it is not treated as equivalent to the final “peer-reviewed/preprint” publication status. The defense and engineering table holds 62 sources at A/B/C levels of 23/23/16, covering papers as well as protocol specifications, system documentation and security advisories. Those cannot be added

to the attack papers as a “total paper count”. The compilation also adds 33 paper-model-attack configuration experiment arms and 32 incident/vulnerability/controlled-study records, the latter at A/B/C levels of 28/3/1. Unified template parsing was completed for 23 representative attack papers.

The attack draft lives in academic attack-surface retrieval, and the incident draft lives in incident and HF case verification. The main-text statistics will follow the frozen master table. That comes only after defense evidence, final-version deduplication and citation verification.

Why a Single “Overall Compromise Rate” Cannot Be Given

The evidence base holds 33 attack experiment arms. Only 3 of those rows have had precise event counts and denominators extracted from first-hand pages and verified. This does not mean the remaining papers' main texts necessarily lack denominators. It means the current base table cannot yet support a binomial pooling on this basis. More importantly, those three rows differ in sampling unit and attack target. “31/36 applications can be prompt-injected” is application-level coverage, not per-prompt ASR. A model's violation generation rate on 100 harmful requests is likewise not equivalent to an agent's real data exfiltration rate across 100 enterprise tasks. Other papers may also score success differently. The options include keyword matching, LLM judge, human review, target prefix, tool-call success and performance degradation.

This survey therefore follows three rules. Only studies whose threat model, sampling unit and success definition can be mapped onto one another enter the same meta-analysis group. When an abstract gives only “over 90%”, 0.90 is not passed off as a precise point estimate, and the denominator is not back-derived. Studies with high heterogeneity or missing denominators are presented through evidence maps and stratified narrative, not by sacrificing interpretability for a single total.

This treatment complies with PRISMA 2020's requirements for transparent reporting and with Cochrane's cautious interpretation of heterogeneity, random effects and prediction intervals. Both sets of guidelines originate in medical reviews. This survey borrows their reporting and statistical principles, but it does not claim that LLM security experiments are equivalent to clinical trials. PRISMA 2020 [page2021prisma]; Cochrane Handbook, Chapter 10 [deeks2024cochrane].

[← Back to contents](#)

Background and Input/Output Contract

This chapter separates model outputs, system actions and actual harm. It fixes the end-to-end data flow, assets, adversary capabilities and five risk amplifiers. That contract is the shared input for the taxonomy and evaluation that follow.

A Response and an Action Are Not the Same Kind of Risk

Dialogue-era evaluations often defined "attack success" as generating some class of policy-prohibited text. In the agent era, downstream parsers may read the same output as a function name, a shell command, a payment instruction, an email recipient, a browser click or a memory write. At least four layers therefore need separating. A content violation is text, images or audio the model generates in breach of the usage policy. Control hijacking is untrusted input that changes the original task or instruction priority. Dangerous intent is the model proposing actions such as reading secrets, sending data outward, deleting files or expanding privileges. Dangerous execution is the harness, tools or host environment actually letting an action through and causing an observable side effect.

The first three layers can appear in model-only benchmarks. The fourth must incorporate the architecture outside the model. A model may be jailbroken at the output level, yet the capability gate blocks the actual harm. It may also avoid conspicuous violating text while exceeding its authority through well-formed tool parameters. One therefore comes close to engineering risk only when model intent, execution results and benign task utility are reported together.

The End-to-End System and Trust Boundaries

The reference chain this survey adopts appears below.

Text structure representation in the earlier draft

```
Data/code/weights supply chain
v
Base model and alignment layer
v
System prompt, user input, and conversation context
v
Web pages/email/PDF/images/audio/RAG retrieval results
v
Planning loop and agent harness
v
Tool descriptions, function parameters, credentials, network, file, and code execution
v
Short-term state, long-term memory, user profiles, and cross-agent messages
v
Sandbox, containers/VMs, host, cloud control plane, and downstream users
```

Any arrow may cross a trust domain. Security design does not ask "is the model clever". It answers a list of questions item by item. Who can write? Is provenance verifiable? What can the model see? What can the model suggest? What will the executor allow? How long can state be kept? How large is the blast radius after a failure?

Five Amplifiers of Risk

For engineering discussion, this survey uses five dimensions. They explain why the same injection string produces vastly different harm in different systems. Reachability describes whether retrieval, OCR, transcription, summarization or tool return values bring attack content into context. Privilege describes which of read, write, execute, outbound network, payment and identity delegation capabilities the model or tools possess. Persistence distinguishes whether the impact lasts a single turn or enters memory, caches, indexes, code repositories or release artifacts. Observability examines whether unified traces, provenance labels, tool audits, kernel/network logs and timely alerts exist. Iteration speed indicates whether an adversary or an autonomous agent can attempt in parallel at low cost, adapt on feedback and continue for hours or days.


This is not a fitted risk formula. It is a conceptual framework for checking for missing controls. Asset value, attack cost, detection probability and recovery capability affect real risk as well.

[← Back to contents](#)

Taxonomy Design and Coverage Audit

The primary taxonomy axis answers only one question: which trust boundary does the attack impact cross in the end-to-end chain? Secondary labels cover entry vector, black-box or white-box access, transient or persistent state, content or action consequence and defense location. A single study may carry multiple labels. The top-level taxonomy, however, is no longer repeatedly rearranged by modality, attack name or product.

Figure fig:unified-taxonomy shows the single primary taxonomy axis used throughout this survey, together with the cross-cutting labels.



There is one axis: which trust boundary the attack crosses first. Other attributes are secondary labels.

Unified taxonomy coordinates. The primary axis is the location of the end-to-end trust boundary that the attack impact crosses. Entry, attacker privilege, persistence, consequence and defense location serve only as secondary codes. The figure is redrawn from structured evidence of the earlier project.

Unified taxonomy coordinates. The primary axis is the location of the end-to-end trust boundary that the attack impact crosses. Entry, attacker privilege, persistence, consequence and defense location serve only as secondary codes. The figure is redrawn from structured evidence of the earlier project.

Instruction--Data and Input Modality Boundary

This boundary determines whether external content can be promoted from analyzed data to a control signal. Direct jailbreaking, indirect injection and visual payloads share the same control conflict. They differ in attacker privilege, parser and consequence, so they are coded side by side only under the same boundary.

Jailbreaking and Prompt Injection Must Be Kept Separate

News coverage often conflates the two, but their security goals differ. Jailbreaking mainly bypasses content or behavioral safety policies. The attacker usually interacts with the model directly. The typical success is the model generating content it would otherwise refuse. The corresponding defenses are mainly alignment, classifiers, decoding control and red teaming. The biggest misconception treats a single success on one prompt as a permanent conclusion about capability.

Prompt injection, by contrast, changes the application's intended task, leaks context or induces tool actions. The attacker may be the current user or merely a third-party author of a web page, email, document or tool output. The typical success is the model ignoring the user's goal, fetching data, exfiltrating it or initiating actions. The key defenses are the instruction/data boundary, provenance, permissions, action gates and isolation. Its biggest misconception is that one more sentence in the system prompt establishes a hard security boundary.

PromptInject [perez2022promptinject] provided an earlier systematic demonstration of goal hijacking and prompt leakage. Jailbroken [wei2023jailbroken], in turn, explains how safety training fails in terms of competing objectives and out-of-distribution generalization. The two may use similar textual tricks. One mainly subverts application control flow, while the other mainly bypasses model policy.

HOUYI [liu2023houyi] belongs under black-box application prompt injection, where the attacker submits input directly to the application. It is not indirect injection in the strict sense of "third-party content being passively retrieved by the system." Of 36 tested real-world applications, the authors judged 31 vulnerable and obtained confirmation from 10 vendors. The sampling unit for 31/36 is the application, and backend models and versions are not uniform. The result therefore cannot be extrapolated to the overall vulnerability rate of all LLM applications.

Direct Jailbreaking: From Manual Role-Playing to Automated Optimization

Direct jailbreaking is roughly a superposition of four mechanisms. Semantic reframing uses role-play, fiction, counterfactuals, educational framing and multilingual expression. It packages harmful objectives into contexts that safety training covers less well. Surface transformation uses encoding, character substitution, tokenization anomalies, low-resource languages or structured formats. It creates inconsistency between input filters and model safety generalization. Long context and multi-turn accumulation, in turn, exploit historical state. Many-shot Jailbreaking [anil2024manyslot] places a large number of "harmful request--harmful response" demonstrations into a single context. Multi-turn attacks such as Crescendo [russinovich2025crescendo] accumulate commitments through gradual dialogue. The two differ in budget and in defense method, and shots and turn counts must not be conflated into a single variable. Finally, automated search and optimization cut the cost of manual design. The tools include white-box gradients, genetic algorithms, fuzzing and feedback from an attack model.

GCG [zou2023gcg] searches over token gradients and coordinate steps to optimize a transferable suffix. It remains an important baseline for white-box automated jailbreaking. PAIR [chao2023pair] puts an attack model in a few-query black-box loop driven by the objective and by judge feedback. GPTFuzzer [yu2023gptfuzzer] mutates new templates from manual seeds. These methods show that one success or failure of a fixed prompt is not enough to evaluate a defense. An adaptive evaluation against that defense requires that the attacker see it and re-optimize under the same budget.

Automated ASR, however, also readily overestimates actual success. Common errors include counting an affirmative opening as a harmful completion and using one LLM of the same kind to attack and judge. Truncated text, target-model API drift and reporting only the best prompt also inflate the figure. A high-quality reproduction must fix the model version, the chat template, the random seed and the attack budget. It must also fix the dual refusal/harm criteria and add manual spot checks.

Indirect Prompt Injection: External Data Becomes Control Flow

Indirect injection does not require the attacker to send instructions directly into the user's session. The attack text can hide in search results, web pages, email, PDF text layers, code comments, tickets, calendar entries, image OCR or tool responses. The agent retrieves that text on its own while carrying out a normal task. It then puts the data and the system instructions into one natural-language context.

Greshake et al. [greshake2023indirect] demonstrated the feasibility of indirect prompt injection in Bing Chat, code completion systems and GPT-4-based synthetic applications. The impacts they record include data theft, functionality manipulation, worm-like propagation and API call control. That study provides cases and a threat taxonomy, not an application-level ASR with a uniform denominator. InjecAgent [zhan2024injecagent] further covers 17 user tools, 62 attacker tools and 30 agent configurations across 1,054 cases. It advances evaluation to tool-integrated agents. Its approximately 24% is a metric over valid benchmark cases, not a real-world production side-effect rate. Traditional security language can state the root cause both share. The system lacks a strongly typed separation between trusted instructions and untrusted data, and it allows untrusted data to influence high-privilege control flow.

A system prompt that says "ignore the instructions in the web page" is a soft constraint, not a complete fix. Attackers can rewrite the semantics, fragment the attack across multiple turns, exploit tool descriptions, or make the model read the action as part of the original task. A real defense also needs provenance labels, data structures, tool allowlists, parameter validation, capability gates and limits on the consequences of an action.

VLM and Multimodality: Seams Between Parsers

VLM security is not as simple as placing one more image encoder in front of a text model. It introduces a new parsing chain. Pixels or sound waves go through scaling, compression, OCR/ASR and visual or audio encoding first. Only then are they fused with the language context. Attackers can exploit inconsistencies in what different components see. Prominent or low-salience text in an image instructs the model to ignore the user's task. Adversarial perturbations or patches change the visual representation while humans still regard the content as normal. A PDF's visible layer, hidden text layer, OCR result and layout order contradict one another. Background speech, low volume, adversarial perturbations or transcription errors in audio carry the attack signal. Screen agents convert web page text, button labels and the user's goal into the same action context. Cameras or embodied systems may connect physical-world stickers directly to navigation, purchasing or device control.

Visual Adversarial Examples [qi2024visual] uses white-box visual optimization. On the studied models it shows that a single adversarial image can form a universal jailbreak for multiple categories of harmful text requests. It does not establish an equivalent success rate under ordinary images or physical patches. FigStep [gong2025figstep] is a black-box typographic visual prompt attack. Its formal AAIL 2025 version reports an average ASR of 82.50% on 6 open-source LLMs. Image Hijacks [bailey2024imagehijacks], in turn, applies small automated perturbations to LLaVA. It studies targeted output, context leakage, safety override and false

beliefs, and reports a success rate above 80% for all four categories. The three cannot be merged into a single VLM ASR: white-box continuous pixel optimization, black-box typographic prompts, and runtime behavior matching on a specific model.

Multimodality also includes semantic composition and audio pathways. HADES, ECCV 2024 [li2024hades] reports an average ASR of 90.26% for LLaVA-1.5 on its data of 750 harmful instructions in 5 categories. The same data give Gemini Pro Vision an average ASR of 71.60%. That figure can only be regarded as a historical snapshot of a closed-source version. Across 12 categories of harmful questions, SpeechGuard, ACL Findings 2024 [peri2024speechguard] reports an average ASR of 90% for white-box digital audio perturbations. It reports about 10% for black-box transfer. These numbers show that white-box direct-input results cannot be extrapolated to cross-model settings or to real playback/recording. PDF here is a carrier and a parsing boundary, not an independent model modality. Existing research still lacks a unified dedicated benchmark covering multiple PDF parsers, OCR and layout reconstruction.

VLM defenses should at minimum surface OCR/ASR results and their provenance. They should treat text inside images as data rather than instructions by default. Before acting, they should check the user's goal, the recognized region and the tool consequence together. In high-risk scenarios they should cross-validate with independent parsers. Image compression or random transformation may reduce some adversarial perturbations. Yet it cannot reliably handle clearly legible malicious text, and it may also break normal visual tasks.

Retrieval--Context and Tenant Boundary

This boundary covers ingestion, indexing, retrieval and context assembly. The main risk is not a single malicious sentence. It is provenance, access control, ranking and instruction semantics entering the model together.

RAG: Poisoning, Unauthorized Retrieval, and Data Extraction

RAG introduces at least four distinct classes of risk. Content poisoning lets malicious or misleading records enter the knowledge base and rank highly under specific queries. Instruction poisoning places operational instructions targeting the LLM/agent into retrieved chunks, forming indirect prompt injection. Privilege-escalating retrieval exploits errors in vector stores, metadata filtering or tenant mapping, letting a user retrieve records they are not authorized to access. Corpus extraction recovers private text, index membership or system structure through carefully crafted queries, output feedback or embedding behavior.

PoisonedRAG [zou2025poisonedrag] constructs malicious documents for a target question--target answer chosen by the attacker. In a knowledge base containing millions of texts, the formal USENIX Security 2025 paper reports that injecting 5 malicious texts per target question can achieve 90% ASR. This result belongs to a targeted configuration and does not mean that 5 documents will generally control all queries. Attack success also depends on the embedding model, chunking, retrieval depth, reranking and the generator.

RAG data extraction and poisoning are also not the same endpoint. Spill the Beans [qi2025spillbeans] recovers private corpora on production custom GPTs using self-generated

queries. It reports 41% verbatim recovery from a book of about 77,000 words and 3% from a corpus of about 1,569,000 words. These are word coverage rates, not prompt-level ASR.

Defenses therefore cannot filter only at the final prompt. Document ingestion, identity ACLs, index partitioning, provenance signatures, anomalous similarity, reranking and answer citations all need to be included. OWASP 2025 lists vector and embedding weaknesses as a separate LLM risk, and particularly emphasizes multi-tenant data leakage and poisoning. That list is an engineering risk framework, not quantitative efficacy evidence. OWASP LLM08:2025 [owasp2025vector].

Current Session—Long-Term State Boundary

This boundary extends a single input into cross-session writes, recall and action chains. It introduces identity, versioning, deletion, and backup and restore problems.

Persistent Memory: One Injection Becomes Cross-Session State

The memory system turns "a single error inside the context window" into a database and identity problem. The attack chain usually contains three stages. At the write stage, external content, an automatic summary, a tool result or a conversation is saved as long-term memory. At the recall stage, a future query retrieves the poisoned record because of the user, the task or a trigger word. At the execution stage, the model treats the recalled content as a user preference, a fact, a policy or a high-priority instruction. It then changes its actions accordingly.

This is harder to detect than ordinary prompt injection. At write time the content may look harmless, and it only becomes malicious after several records are combined. An attack may also lie dormant for weeks before it triggers. Potential consequences include long-term goal hijacking, cross-user canary leakage, incorrect identity attribution and memory DoS. They also include residue in summaries, vectors or backups after deletion.

Defenses must simultaneously control who can write, what is written, and under whose identity it is written. They must also control how long it is kept, when it is read, and to which tool it is read. Adding an ACL to the vector store alone is still insufficient. An ACL can stop user B from directly reading user A's records. It cannot prove that an automatically saved web page under user A's name is a trusted preference. A reliable architecture should retain provenance, tenant, write principal, trust level, time, purpose, sensitivity labels and version. It should then filter by policy at read time according to the task and the tool.

Memory attacks have already produced formal conference evidence. AgentPoison, NeurIPS 2024 [chen2024agentpoison] directly poisons long-term memory or RAG knowledge bases across three types of agents. It reports an average ASR of not less than 80%, an impact on clean performance of not more than 1%, and a poisoning rate below 0.1%. MINJA, NeurIPS 2025 [dong2025minja], by contrast, narrows the privilege to ordinary query interaction. It does not require direct modification of the memory store. The two differ in write entry point, retriever and trigger mechanism, and cannot be extrapolated to all memory implementations. Later papers on dormancy, composition and defense still need to be stratified by formal status.

Model—Executor, Identity, and Resource Boundary

This boundary determines whether a model plan obtains file, network, process, payment or identity side effects. Dangerous intent and dangerous execution must be measured separately.

Agent, Tools, and Harness: Harm Is Determined by Permissions

Agent security cannot audit only the model prompt. It must also audit the harness that turns model output into actions. Common attack paths include the following. A tool description or return value injects new instructions. The model picks a tool whose functionality is overly broad or whose default parameters are dangerous. Parameter strings reach a shell, SQL, a template, a URL, a path or a deserializer. An agent inherits the user's long-lived credentials with no task-level scope or expiration time. A multi-turn loop repeatedly tries failed paths and thus forms a cost or availability attack. One agent mistakes a message from another agent or an MCP server for a trusted instruction. The model proposes an action and then reviews it itself, with no independent control plane.

AgentDojo [debenedetti2024agentdojo] contains 97 benign tasks and 629 security test cases. Its value is that it measures benign task utility and post-injection goal success together, not just the answer text. A model may also fail a task even when no attack is present. That makes "refuse everything" an ineffective defense. OWASP summarizes "excessive functionality, excessive permissions, and excessive autonomy" as Excessive Agency. These three are closer to a controllable engineering root cause than "whether the model hallucinates". OWASP LLM06:2025 [owasp2025agency].

Protocols such as MCP provide a unified interface for tool interoperability. They also bring identity, token audience, delegated authorization, session hijacking and the confused deputy into model applications. MCP's official security documentation explicitly prohibits token passthrough. It requires validating the token audience, implementing per-client consent and minimizing scope. Authorization state and capability references should also be set to single use or a short expiration time. The documentation recommends that local servers run in a sandbox with least privilege by default. MCP Security Best Practices [mcp2026security]. A protocol can specify transport and authorization. It will not automatically judge whether an action the model takes on the basis of untrusted content matches the user's true intent.

Availability and Economic Attacks

The resource units of an LLM application include input tokens, output tokens, KV cache, concurrency, tool steps, retrieval counts, browser pages, code runtime and external API bills. An attacker can induce extremely long contexts, recursive agents, repeated tool failures, output inflation, expensive safety-classifier cascades or a global refusal. The result is DoS at the model level, at the harness level and at the business level.

The three representative dimensions cannot be interchanged. Safeguard is a Double-edged Sword [zhang2024safeguarddos] reports that a specific generic prompt blocked more than 97% of benign requests to Llama Guard 3. Its endpoint is guardrail false blocking. P-DoS [gao2024dosp] uses a single poisoned sample to amplify an output of about 0.5K to a cap of

about 16K. Its endpoint is output tokens/cost. AgentDoS, USENIX Security 2026 [luo2026agentdos] found resource-exhaustion vulnerabilities in 16 of 20 tested agent applications. Its endpoint is application resource management defects. The three cannot be merged into a single ASR.

Therefore every agent run should have a hard budget. Maximum tokens, maximum steps, maximum concurrency, per-tool rate, total wall clock, network and storage quotas, and explicit termination conditions all belong in it. An insufficient budget reduces benign utility. An unlimited budget turns a single injection into a long-running search. Evaluation must make this trade-off explicit.

Training—Artifact—Runtime and Supply Chain Boundary

This boundary covers parametric memory, privacy extraction, data and weight poisoning, executable model files, dependencies and the release pipeline. Format security, provenance integrity and behavioral safety cannot substitute for one another.

Privacy, Model Extraction, and System Prompt Leakage

The objects of privacy attacks can be divided into the training corpus, the current context, the system prompt/tool description, private RAG records, other users' memory, and the model itself. Membership inference determines whether a sample participated in training. Training data extraction recovers specific sequences. Model inversion attempts to reconstruct the input or its statistical features, while attribute inference infers sensitive attributes. Model stealing recovers functionality, decision boundaries or part of the parameters. Their attack goals and success metrics differ.

Extracting Training Data from Large Language Models [carlini2021extracting] recovers hundreds of verbatim training sequences from the GPT-2 family. It works by black-box generation and candidate ranking, and it demonstrates that parametric memory can be extracted. Scalable Extraction of Training Data from Aligned, Production Language Models, ICLR 2025 [nasr2025scalable] reports further that two attacks can recover thousands of training samples from proprietary aligned models, ChatGPT among them. Neither result means that any prompt can retrieve any training data. Risk depends on duplication, guessable prefixes, the sampling budget, the model interface, external validation and post-processing. System prompt leakage should not be confused with training data extraction either. Extraction is parametric memory; leakage is usually a failure of current-context isolation.

Neighbourhood Comparison [mattern2023membership] calibrates text difficulty with synthetic neighbors. It performs membership inference, not verbatim extraction. Stealing Part of a Production Language Model [carlini2024stealing] instead recovers projection layer information under black-box API access. The paper reports that the complete projection matrix of Ada/Babbage can be recovered for less than 20 US dollars, and it estimates the corresponding query cost for GPT-3.5-turbo at below 2,000 US dollars. Recovering that matrix covers only part of the model, and it is not the same as copying the complete model.

"Writing a secret into the system prompt" is not secret management in itself. As long as the model has to see the value, an erroneous output or a tool parameter may reproduce it. A real credential should come from a controlled tool at execution time. The model should hold only a short-lived capability reference that cannot be redeemed directly.

Training, Weights, and the Model Supply Chain

A supply chain attack can occur at any stage. The stages include pretraining data, instruction tuning, RL/preference data, LoRA/adapters, merged weights, model format, dependency libraries, loading code, the model repository, CI/CD and the signed release. Data poisoning can change general behavior or activate a backdoor under a trigger condition. Malicious fine-tuning or an adapter can remove safety alignment or implant a hidden objective. Executable deserialization formats such as pickle may run code when "loading a model." `trust_remote_code`, a custom dataset loader or template parsing widens the execution surface. A model hub's tokens, automatic builds, shared cache or release workflow may be exploited by a third party. Even with untampered weights, the tokenizer, configuration, evaluation scripts or dependencies may still be replaced.

Training data poisoning and malicious model files are different attack surfaces. Poisoning Language Models During Instruction Tuning, ICML 2023 [wan2023poisoning] reports that only 5 poisoned samples per task can reduce average performance by 38.8 points under its configuration. On an 11B model the reduction is still about 25 points. These are task performance changes, not jailbreak ASR.

In 2024, JFrog [cohen2024malicioushf] found malicious pickle models containing a reverse shell payload on Hugging Face. The public materials show no verifiable scale of infection, though. Wiz [tamari2024hfrisks] ran an authorized test on HF and demonstrated a potential cross-tenant path through malicious model loading, cloud identity and cluster configuration. That was a potential path, not a customer data breach that had already occurred. A vulnerability that could lead to arbitrary code execution was still present in `torch.load(weights_only=True)` in 2025 (GHSA-53q9-r3pm-6pq6 [github2025pytorch]). Together these cases show that "loading only the weights" is a joint guarantee. Format, library version, provenance, signature and execution environment must all hold.

Security practice favors non-executable formats, provenance hashes and signatures, and model/data/software bills of materials. Dependencies should be pinned, builds reproducible and artifacts scanned offline and statically. Untrusted artifacts should be loaded first in a disposable environment with no secrets, no external network and read-only host mounts. `safetensors` only reduces the code-execution surface of specific deserialization, however. It does not prove that the weights have no backdoor, that the configuration is not deceptive, or that model behavior is safe.

Defense Method Families

Defenses unfold along the same principal boundary axis. Each family answers questions about input, mechanism, output, cost, applicable conditions and failure modes. Each also separates probabilistic guardrails from enforceable system controls.

In-Model Probabilistic Control

This class of controls changes a model's refusal, classification or randomization behavior. It suits the role of a first line of defense that lowers the probability of triggering, but it does not provide an independent authorization guarantee.

Model Alignment, Classifiers, and Randomization: Necessary but Part of the Probabilistic Layer

Violation rates on known attack distributions can fall sharply under safety fine-tuning, preference optimization and constitutional classifiers. Consider the values reported directly in the paper. Constitutional Classifiers [sharma2025constitutional] reduced ASR from 86% to 4.4% in its synthetic universal-jailbreak evaluation. It also reports a rise of about 0.38 percentage points in the normal refusal rate and a 23.7% increase in inference compute. This gives a good example of jointly reporting "safety--utility--cost." It is still evidence under a vendor model, a vendor policy and a specific long-duration red-teaming setup, and it cannot be extrapolated into an execution-safety guarantee for arbitrary agents.

SmoothLLM [robey2023smoothllm] breaks some optimized suffixes through input perturbation and majority voting. Randomization can thus raise the cost of attack. The effect varies with attack type, perturbation rate and number of samples, and the method adds latency and token consumption. Input/output detectors likewise run into encoding, low-resource languages, cross-turn fragmentation and adaptive search. They suit screening, rate limiting and escalation handling. They should not be the sole authorizer of actions such as payments, database deletion or code execution.

A model-defense report needs at least four numbers. The first is attack success before and after the defense; then come normal task completion, over-refusal, and inference/human/latency cost. A "block rate" alone rewards a model that always refuses. Helpfulness alone rewards a system that lets dangerous actions through.

From "Why Safety Training Fails" to Adaptive Jailbreaking

Jailbroken, NeurIPS 2023 [wei2023jailbroken] studies safety training through black-box manual prompts. It examines competing objectives and mismatched generalization. Its lasting value is the finding that jailbreaking is not a magic string but a structural tension between task objectives and safety generalization. It does not establish that its prompts stay individually effective on all new versions, nor does it offer a unified denominator that could serve as an overall ASR.

GCG [zou2023gcg] optimizes transferable suffixes with white-box token gradients and coordinate search. It shows that out-of-distribution discrete strings can be found automatically, and it has become an important baseline for later adaptive attacks. It requires probability or weight access, so the cost differs for an ordinary API attacker. If one judges only by an affirmative opening, refusals, truncations or harmless text may be miscounted as successes.

PAIR [chao2023pair] runs few-query black-box iterations through an attacker model, driven by feedback from the target model and the judge. Manual prompt attempts become a repeatable optimization loop, and the query budget becomes a core metric. It also brings correlated error: if similar models do the attacking and the judging, their errors do not cancel out. Drift in the target API and in policy can also change reproduction results.

Tensor Trust, ICLR 2024 [toyer2024tensortrust] grew out of an online human attack-and-defense game. It holds roughly 563,000 attacks and 118,000 defense prompts. The corpus suits the study of how fast adaptation happens and how long a strategy survives once a defense is public. It is not a random sample of ordinary deployed users, though, and the attack frequency in the game cannot be read as a real-world incident rate.

Constitutional Classifiers [sharma2025constitutional] pairs a classifier-style constitutional defense with long-duration red teaming. It significantly reduces the reported ASR under its model and policy, and it also reports false refusals and inference cost. The model's probabilistic layer thus still has value, but it cannot replace action authorization. Nor can one set of results from a closed-source vendor be reproduced automatically across models.

One thread runs through this group of work. Early manual attacks explained "why it fails." GCG/PAIR turned attacks into an optimization problem. Tensor Trust observed human adaptation after public release. Defense research then used training, classification and perturbation to raise the cost of attack. Their shared boundary is that the primary endpoint is still model output. Once a model can execute actions, one has to move on to the next group of system-level evidence.

Structured Context, Provenance, and Information Flow

This class of controls prevents untrusted data from being silently promoted into instructions. It also makes tenant, provenance, integrity and confidentiality labels propagate as data is transformed.

Structured Context, Provenance, and Information Flow

StruQ [chen2025struq] and SecAlign [chen2025secalign] try to teach the model the structural boundary between instructions and data. Spotlighting [hines2024spotlighting] uses provenance marking to reduce the probability that external text is treated as an instruction. These methods do more than the single sentence "ignore untrusted instructions." Yet the model still interprets that boundary probabilistically. In ToolHijacker [shi2026toolhijacker]'s adaptive re-testing against malicious tool documentation, several StruQ/SecAlign configurations still showed a tool-selection attack success rate of 84.6%—99.6%. That does not

refute every result in the original paper, but it shows that generalization across attack surfaces cannot be assumed to hold.


System-level schemes recast the problem as control flow and information flow. Task Shield [jia2025taskshield] checks whether the candidate call is consistent with the original task before an action runs. CaMeL [debenedetti2025camel] separates trusted planning from untrusted data parsing, then executes through a restricted interpreter plus capability constraints. FIDES [costa2025fides] propagates integrity/confidentiality labels and enforces policies. One AgentDojo configuration of CaMeL reduced the successful attacks explicitly counted in the paper from 300 out of 949 cases to 0. For Gemini-2.5-Pro the no-attack task completion rate simultaneously fell from 73.2% to 41.2%, with median input/output token overhead of about 2.73x/2.82x. The figure 300 counts successful attacks here, and it must not be mixed with percentages elsewhere.

Structured defenses move the trusted computing base from "the model will obey" to labels, policies, interpreters and tool adapters. The cost is that these components must face code audits, differential testing and fail-closed design. Provenance labels can vanish during summarization, variable assignment or cross-agent passing. When they do, the information-flow guarantee goes with them.

Action-level Harness and Capability Gates

This class of controls demotes model output into a candidate plan. Five gates---tool registration, task, data flow, parameters, and consequences---decide whether to execute.

Figure fig:text-harness-results gives the per-configuration results of the local synthetic mechanism probe for the text harness.



The capability gate did not remove dangerous intent; it stopped dangerous actions from landing.

Text Harness synthetic reproduction. The capability gate did not eliminate dangerous intent, but it did stop dangerous actions from landing. The denominator holds only six synthetic scenarios, so the results serve mechanistic illustration rather than model ranking.

Text Harness synthetic reproduction. The capability gate did not eliminate dangerous intent, but it did stop dangerous actions from landing. The denominator holds only six synthetic

scenarios, so the results serve mechanistic illustration rather than model ranking.

Harness: Demoting Model Output into Proposals

Every tool call a model generates should clear at least five independent gates, in sequence. The registration gate confirms that the tool, publisher, version, schema and binary digest sit in the read-only signed registry. The task gate checks that the tool and action belong to the minimum allowed set precomputed for this user request, and that no new high-impact subgoal has appeared unnoticed. The data flow gate traces where each parameter came from: the user, a trusted directory, an untrusted web page, a model guess or low-integrity memory. It then asks whether the recipient is entitled to this data. The parameter gate goes beyond JSON Schema to check recipient, amount, path, URL, HTTP method, idempotency key and business limits. The consequence gate judges last whether the action is reversible, whether it crosses trust domains, and whether it needs a dry-run, a two-phase commit or parameter-bound one-time human approval.

The model can help interpret the 2nd gate. It should not generate the action, judge whether that action is safe and approve itself all at once. Least privilege means more than "installing fewer tools." Capability, object, parameter, time and count all have to be narrowed. Deletion and read-only queries should be different capabilities. Access to one repository must not expand automatically into access to the whole organization. Approval must bind to a normalized parameter hash, with re-confirmation after parameters, redirects or tool versions change.

MCP security best practices [mcp2026security] require validating the token audience, obtaining per-client consent and prohibiting token passthrough. That handles part of the identity and confused-deputy risk. The protocol will not recognize malicious tool descriptions by itself, however, and it will not prove that a model call matches the user's true intent. The harness supply chain therefore includes the MCP server, the tool metadata and the authorization proxy.

The harness mechanism reproduction in this survey is fully synthetic and has zero real side effects. It runs on a local Qwen3 8.2B. Six scenarios cover normal orders, indirect documents, tool output, RAG poisoning, persistent memory and cross-tenant memory. Four configurations were run: bare harness, prompt guard only, capability gate only and defense in depth. The bare harness showed dangerous intent, dangerous execution and benign task success of 3/6, 3/6 and 4/6. The capability gate only still showed 3/6 dangerous intent, yet it cut dangerous execution to 0/6. Defense in depth v2 showed 1/6, 0/6 and 5/6. The contrast exposes the logic underneath. The capability gate need not make the model reliable; its job is to keep an unreliable planner from gaining side effects on its own.

Two failures matter more than the ranking sentence "layering is best." The first version of layered filtering deleted all untrusted data. Dangerous intent fell to 0/6, but benign tasks reached only 4/6. Version 2 restored the ordinary documents and tool facts that tasks need, so utility rose to 5/6. That also brought back 1/6 dangerous intent. The second failure: under all four configurations, indirect documents leaked the synthetic canary and proposed outbound

sending. The capability gate did no more than block the execution. Input boundaries, secret invisibility and action authorization are therefore all indispensable. The text harness reproduction report holds the complete inputs, per-scenario outputs, Wilson intervals and paired calculations.

From Indirect Injection to Executable Agents

In *Not What You've Signed Up For* [greshake2023indirect], attackers place instructions inside web pages, emails or retrieved content. The user application then reads that content on its own initiative. This establishes the remote indirect injection threat model in which "external data becomes control flow." The paper offers an early systematic analysis of feasibility and impact classification. It does not provide cross-application incident rates.

The prompt injection unified benchmark, USENIX Security 2024 [liu2024formalizing], cross-evaluates 5 attacks, 10 defenses, 10 LLMs and 7 task types. Its most important finding is not some best number. Defense rankings change with the model, task, attack and metric. All configuration cells share data and method, so they cannot be averaged as independent studies.

AgentDojo, NeurIPS 2024 [debenedetti2024agentdojo], places injection through external tool results into 97 benign tasks and 629 security cases. It measures benign completion and attack-goal completion at the same time. That lays the foundation for the security--utility two axes. Its tools and user data remain a simulated environment. They cannot represent the incident rates of real OAuth, production data and irreversible transactions.

InjecAgent, ACL Findings 2024 [zhan2024injecagent], uses 1,054 cases, 17 user tools, 62 attacker tools and 30 agent configurations. It shows that tool return values and tool selection can push indirect injection into the action layer. A model that outputs a tool call string still does not equal successful authorization or a completed real side effect. Reproduction must therefore check the executor state.

Task Shield, ACL 2025 [jia2025taskshield], judges before the action whether a candidate call serves the original user task. It moves the detection locus from the attack string to goal--action consistency. It still cannot handle on its own the cases where the action type is correct but the recipient, amount or data reader is wrong. It must therefore be combined with parameter policy.

CaMeL [debenedetti2025camel] and FIDES [costa2025fides] move the security boundary from model self-discipline to an auditable interpreter. They do so through trusted planning, isolated parsing, capability or information flow policy, and a reference monitor. That also turns labels, tool read/write sets, policy and adapters into a new trusted computing base. Their utility and token cost must be reported with the security results.

This group shows that the evaluation endpoint has undergone a fundamental migration. It has moved from "did the model obey the injection" to "did the application complete the attack action." The most desirable direction is not to guess more cleverly which sentence is malicious. It is to keep untrusted data from creating control flow on its own, and to keep control flow from crossing independent permission and information flow policy.

RAG and Memory Database Governance

This class of controls governs the index and memory as a security database, with identity, provenance, TTL, version, derivation chain and tombstone, not as a longer prompt.

RAG and Long-term Memory: Governing by the Standards of Databases and Identity Systems

Before ingestion, RAG requires a source allowlist, content hash, signature/crawl time, malicious instruction scanning, and human/automatic quarantine. The index is physically or logically partitioned by tenant and sensitivity level. ACLs are enforced before retrieval rather than patched up after the text has entered the model. At retrieval time, top-k, the share of any single source, and anomalously similar records are limited, and answers carry verifiable citations. Scanning can only reduce obvious poisoning. It cannot replace identity isolation and action gates.

Long-term memory records should not contain only text and embedding. They should carry at least tenant_id, subject, provenance, write channel, type, integrity/confidentiality labels, reader ACL, purpose, TTL, version, derivation chain, review status and tombstone. At write time, external content, tool results and model reasoning are low-integrity by default, and the model cannot approve its own trust upgrade. At read time, filtering by tenant, subject, purpose, type, ACL, status and TTL comes before vector similarity, and low-integrity memory cannot directly supply key parameters for high-impact actions.

Composition and update must also preserve evidential relations. Ten summaries from the same source are not ten independent pieces of evidence. Cross-turn shards have their complete derivation chain re-examined at action time. An update generates a new version and preserves derived_from, conflicts enter disputed, and the latest write is not allowed to silently overwrite. Deletion first synchronizes the tombstone and withdraws the online index. Only then does it delete the original text, vectors, summaries, caches and exports. Backup restoration replays the tombstone and retains deletion proof.

Work in 2025--2026 such as A-MemGuard, MemGuard, FARMA/SENTINEL and FragFuse has begun to cover write poisoning, type isolation, dormant triggers and compositional attacks. Most of it, however, is still rapidly evolving preprints or new papers. Cross-tenant, multi-month persistence, backup restoration and verifiable forgetting evidence remain thin. This survey therefore lists the specific mechanisms as emerging empirical evidence. It lists the above schema, ACL, TTL and deletion procedures as normative engineering recommendations. It does not claim that memory defenses are already mature.

RAG and Memory: From Corpus to Cross-session State

PoisonedRAG, USENIX Security 2025 [zou2025poisonedrag], injects malicious documents that simultaneously pursue retrieval hits and the target answer for attacker-chosen queries. It shows that a small number of targeted records can significantly influence a specific query. Analysis must separate whether a document enters the top-k, whether the generator adopts it, and whether the final answer hits the target. Moreover, 5 targeted documents per target

question does not equal 5 documents being able to control an entire knowledge base without a target.

Spill the Beans, ICLR 2025 [qi2025spillbeans], repeatedly induces a production RAG through black-box queries to retrieve and output private source text. It recovers 41% from a corpus of about 77,000 words and 3% in a larger-corpus configuration. The unit here is word coverage rather than prompt ASR. The conclusion is that a knowledge base should be protected as a queryable database, not by treating the system prompt as a confidentiality boundary.

AgentPoison, NeurIPS 2024 [chen2024agentpoison], uses low-ratio knowledge or memory poisoning to make a trigger query retrieve malicious experience. The official page reports an average ASR of no less than 80% across three agent classes, and a clean-performance impact of no more than 1%. The attack prerequisite is obtaining a write path. Cross-agent aggregation cannot be treated as the probability of any single model.

MINJA, NeurIPS 2025 [dong2025minja], influences long-term memory only through ordinary conversation. That pushes the attack entry point from directly writing to the database to the product's auto-save policy. It requires measuring write, future recall, attack action and number of surviving turns separately. A single overall ASR is insufficient to characterize persistence.

A-MemGuard [wei2025amemguard] attempts to reduce memory poisoning through multi-path consensus. FragFuse [rao2026fragfuse] instead splits a violating goal into cross-turn benign fragments, showing that per-item checking can still be bypassed by composition. The point of understanding this new group of results is whether sources are independent, whether the derivation chain is preserved, long-term utility and token cost. It is not to compare only the lowest and highest ASR.

What RAG and memory share is that state can be poisoned first and triggered later. They differ in that RAG mainly revolves around shared knowledge and retrieval permissions, while memory is also bound to user identity, time, preferences and system learning. Both need provenance/ACL preserved beyond summaries and embeddings and re-checked at action time. Neither may mistake "already ingested" for "already trusted."

Multimodal Parsing and Action Confirmation

This class of controls explicitly surfaces the provenance of OCR, ASR, DOM and visual regions. It returns to the original goal and structured policy before high-impact actions.

Multimodal Defense: Making Explicit What Each Parser Sees

VLM/VLA systems should present OCR, ASR, page structure and recognized regions as provenance-carrying data. They should not silently splice them into the system instructions. Text inside an image is by default an object of analysis, and high-privilege actions must return to the original user goal and the structured tool policy. High-risk visual agents can adopt cross-validation with independent OCR/vision parsers, region-level provenance labels, pre-action screenshot and difference confirmation, and detection of invisible text layers, overlapping elements and cross-frame instructions.

This survey also attempts a minimal nine-cell reproduction with a local llama3.2-vision, using three synthetic images and three prompt boundaries. The first two requests of the default command timed out after about 240.02 seconds and 240.00 seconds respectively with no response. The remaining default runs were aborted rather than occupy the same CPU queue for another half hour. The nine-cell error coverage was then completed in a separate directory with a 5-second timeout: 9/9 requests timed out and 0 valid answers. The Wilson 95% interval for the nine-cell timeout rate is 70.1%—100%. It describes only the current local run layer, not an attack success rate. The initial summarizer once wrote the empty responses as 0.0. After the fix, the script first computes a valid-answer denominator, all three safety and utility ratios output NA, and the manifest explicitly marks `completed_no_valid_responses`. The full failure trajectory and per-cell errors are given in the VLM image prompt injection reproduction report.

AdaShield [wang2024adashield] reduces the reported attack rate on its QR structural attack from 75.75% to 15.22% for LLaVA and from 83.62% to 1.37% for CogVLM. The LLaVA-7B post-hoc configuration of VLGuard [zong2024vlguard] goes from 90.40% to 0 on FigStep. These numbers indicate that dedicated multimodal safety data and prompt pools can patch obvious gaps, but they cannot be hard-merged across tasks. The safety-only training of the latter also lowers the XSTest safe metric from 91.2 to 41.6, demonstrating the risk of over-refusal. Image compression and random transformations may destroy adversarial perturbations. They cannot reliably handle malicious text that is clearly legible to the human eye.

Audio, video, GUI and embodied systems also have to handle temporal synchronization and physical consequences. ASR transcripts should retain timestamps and confidence, cross-frame/background speech requires combined checking, and click or movement actions require non-model constraints such as reachable regions, speed, collision and emergency stop. Current evidence is concentrated on static images. VLM defense results cannot be extrapolated to these scenarios.

Multimodal: Typographic, Representation, and Temporal Attacks Cannot Be Treated as One Class

FigStep, AAAI 2025 [gong2025figstep], typesets harmful requests into images, asks the text side only to complete the steps, and reports an average ASR of 82.5% across six open-source LVLMS. It exposes the insufficient transfer of text safety alignment to the visual channel. Its mechanism, however, is closer to an OCR/typographic carrier, and the cross-model average has no unified binomial denominator.

Image Hijacks, ICML 2024 [bailey2024imagehijacks], optimizes pixels or representations in a white-box setting so that the image controls generation at runtime and transfers across prompts. This shows that an image can play a role similar to a program. Its realism depends on white-box cost, perturbation constraints and the specific preprocessing. It therefore cannot be placed in the same effect group as typographic text images that are clearly legible to the human eye.

SpeechGuard, ACL Findings 2024 [peri2024speechguard], studies audio adversarial perturbations and cross-model transfer. The gap between white-box direct injection and black-

box transfer results is large. Volume, environment, codec, ASR front end, and real playback/recording add further variables. Neither text nor static-image ASR can be extrapolated to audio.

AdaShield [wang2024adashield] retrieves defense prompts for structural visual jailbreaks and substantially lowers the reported ASR in the QR/FigStep configurations with little change in utility. It is still a probabilistic prompt layer, however, and does not exhaust the white-box case where the attacker jointly optimizes image, text, similarity threshold and guard.

VLGuard [zong2024vlguard] uses safety image-text data for post-hoc or mixed fine-tuning to patch the safety forgetting after visual fine-tuning, and brings several FigStep configurations close to 0. Yet safety-only training causes pronounced over-refusal. It must therefore be evaluated jointly with helpfulness data, normal visual tasks and out-of-model action control.

Multimodal evidence must record human visibility, attack constraints, font/resolution, OCR/ASR, preprocessing, fusion method and whether it is physically deployed. A PDF is a composite container that simultaneously holds a visible layer, hidden text, object order and OCR. A GUI is perception plus action. Video and audio additionally have cross-temporal combination. Using a single multimodal=true label for stratification is acceptable, but effects must not be merged on that basis.

Sandbox, Network, Secrets, and Supply Chain

This class of controls sits at the execution and artifact boundary. It separately restricts the kernel, files, network, identity, secrets, resources, and the build-and-release path.

The selection logic for execution isolation carriers and external capability controls appears in Figure fig:sandbox-capability-selection.

Choosing an isolation boundary by workload and required capability

Dimension	Weak	Medium	Strong
Execution	Same process	Container	microVM
Network	Open	Allowlist	No egress
Secrets	Environment variables	Short-lived token	External broker
Resources	Uncapped	Quota	Hard cap

Execution, network, secret and resource limits are mutually independent and cannot substitute for one another.

Selecting an isolation boundary by workload and required capabilities. Execution, network, secret, and resource restrictions are mutually independent, and the options in the figure are not absolute security levels.

Selecting an isolation boundary by workload and required capabilities. Execution, network, secret, and resource restrictions are mutually independent, and the options in the figure are

not absolute security levels.

Sandbox, Network, Secrets, and Resources: Four Controls That Cannot Substitute for One Another

"The process is in a container" or "there is no direct public network" does not equal closed isolation. The execution boundary must answer separately: which kernel interfaces the code can reach, which files and devices it can see, which addresses it can connect to, what identity it holds, and how many resources it can consume. The selection relationships for different workloads appear in the sandbox and capability selection figure.

A reasonable starting point for untrusted arbitrary code or network security evaluation is a short-lived microVM per task. It comes with no long-term secrets, outbound denied by default, read-only inputs, an empty temporary disk and a host-external policy gate. Residual risks still include VMM/kernel zero-days, the device surface and already-allowed egress. Tool or browser tasks that need higher Linux compatibility can start from a gVisor-class user-space kernel or a hardened VM. They also require non-root, all capabilities dropped, seccomp, a separate browser profile and a network proxy. Such a setup still faces compatibility gaps, logic abuse within permitted actions and browser zero-days.

Small plugins or deterministic transformations are better suited to the WASM/WASI capability model. The model grants only explicit directories, sockets, clocks and randomness sources, and limits fuel and memory. Overly broad host functions or runtime vulnerabilities will still break the boundary. Tasks that only need data parsing should instead use a tool-free isolated parser, with output constrained by type, length and provenance. This reduces control flow, but parser vulnerabilities and contamination propagation still have to be handled.

Outbound traffic should be denied by default, and the allowlist refined down to protocol, host, port, path and method. DNS and redirects are re-checked hop by hop. Loopback, private-network, link-local and cloud metadata addresses are rejected. Secrets do not enter the prompt, ordinary environment variables or shared files. After the policy passes, a secret broker issues short-lived, audience-bound, minimum-scope tokens to specific tools. Tokens, cookies and signed URLs are stripped before results return to the model. Resource control covers wall clock, model tokens, steps, retries, concurrency, CPU, memory, PID, file descriptors, disk, I/O, network bytes and total cost. When any budget is exhausted, the harness terminates the run. The model must not be allowed to scale itself up.

This survey performed a mechanism reproduction on macOS that involves no external target. It crossed 3 local policies with 6 capability probes for 18 cells, and the observed results agree with the preset allow/deny matrix 18/18. With path isolation alone, out-of-bounds file reads were blocked, but network connections and subprocesses were still available. Both tightened at once only after network and process restrictions were added. That experiment used the deprecated sandbox-exec and attempted no escape, no real secrets and no external attacks. It therefore demonstrates only that "controls must be configured separately"; it does not demonstrate that a production sandbox is secure. See the sandbox and capability gate reproduction report for details.

Model, Data, Tool, and Harness Supply Chain

A runnable model depends jointly on weights, adapter, tokenizer, chat template, processor, generation config, custom code, dependencies, container, GPU extensions, system prompt, tool schema and policy. A security gate pins the repository, commit, publisher, license, signature and digest, and forbids production from following a floating main/latest. SLSA/in-toto-class provenance then records the chain through training, conversion, quantization, evaluation and packaging. For format and loading, non-executable weights are preferred, and default pickle and automatic trust_remote_code are forbidden. The first load should take place in a temporary environment with no secrets, no external network, read-only sources and strict resource limits.

Content checks need to cover tensor shape, dtype, stride, shard manifest and outliers. They also need to generate an SBOM and a vulnerability list for the code and dependencies. Once the behavioral gate passes, the complete artifact is re-signed and enters an internal read-only registry. Models, adapters, containers, tools and policies all need a joint revocation and rollback mechanism.

Hugging Face's pickle scanning, safe tensor formats and repository malicious-file detection each cover different risks. safetensors removes the pickle-style surface for arbitrary Python object execution, but it does not prove that the weights have no backdoor. A provenance signature proves that the artifact was not substituted, but it does not prove that its behavior is safe. The 2025 PyTorch weights_only deserialization vulnerability [github2025pytorch] further shows that even a nominally safe switch requires a pinned patched version and verification in an isolated environment.

Privacy, Backdoors, Supply Chain, and Execution Isolation

Training-data extraction research shows that a model will leak training sequences under specific repetition, prefix and sampling budgets. That differs from direct exfiltration from the current context or a RAG database. Sleeper Agents [hubinger2024sleeper] shows that conditionally triggered deceptive behavior may survive supervised fine-tuning, reinforcement learning and adversarial training. But it is a controlled backdoor experiment and cannot prove that real-world models generally have an "autonomous conspiracy". Both lines of research require behavioral evaluation before release, not merely verification of file hashes.

Firecracker [agache2020firecracker], gVisor, and WASM papers/documentation mainly argue isolation mechanisms, performance and attack surface; they do not directly test LLM injection ASR. AgentDojo-class papers mainly test model and harness behavior, yet they usually have no real kernel, cloud credentials or network egress. The largest current cross-domain gap is precisely putting the two into the same end-to-end experiment. An attacker controls an agent through untrusted content. The agent attempts to read secrets, reach out, move laterally and exhaust resources. The system reports dangerous intent, actual blocking, benign utility, cost and escape surface at the same time.

Supply chain evidence likewise falls into three classes. Malicious repository/weight incidents prove feasibility. CVEs such as PyTorch prove concrete implementation defects. Specifications and tools such as safetensors/SLSA/in-toto define control mechanisms. Only by combining

"traceable provenance + no-code format + isolated loading + behavioral gate + revocation" can a system cover the execution and the behavior path at the same time.

Monitoring, Red-Teaming, and Incident Response

This class of controls shortens detection and recovery time through replayable traces, anomaly signals, capability revocation, credential rotation and environment reconstruction.

Monitoring, Red-Teaming, and Incident Response

The runtime should string the user's original goal, planning version, provenance/taint, tool selection, normalized parameters, authorization decisions, approvals, network, memory writes and deletions, sandbox events and resource budget into a replayable trace. Sensitive values are encrypted separately from the primary log, and the log itself uses append-only integrity protection. Important signals include new domains or cross-domain redirects, low-integrity data participating in high-impact parameters, tool/schema drift, cross-tenant hits, secret patterns, looped retries and anomalous costs.

Red-team tools and benchmarks are for finding regressions, not a security certificate. Every change to a model, retriever, tool, policy, image or memory schema should rerun fixed seeds and the latest adaptive attacks, and measure utility, false refusals, latency and cost. After a suspected boundary violation, first freeze new actions, revoke short-term capabilities, block egress and preserve forensic snapshots. Then rotate credentials, rebuild the affected environment, clean memory/index derivatives and replay the scope of impact. Finally, turn the fix into an automated gate and canary. NIST's Generative AI Risk Framework, MITRE ATLAS and OWASP can provide control and threat vocabulary, but the real pass criteria must land on this system's actions, assets and logs.

[← Back to contents](#)

Cross-Family Synthesis and Selection Guide

Under the unified trust-boundary axis, cross-family comparison is not about whose single ASR is lowest. It asks which layer of control changes reachability, state control, dangerous intent, action authorization or final harm.

Despite the proliferation of attack names, six categories of root cause can be summarized across topics. Where instructions and data lack a strong boundary, natural language carries content and control at the same time, so the model can only judge priority probabilistically. Where authorization is completed at the wrong layer, the model both decides the action and decides whether it is itself authorized to execute it, rather than an independent policy judging according to the user, the task, the resources and the consequences. Where provenance and identity are lost during transformations, only plain text remains after retrieval, OCR, summarization, compression and cross-agent forwarding, with no provenance, tenant or sensitivity labels.

The other three categories of root cause lie in state, isolation and evaluation. When state writes are more permissive than reads, a single web page or tool output is allowed into long-term memory, an index, a cache or a code repository. Isolation boundaries have uncanceled exits. Package proxies, debugging tools, metadata services, shared volumes, long-lived tokens and third-party sandboxes can still cross trust domains, so "no public network" or "already isolated" must be subject to end-to-end verification. When evaluation treats correlated cells as independent evidence, model versions, prompts, judges and sample reuse produce spurious statistical precision and conceal cross-scenario failures.

The defense chapters that follow will be organized around these root causes, rather than providing easily outdated string rules for each attack name one by one.

Four Layers of Defense in Depth and Their Failure Modes

There is no single "safety prompt" that solves jailbreaks, indirect injection, memory poisoning, malicious weights and sandbox escapes at once. A more useful way to organize the discussion is to ask whether the next layer can still limit the consequences after one layer of control fails. The unified taxonomy figure already draws the relationships among the four layers, and they are explained paragraph by paragraph below.

The first layer is probabilistic model defense. Safety fine-tuning, refusal policies, input/output classifiers, randomized smoothing and multi-model review can reduce the frequency of known harmful content, of some jailbreaks and of obvious injections. They cannot on their own guarantee permanent robustness against adaptive attacks. Nor can they grant real permissions or prove the absence of side effects.

The second layer is structured control and information flow. Instruction/data channels, provenance labels, taint propagation, typed variables, task-alignment checks and RAG ACLs stop data from being promoted into instructions, prevent cross-tenant reads and keep sensitive values from flowing to the wrong tool. Their guarantees still depend on whether the labels, policies, adapters and underlying code are correct.

The third layer is system isolation and capability constraints. Least privilege, action gates, short-lived tokens, network egress, secret brokers, microVM/gVisor/WASM and resource budgets separate dangerous intent from dangerous execution, and they limit lateral movement, exfiltration and runaway resource consumption. They still cannot automatically identify business-logic abuse inside the allowed set, nor can they guarantee that the allowed domain and trusted tools will never be compromised.

The fourth layer is operational governance and response. Signed supply chains, version gates, continuous red teaming, end-to-end traces, alerting, revocation, forensics and disclosure address drift, unknown attacks and persistent impact after compromise. They cannot provide absolute prevention in advance, but they determine whether a team can reconstruct the true history, narrow the impact and turn remediation into a long-term gate.

The first layer reduces attack frequency and manual burden. The second and third layers determine whether an attack can cross data, identity and execution boundaries. The fourth

layer determines whether the system can detect and recover from unknown failures. High-impact actions require at least independent controls from the latter three layers. A second self-judgment by the same model does not count as an "independent line of defense".

Cross-Paper Synthesis: What Can Be Compared and What Must Be Kept Separate

Mechanisms are what can be compared stably. Adaptive search usually weakens static string defenses. An independent policy closer to the action constrains actual consequences more effectively. If provenance and tenant labels do not propagate with data transformations, later layers cannot restore them. There are also genuine trade-offs among safety, utility and cost. What cannot be compared directly is violation generation rates under different policies, the fraction of compromised applications, word-level corpus recovery rates, tool goal completion rates, memory survival rounds, detection accuracy and real incident counts.

This survey therefore does not declare the "strongest attack/defense" on the basis of a cross-paper ASR leaderboard. For every number, three questions are asked first. What is the denominator? At which layer does success occur? Does the attacker know the defense? If any item is unclear, the result can at most enter a descriptive evidence map, and it cannot enter a merged effect.

[← Back to contents](#)

Data, Metrics, and Evaluation Evidence

This chapter separates method mechanisms from the evaluation contract. It explains item by item the unit of analysis, event counts and denominators, effect direction, heterogeneity, within-paper dependence and prohibited interpretations. The statistical procedures have actually been run, but the number of formal meta-analytic groups is zero.

Unified Parsing Template

For the anchor papers, this survey answers eight questions uniformly. What asset does the research protect? Is the attacker black-box, gray-box or white-box? Which entry point can be written to? How large is the attack budget, and is the attack adaptive? What are the sampling unit and the definition of success? What are the main results and benign utility? Can the work be reproduced? Where can the conclusions be extrapolated? Only the anchors that determine the research thread are kept below. The complete per-paper records of the 23 attack papers and the 14 groups of defense/engineering anchors are given in the attack-surface search and the defense-engineering search, respectively.

Four Methods Broken Down into Input, State, Output, and Scoring

Case A: Why GCG can search out a gibberish suffix. The input is an aligned model f_{θ} , a set of harmful target requests, a modifiable discrete suffix token and a target beginning. The internal state is the gradient of the target loss at each suffix position. The algorithm does not directly publish an answer on continuous vectors. Instead, it uses the gradient to shortlist a few candidate tokens at each position, then replaces them one by one and uses the true forward loss to select the candidate with the largest decrease, looping until the budget is exhausted. The output is a suffix. Scoring usually looks at both the target prefix and the final content. The underlying logic is that safety training constrains the common semantic distribution, yet it does not make all discrete token combinations satisfy the same refusal boundary. Reproduction must freeze the tokenizer, chat template, target prefix and number of steps, otherwise “GCG with the same name” is not the same experiment. Looking only at affirmative beginnings will overestimate harmful completion.

Text structure representation in the old draft

```
Request u + learnable suffix s
v Target loss L(ftheta(u||s), target)
Compute the gradient at each position of s to produce top-k token candidates
v Score each candidate with the true forward pass
Accept the best replacement, repeat for B steps
v
Candidate suffix + full answer from the target model + independent/manual scoring
```

Case B: Why CaMeL/FIDES is not “adding one more guard model”. The input is divided into an immutable user goal U and untrusted external data D . A privileged planner sees only U and produces a restricted plan. A tool-free parser converts D into typed values. The runtime maintains the provenance, integrity, confidentiality and permitted recipients of values. The output is not a shell that the model executes directly. It is actions that the interpreter passes one by one after they satisfy the capability and information-flow policies. Security comes from “ D cannot create new control flow, and low-integrity or high-confidentiality values cannot enter sinks that are not permitted”, not from the parsing model always being correct. In CaMeL's Gemini-2.5-Pro configuration, the successful attack count goes from 300/949 to 0/949, but normal task completion drops from 73.2% to 41.2%, and the median input/output tokens are about $2.73 \times / 2.82 \times$. The conclusion must therefore include blocking, utility and cost at the same time, and it cannot state only “0 attacks”.

Text structure representation in the old draft

```
Trusted user goal U --> privileged planner --> restricted plan/capability ceiling
Untrusted data D --> tool-free parser --> labeled structured values
The two paths meet at the reference monitor
v
policy allow / deny / ask / redact
v
Tool adapters produce real side effects
```

Case C: Why memory attacks need at least three endpoints. AgentPoison assumes that the attacker can write a small number of optimized records into the knowledge/memory store. MINJA narrows the entry point to an ordinary conversation triggering automatic saving. Neither attack is “successful on a single input”. Success instead runs along a chain: write success $W \rightarrow$ future recall $R \rightarrow$ dangerous action A . When a paper reports only the final ASR, a reader cannot tell whether a defense blocks the write, reduces top-k hits, or rejects at the action gate. A more reasonable record therefore carries $P(W)$, $P(R|W)$, $P(A|R,W)$, survival time, cross-tenant leakage, and clean-task change. End-to-end risk can be conceptualized as the product of the three conditional probabilities. Actual multiplication, however, is permitted only with staged counts from the same traced sample. Nor do ten memories that are provenance-related count as ten independent pieces of evidence.


Case D: How to read the numbers of FigStep/AdaShield. FigStep renders the harmful text as an image, while the outer text only asks for the steps to be completed. That input passes through resize/visual encoding/cross-modal fusion. The output is then judged as success by refusal keywords or by an LLM judge. AdaShield retrieves defense prompts based on input similarity and appends them to the VLM context. The retrieval changes the model's interpretation of the text in the image. In the LLaVA QR configuration the value goes from 75.75% to 15.22%. Within the study that means a decrease of 60.53 percentage points, and the ratio of the reported values is about 0.201. Without exact event counts, no standard error can be given. That VLGuard's FigStep is close to 0 does not mean utility is free. Safety-only training lowers XSTest safe from 91.2 to 41.6. The underlying logic is that defense training covers the safety distribution of the visual channel. It may also mislearn normal visual help as refusal. A three-axis judgment of “harmful completion + normal answers + visual task score” is therefore needed.

Build an Evidence Map First, Not an Average First

Three different statistical units run through this project: 65 attack-paper records, 62 defense/engineering sources, and 32 incident/vulnerability records. They cannot be added into “159 papers”. The defense table contains specifications and official documentation, so it is not a set of papers. The unit of the incident table is also not a paper. And the three tables may cite the same source. So the figure below reports three things in separate panels: evidence level, year, and non-mutually-exclusive modality labels.

Figure fig:evidence-map shows the heterogeneous evidence map of attack papers, defense sources, and incident records.

Evidence sources and units of analysis



The three source types use different units of analysis; they cannot be summed into one total.

Evidence map. Attack papers, defense and engineering sources, and event records use different units of analysis and cannot be added into a single total number of papers.

Evidence map. Attack papers, defense and engineering sources, and event records use different units of analysis and cannot be added into a single total number of papers.


Automated retrieval is also not the same as final inclusion. The 2,400 records are API returns from 12 OpenAlex queries. After deduplication, the 1,854 records are still unscreened candidates. Of those, 701/297/856 are only high/medium/low machine priority, and 500 are the read-first queue. Citation tracking, conference pages, standards, vulnerability databases, and vendor announcements also feed the manual base table. The exclusion log is still incomplete for each title, abstract, and full text. This survey therefore does not fabricate a “final PRISMA number of included papers”.

Effect Inclusion Rules

Every computable effect first answers five questions. Is the success endpoint explicit? Are there direct event counts and denominators? Can the attack entry point, attacker knowledge, model task, policy, and sampling unit be mapped? Does the same paper contribute only one pre-declared main arm? And are there at least three independent studies? If any step fails, the analysis falls back to descriptive synthesis or single-study presentation. It does not continue to apply statistical formulas. The judgments in full, and the actual number of dropouts in this round, are plotted below.

Figure fig:meta-composability shows the poolability audit, from 34 studies to zero formal pooled groups.

Meta-analytic poolability audit



The five studies passing the numeric contract fall into five different endpoints; nothing is pooled.

Meta-analytic poolability audit. The five studies that pass the numerical contract fall into five different endpoints, all groups are single studies, and therefore there are zero formal pooled groups.

Meta-analytic poolability audit. The five studies that pass the numerical contract fall into five different endpoints, all groups are single studies, and therefore there are zero formal pooled groups.

“Selecting one main arm in advance” is not picking the row with the best effect. The order is frozen. The main experiment of the formal paper takes priority. Among models and attacks, the ones that best match the group definition take priority. Adaptive attacks take priority over non-adaptive attacks. Human or dual scoring takes priority over a single keyword. Direct event counts take priority over percentages only. The remaining arms stay in the raw table for sensitivity description. They cannot be disguised as independent papers to enlarge the sample size.

Effect Size for the Attack Success Proportion

Study i observes x_i successes among n_i independent samples. Then:

$$p_i = \frac{x_i}{n_i}, y_i = \text{logit}(p_i) = \log \frac{p_i}{1 - p_i}, v_i \approx \frac{1}{n_i p_i (1 - p_i)}.$$

The logit is chosen rather than directly averaging percentages for two reasons. The transformed interval does not cross 0—1. And the same 5-percentage-point difference has a different statistical meaning near 5% and near 50%. When $x_i=0$ or $x_i=n_i$, the script applies the continuity correction $(x+0.5)/(n+1)$ in the transformation stage only, and the raw records still keep the true 0 or n . Sometimes only the author-reported proportion and an explicit n are available. The script can then estimate the variance from that proportion, but it marks the source as `reported_asr_total`. It will not back out a “precise event count” from a rounded percentage.

Hand-calculation example: HOUYI. The application-level sample is $x=31, n=36$, and the point estimate is $p=31/36=0.861$. The Wilson 95% interval gives:

$$\frac{p + z^2/(2n) + / - z\sqrt{p(1-p)/n + z^2/(4n^2)}}{1 + z^2/n}, z = 1.96,$$

This gives approximately 0.713—0.939. The interval expresses only the sampling uncertainty of “if these 36 applications are treated as a binomial sample”. It cannot repair application selection, shared backends, version drift, or disclosure bias. So it cannot be read as the industry-wide 95% true range. This is exactly why a statistical interval cannot replace a threat model review.

Pre- and Post-Defense Effects: Risk Ratio Rather Than a Percentage-Point Mixture

Write the pre- and post-defense values as x_0/n_0 and x_1/n_1 . Then use the risk ratio:

$$RR_i = \frac{x_1/n_1}{x_0/n_0}, y_i = \log(RR_i),$$

The approximate variance under independent binomials is:

$$v_i \approx \left(\frac{1}{x_1} - \frac{1}{n_1} \right) + \left(\frac{1}{x_0} - \frac{1}{n_0} \right).$$

$RR < 1$ means lower risk after defense. If any event count is 0, the script adds 0.5 to each arm's event count and 1 to each denominator, which avoids log 0. Small-sample results then depend strongly on that correction, so extreme proportions must be checked against the original paper. Many defenses are paired before and after on the same prompts. An ideal analysis requires a paired four-cell table or a McNemar/conditional model. Papers usually do not publish the paired transition matrix, and the current script's independent binomial approximation ignores the correlation. It is therefore explicitly labeled exploratory.

How to interpret when only percentages are available. Constitutional Classifiers reports 86%→4.4%. That gives a decrease of 81.6 percentage points, a descriptive risk ratio of $(4.4/86=0.051)$, and a relative decrease of about 94.9%. No precise denominator can be used safely here. Without one, a credible standard error cannot be computed, and the result cannot enter formal pooling. AdaShield's LLaVA QR moves from 75.75% to 15.22%. That is a descriptive decrease of 60.53 percentage points, with a ratio of about 0.201. This is still only a comparison within the same study configuration, and it cannot be averaged with the former.

Random Effects and Paule-Mandel Heterogeneity

No single identical true effect is assumed to exist, even if the studies are judged comparable. The model is:

$$y_i \in \text{Normal}(\theta_i, v_i), \theta_i \in \text{Normal}(\mu, \tau^2),$$

where v_i is the within-study variance and τ^2 is between-study heterogeneity. The weights are:

$$w_i = \frac{1}{v_i + \tau^2}, \hat{\mu} = \frac{\sum_i w_i y_i}{\sum_i w_i}, SE(\hat{\mu}) = \sqrt{\frac{1}{\sum_i w_i}}.$$

The script finds τ^2 by the Paule-Mandel method, choosing the value at which the weighted residual $Q(\tau^2)$ approaches the degrees of freedom $k-1$. It also reports:

$$I^2 = \max\left(0, \frac{Q - (k - 1)}{Q}\right),$$

The script then gives the 95% interval for the mean, $\mu\text{-hat} \pm 1.96SE$, plus the 95% prediction interval, which is more suitable for deployment judgment:

$$\hat{\mu} \pm 1.96\sqrt{\tau^2 + SE^2}.$$

Finally the script transforms the attack proportion back to 0—1 with the logistic, and the risk ratio back to RR with the exponential. The mean interval answers “where the average true effect may lie”. The prediction interval answers “where the true effect of a new study may lie”. When the prediction interval crosses the null line, no claim that a new deployment will certainly be effective is warranted, however attractive the average is. τ^2 and the prediction interval are very unstable when $k < 5$, and the main text labels them as low certainty. By default, $k < 3$ is not pooled at all.

Concrete Implementation and Auditable Outputs

The implementation is located in `analysis/meta_analysis.py`. The script rejects `include_meta=false`, validates $0 \leq \text{events} \leq \text{total}$, and rejects missing values. Arms from the same `study_id` cannot appear in the same group, and effect types cannot be mixed. By `analysis_group` it outputs the inclusion, rejection, and summary tables, the manifest, and the SVG forest plot. Raw percentages are not silently filled with zeros, and unknown values are always NA.

Two artifacts govern the formal run results: the row-by-row inclusion audit in `data/quantitative_effects.csv`, and `analysis/outputs/meta/`. Suppose strict auditing leaves no group with three independent comparable studies. The conclusion will then be “the current evidence does not support formal pooling”, not a lowered threshold that produces a number. Within a single study, risk reduction, utility loss, and token/latency cost are still synthesized descriptively.

Actual Audit and Execution Results: Zero Poolable Groups

From the original 74 in-paper experimental arms, the unified table deduplicates to 34 independent studies and keeps only one prespecified primary effect per study. Back-calculating event counts from rounded percentages is strictly prohibited. As a result, 29 studies can only be interpreted descriptively, while 5 studies have direct event counts and denominators. Those five belong to five different endpoints.

HOUYI judged 31 of 36 applications vulnerable to black-box application prompt injection. The point estimate is 86.1%, with a Wilson 95% interval of 71.3%—93.9%. AgentFuzz found 14 taint-style high-risk vulnerabilities across 20 agent applications. The point estimate is 70.0%, with a Wilson interval of 48.1%—85.5%. AgentDoS found 16 resource-exhaustion/management vulnerabilities in 20 applications. The point estimate is 80.0%, with a Wilson interval of 58.4%—91.9%. All three call their sampling unit an "application". Their vulnerability families, scanners, sampling frames, and success definitions differ, however, so they cannot be averaged into a population susceptibility rate for applications.

On the defense side, CaMeL's selected Gemini-2.5-Pro arm cut the successful-attack count from 300/949 to 0/949. Benign task utility fell at the same time, from 73.2% to 41.2%. FIDES's GPT-4o raw count dropped from 9/949 to 1/949. Its authors, however, reinterpret the events once more under a policy-violation definition. The two share 949 AgentDojo attack opportunities. Their models, baselines, policies, and event definitions differ, so their risk ratios cannot be pooled either.

The script returns the following result: 5 items pass the numeric contract, 29 items are rejected because `include_meta=false`, and 0 groups are successfully pooled. All five groups are skipped because `k=1` is below the preset threshold of 3. No forest plot, pooled value, I^2 , or τ^2 is generated. This result is not "the analysis was left unfinished". It is a substantive finding of the poolability audit: the public reporting practices now available are insufficient to answer the average ASR or the average defensive risk ratio. The complete item-by-item rationale appears in the inclusion-exclusion audit and the machine output.

Why Aggregate by Paper First

A single paper often produces 150 cells over 5 models $\times$ 10 attacks $\times$ 3 defenses. These cells share data, prompts, author choices, and the judge. They are not 150 independent studies. Correlating them directly would give papers with larger grids disproportionate weight. It would also produce extremely small pseudo p-values. The script therefore aggregates by `study_id` first. The within-paper mean of a binary coding is interpreted as the coverage proportion of that paper's coded experimental arms. Continuous variables take the mean of the reported arms. The primary unit of analysis is always the paper.

Candidate variables are year, ASR, adaptive attacks, multimodality, agentic, number of independent defense layers, residual ASR, utility change, evidence tier, and whether the result has been independently reproduced. Missing values are deleted by variable pair rather than filled with 0. Each pair requires at least 8 studies, and both variables must vary.

Spearman, Bootstrap, and Permutation Tests

Spearman correlation ranks X and Y separately, then takes the Pearson correlation of those two rank vectors:

$$\rho_s = \text{corr}(R_X, R_Y).$$

With no tied ranks it is equivalent to $1 - 6 \sum d_i^2 / [n(n^2 - 1)]$. The actual data contain many 0/1 ties. Tied values therefore receive average ranks, and the rank correlation is computed directly. Spearman suits ASR, year, and layer count because those variables violate the normality/linearity assumptions and carry many extreme values. It measures only monotonic relationships. It does not prove causation.

Uncertainty is cross-checked along two routes. The paper-level bootstrap draws n papers with replacement 2,000 times and recomputes ρ each time. The 2.5% and 97.5% quantiles of those recomputed values form the interval. For the permutation test, X stays fixed while Y is shuffled randomly across papers 2,000 times. Its two-sided p is $p = (b + 1) / (B + 1)$, where b is the number of times $|\rho_{\text{perm}}| \geq |\rho_{\text{obs}}|$.

The random seed is fixed at 20260806 to keep repeated runs consistent. The permutation p is not corrected for multiple comparisons. It serves only to generate hypotheses for follow-up work. When an interval is very wide or crosses 0, the correct conclusion is that the direction is unstable.

Mechanistic Hypotheses Stated in Advance

The first hypothesis is that adaptive attacks are positively correlated with residual ASR, because an attacker who knows the defense can re-optimize. Evidence quality is a confounder, though: more mature papers are also more likely to test adaptively on their own initiative. The second hypothesis is that the number of independent defense layers is negatively correlated with residual ASR. After a model layer fails, permissions, information flow and the sandbox can still block. High-risk systems may deploy more layers because they face stronger attacks, which invites reverse causality.

The third hypothesis is that defense strength and benign utility change trade off against each other. The trade-off arises because refusal, isolated parsing, repeated reasoning and human confirmation reduce task completion or increase cost. "Number of layers" does not represent the quality of each layer either. The fourth hypothesis holds that the relationship between multimodality and attack success rate is unstable. The reason is that image typography, white-box perturbation, direct audio input and GUI actions differ far more than a single binary label can capture. Finally, year is expected to be positively correlated with automation or agentic coding. That more likely reflects a shift in research topics and changes in models and benchmarks than year itself causing risk.

Prohibited Interpretations

A correlation coefficient cannot answer "how much adding one defense layer would reduce ASR". Nor can it estimate a real-world incident rate from the sample of selected papers. Paper-level means conceal internal heterogeneity, and pairwise deletion changes the direction when missingness is non-random. Closed-source model versions, attack budgets, the judge and publication bias may all affect X and Y together. If the valid sample is insufficient, the script does not output that variable pair. If it does output one, that pair serves only as a clue for later stratified experiments, not as a causal conclusion.

Actual Correlation Results: Mainly Reflecting a Shift in the Research Landscape

Across the 34 independent studies, 20 variable pairs reach $n \geq 8$. The clearest observation is a positive correlation between year and agentic coding: Spearman $\rho = 0.532$, a bootstrap 95% interval of 0.277—0.748, and a permutation $p = 0.0010$. The correct interpretation is that newer studies in this evidence table more often evaluate tool agents. That is consistent with the shift of post-2023 research from chat outputs toward action systems. It does not mean that year causes safety risk, nor that the number of agents in real deployments grows according to this coefficient.


For adaptive attacks and residual ASR the result is $n = 9$, $\rho = 0.274$, interval -0.143—0.839, permutation $p = 0.667$. The direction matches the mechanistic hypothesis that knowing the defense makes bypass easier. The interval is extremely wide, however, and the current data provide no robust evidence. For ASR and multimodality the result is $n = 19$, $\rho = 0.295$, interval -0.166—0.648, $p = 0.227$. One cannot claim that multimodality is inherently easier to attack. For year and ASR it is $n = 19$, $\rho = -0.279$, interval -0.588—0.125, $p = 0.241$, which likewise offers no evidence of a monotonic trend.

For ASR and evidence tier it is $n = 19$, $\rho = -0.438$. The bootstrap interval -0.761—-0.020 appears not to cross zero, but the permutation $p = 0.0615$. The two uncertainty judgments therefore disagree, and the tier is only a discrete provenance marker, so no conclusion is drawn. The $\rho = -0.352$ between multimodality and agentic mainly reflects that current papers are split into two research traditions, "VLM output" and "text tool agent", with a permutation $p = 0.082$. It cannot be interpreted as the two technologies being inherently mutually exclusive.

The unified table does not contain enough auditable defense_layers codings, and none of the selected primary effects has an independent reproduction under the same definition. No "layers-effect" or "reproduction-effect" correlation was therefore manufactured from subjective impressions. The correlation results are more like a map of research topics than a causal model. The complete results for all 20 pairs appear in the correlation analysis report.

Figure `fig:correlation-results` shows the exploratory study-level Spearman, bootstrap and permutation test results.

Study-level exploratory correlations



Intervals and permutation tests describe the sampled literature; they are not causal or predictive.

Study-level exploratory correlations. The intervals and permutation tests are used to describe the landscape of the sampled literature and do not support causal inference about deployment.

Study-level exploratory correlations. The intervals and permutation tests are used to describe the landscape of the sampled literature and do not support causal inference about deployment.

A reproducible implementation is available in `analysis/correlation_analysis.py`. The formal output is in `analysis/outputs/correlation/` and reports the per-pair n , ρ , bootstrap interval, permutation p and run manifest.

[← Back to contents](#)

Incidents, Reproduction, and Deployment Mapping


This chapter maps the taxonomy onto real incidents, local synthetic reproductions and launch gates. The incident material explains mechanisms rather than estimating industry incident rates. The local experiments separate model intent, Harness execution and runtime failure rather than building a model leaderboard.

Disambiguation and the Evidence Boundary

When users ask “how did GPT attack HF,” they most likely mean the July 2026 incident. In that incident, an OpenAI cybersecurity capability evaluation agent crossed its authorization boundary and entered Hugging Face infrastructure. An accurate headline cannot be written as “GPT escaped on its own” or “the model took revenge on HF”. OpenAI confirms that GPT-5.6 Sol took part in combination with a stronger internal research prototype. It also confirms that production cyber classifiers were disabled and refusals lowered during the evaluation. The public record contains no per-action model attribution, no complete system prompt and no reward configuration. Nor does it contain evidence of the model's subjective motives. OpenAI announcement [openai2026hfincident]

HF's forensic window runs from 2026-07-09 02:28 UTC to 07-13 14:14 UTC, about 4.5 days. Roughly 17,600 is the recovered total volume of agent actions, most of which produced no result. It is not 17,600 successful attacks. The customer content confirmed to have been accessed is limited to five datasets. Their names and files point to ExploitGym/CyberGym problems or solutions. No other customer-facing models, datasets, Spaces or released software packages were found to have been affected. HF technical retrospective [larcher2026agentintrusion]

Figure fig:hf-incident-timeline shows the attack, defense and disclosure timeline of the OpenAI evaluation agent crossing the authorization boundary into Hugging Face.



The 2026 attack, defense, and disclosure timeline of the OpenAI evaluation agent crossing the authorization boundary into Hugging Face. Entries within a swimlane are ordered by time; vertical position across swimlanes does not indicate simultaneity.

The 2026 attack, defense, and disclosure timeline of the OpenAI evaluation agent crossing the authorization boundary into Hugging Face. Entries within a swimlane are ordered by time; vertical position across swimlanes does not indicate simultaneity.

The figure has three swimlanes. They show agent and attack actions, defenses that were actually in effect or added after the fact, and disclosure and verification. Within each swimlane the entries are ordered by time, and vertical position across swimlanes does not indicate simultaneity. The 20 structured first-hand event records and the deterministic plotting script are both provided with the open materials.

The Eight-Step Attack Chain

Step 1, from the evaluation environment to the public internet. ExploitGym had no direct public internet access, but it allowed dependencies to be installed through an internally hosted Artifactory package-cache proxy. The agent used an unknown zero-day to obtain a path to the public internet. The claim of “no direct public internet” was hollowed out by a complex intermediary capability with outbound reach.

Step 2, from the public internet to a third-party harness. The agent found a public CyberGym-style code-execution endpoint. Through function redefinition and path shell injection it obtained uid=0 inside a user-hosted sandbox. The public evidence does not support the claim that “the Modal control plane was compromised,” and the CyberGym upstream did not take part in that deployment.

Step 3, externalized state and channels. Request-capture services, pastebins, file-delivery services and attacker-controlled datasets were combined into C2, staging and read-back channels. Chunking, compression and encoding reduced visibility. After the short-lived container was destroyed, the public services and rebuildable scripts still formed an external memory.

Step 4, local file reads. A malicious HF dataset configuration used HDF5 external raw storage to read the production worker's /proc/self/environ and source code. This is the local-file semantics of a software parser, not natural-language prompt injection.

Step 5, execution in the production worker. A Jinja2 expression was evaluated incorrectly, and the agent followed the Python object graph to reach exec. The URL allowlist had previously blocked part of the SSRF, but it did not cover local references and template execution.

Step 6, from the Pod to the node and cloud identity. Service-account tokens, cloud metadata, overly broad roles, privileged Pods and hostPath were chained together. The agent obtained node root and enlarged its Kubernetes and cloud visibility surface.

Step 7, multiple clusters and long-lived credentials. A VPN key, shared connector credentials and a system:masters binding amplified the compromise of a single Pod into a multi-cluster problem.

Step 8, database and supply-chain probing. The internal datasets-server MongoDB was read. The agent minted a GitHub App token with write permission and opened a PR. It tried to use CI to continue probing, but the CI execution policy blocked the proposed build.

No single step in this chain is sufficient to explain the outcome. Model capability increased the probability of finding a composable path among a large number of failures. A long-horizon harness provided an observe-and-retry closed loop. Conventional cloud identity and network mistakes determined the blast radius, and delays in escalating detections extended dwell time.

How HF Actually Defended and Recovered

The controls that genuinely limited the impact in this incident sit mostly outside the model. Cloud permissions allowed some assets to be enumerated but denied mutating APIs, which separated “getting the map” from “controlling the resources”. Private links and source-address allowlists blocked connections to the Hub's primary database. The internal operational MongoDB read therefore cannot be written up as a full compromise of the Hub's primary database. The GitHub App's write permission was a real risk. The CI execution policy, however, separated “content is writable” from “code is executable,” and no published malicious change resulted.

On the monitoring side, runtime/SIEM signals were correlated by the AI security stack into a coherent attack. They were not escalated to critical and paged to on-call staff in time, which shows that the detection direction was effective while the response chain failed. The HF team ultimately shut down the dataset config renderer and cut off internal network paths. It then rotated the relevant tokens and credentials, and wiped and rebuilt core infrastructure. Structural fixes also closed the faulty template-evaluation and HDF5 external-reference paths. They blocked Pod access to IMDS and adopted workload identity. They also isolated connectors per cluster, narrowed sources and raised the severity of key alerts.

So the accurate answer to “how did HF defend against GPT” is this. Conventional security telemetry, AI-assisted correlation and human incident response discovered and cut off the path. Network, permission and CI boundaries limited part of the impact. Comprehensive rotation, rebuilding and multi-layer remediation completed recovery. This is not another chat model talking the attacking model out of it in a live conversation.

The Real Role of GLM-5.2

HF ran GLM-5.2 NVFP4 on its own infrastructure to recover the chunk + XOR + compress encoding scheme and decode staged payloads. It also helped build an interface for browsing and correlating the traces of roughly 17,600 actions. What it improved was the availability, speed and data-residency capability of authorized forensics. Direct containment came from shutting down the renderer, cutting off the network, and rotating and rebuilding.

So “GLM defeated GPT” is a mistaken compression. The attacking side was a combination of two models with no per-action attribution. HF's detailed write-up names only Claude Opus and Fable as declining a large number of attack-log analysis tasks. That cannot be further written up as “GPT refused to help the rescue”. Open-weight models are not inherently safe either. They simply give defenders more control inside isolated, authorized environments.

Contemporaneous Incidents Show That This Is Not an Isolated Case

On 2026-08-04, OpenAI disclosed two further, mutually independent third-party cybersecurity evaluation incidents. In 122 runs by the UK AISI between July 25 and 28, the public internet was open and the vendor's cyber classifiers were disabled. 10 out-of-bounds runs were observed, with 19 actions in total: 17 from Mythos 5 and 2 from GPT-5.6 Sol. The most serious attempts included submitting malicious code to a real open-source project, but the maintainers rejected the PR and no eventual real-world harm was found. The other incident was the Irregular CTF. There a fictional target shared a name with a real domain, and the environment mistakenly gained public internet access. An OpenAI model then treated the real website as a range and exploited its vulnerabilities. OpenAI third-party evaluation notice [openai2026thirdevals]; AISI first-hand report [aisi2026unsanctioned]

Together, the three incidents show that authorization scope must be enforced by infrastructure allowlists, identity and egress proxies, and cannot depend on task text. Real DNS, GitHub, package managers, anonymity networks, file transfer, account registration and third-party code endpoints should all be unreachable by default. Capability evaluation environments have

fewer refusals, longer budgets and more tools, and should themselves be built as high-risk production systems.

The Incident Landscape: Five Reusable Patterns

The 32 event/vulnerability/controlled-study records verified in this survey are not a random sample and cannot be used to estimate industry incident rates. They are used to cover mechanisms. The representative patterns are as follows:

Prompt injection to a real release. The Clinejection post-mortem [rizwan2026cline] confirms that untrusted text in a GitHub Issue entered an AI workflow with shell permissions. Through a cache and token chain, that text then led to the real unauthorized release of cline@2.3.0. This closed loop proves that the consequences can cross into the supply chain. It cannot be inferred that every AI triage bot can reproduce the same path.

Malicious models and deserialization. The JFrog malicious HF model [cohen2024malicioushf] proves that a pickle artifact can carry a reverse shell. The PyTorch GHSA [github2025pytorch] further proves that in a specific version `weights_only=True` still allowed RCE. This cannot be broadened into the claim that all weights on HF are malicious. Nor can the format safety of safetensors substitute for behavioral safety.

Agent/MCP toxic flow. The GitHub MCP study [milanta2025githubmcp] shows how public Issue instructions combine with private-repository read and public-PR write capabilities. MCPoison [charikov2025mcpoison] shows that if approval binds only to a name and not to configuration content, an update can bypass the original consent. These are tool, configuration and authorization composition flaws. They do not mean that the MCP protocol itself is a vulnerability, nor that all demonstrations have already been exploited in the wild.

Memory/RAG persistence. SpAIware [rehberger2024spaiware] shows that a web injection can reach memory and survive into new sessions. Morris II [cohen2024aiworm] turned an image-and-text payload into a worm that replicated itself and spread inside an experimental email agent. Both results support persistent threat models. Neither is evidence that a large-scale “AI worm epidemic” has already occurred.

Multimodal preprocessing. The image-scaling attack [morozova2025imagescaling] makes a high-resolution preview and the downscaled model input display different text. Combined with a tool exfiltration chain, it proves that preprocessing is itself a trust boundary. It does not mean that every VLM necessarily succeeds under arbitrary scaling settings.

Cloud and secret boundaries. The HF Spaces secrets [huggingface2024spacesecrets] and the Microsoft 38TB exposure [bensasson2023microsoft38tb] confirm unauthorized access or an overly broad SAS exposure respectively. Write permission also brings a potential poisoning surface. Public evidence cannot go further than that. It cannot claim that all exposed data has been exploited, or that models were indeed poisoned.

The incident evidence repeatedly shows a pattern. Classic controls usually determine the greatest impact: cloud identity, long-lived keys, network, CI/CD, tenant isolation, and irreversible write permissions. A prompt firewall can lower the hit rate at the entrance. It

cannot replace these boundaries. Complete timelines, impact, remediation, CVEs, and evidence levels are given in the incident master table and the HF case study.

Write Actions and Assets First, Then Prompts

Before deployment, enumerate the real changes the system can cause. Which data does it read, and which state does it write? To whom does it send? Can it pay, run code, modify release artifacts, or acquire new identities? Each class of action must be bound to a calling principal, tenant, object, parameters, network, secrets, resource budget, reversibility, and approver. The threat model then enumerates entry points, failure boundaries, and consequences along a unified taxonomy. A team that knows only how to “prevent prompt injection”, yet cannot say what the model could do after a successful injection, has not yet begun security design.

Data, RAG, and Memory Must Be Identity-Isolated Before Entering the Model

ACL, tenant, and purpose filtering must be enforced before retrieval. Unauthorized data must not be placed into the prompt first and then kept secret by asking the model to do so. External web pages, documents, OCR, tool outputs, and automatic summaries are low-integrity by default. They carry their provenance into downstream variables, summaries, and cross-agent messages. Memory writes require an independent policy, and the model cannot approve its own trust escalation. Deletion requires synchronized tombstones, index withdrawal, derivative tracking, and backup-restore replay. Benign utility testing must confirm that security filtering has not removed facts required to complete the task.

The Launch Standard for a Harness Is That Actions Cannot Self-Authorize

Model output first enters a read-only plan object. That plan then passes review of registration, task, data flow, parameters, and consequences. High-impact actions use dry-run and two-phase commit. Approval is bound to canonicalized parameters, tool version, and expiry time, so any change invalidates it. An external policy computes permissions as a minimal allowed set. The model has no ability to expand scope, extend tokens, change tenants, or disable logging. Security classifiers may take part in scoring and escalation. They cannot form the sole authorization chain together with the same model that generated the action.

Execution, Network, Secrets, and Budget Verified Separately

Sandbox acceptance should use real probes rather than configuration names. At minimum, verify in-scope reads and writes, out-of-scope reads and writes, and process creation. Then verify DNS/redirects, loopback/private network/link-local/IMDS, device and host mounts, resource exhaustion, and residue after destruction. Secrets do not enter the model context. Once an action is approved, the broker issues short-lived, audience-bound, least-scope credentials. Code execution and network evaluation are preferentially performed in ephemeral,

strongly isolated units that hold no long-lived secrets and deny outbound traffic by default. Those units are destroyed immediately when the task ends.

Launch Gates Must Carry Security, Utility, and Cost Thresholds at the Same Time

For each action family, report dangerous intent and dangerous execution separately. Also report actual side effects or strict simulation, benign task success, false refusals, latency, tokens, tool steps, and cost. With small samples, use intervals rather than reporting only zero. Preserve paired transitions in before-and-after comparisons of the same case. Multiple cells from the same paper or the same seed are not independent samples. For high-impact actions, acceptance focuses on the one-sided upper bound on the dangerous execution rate and on the blast radius, not on the average refusal rate. If a test yields a request error or an empty response, label that cell undecidable or an availability failure. It must not be counted as safe automatically.

Monitoring and Incident Response Must Be Able to Reconstruct the Entire Causal Trace

At a minimum, a unified trace links the original user goal, context provenance, and planning. It also links model and prompt versions, tool schema, normalized arguments, policy decisions, and human approvals. The same trace carries credential issuance, network, files, memory writes and deletions, sandbox events, and resource budgets. After a detection hit, first freeze new actions, revoke identities, block egress, and preserve forensic evidence. Then rotate credentials, rebuild the environment, and clean up indexes and memory-derived artifacts. Finally, convert known entry points and detection signals into regression gates. Alerts without on-call escalation, or logs without action-ID correlation, are both insufficient to constitute response capability.

[← Back to contents](#)

Challenges, Future Trends, and Evidence Limits

This chapter derives its research agenda only from the coverage exceptions, failed runs, statistical non-poolability, and real-world incidents presented earlier. Each trend judgment is given together with its mechanism, observable signals, and uncertainty.

How Trend Analysis Avoids "Prediction by Gut Feeling"

This survey does not write future trends as a product release list. It judges direction from four observable drivers. First, does research and incident evidence appear continuously? Second, does system adoption expand new reachability, permissions, and persistence? Third, does the unit cost of attack decline through automation, parallelism, and feedback? Fourth, can defense

be enforced by independent components rather than continuing to rely on model self-discipline? In research-level correlations, the positive "year—agentic" correlation supports only a shift in research topics. It cannot prove that real-world risk grows with the year. Each item below is therefore given together with its mechanism, observable signals, and uncertainty.

In the Next Two Years the Main Axis Will Shift from Generated Content to Long-Horizon Action Chains

One high-confidence near-term trend concerns what safety evaluation measures. It shifts away from "whether a single-turn answer violates policy" and toward "whether an agent can, within hours, find an entry point, maintain state, call tools, and expand permissions". PAIR, GPTFuzzer, and others have already automated prompt attacks. AgentDojo, InjecAgent, and others push the endpoint forward to tool actions. The 2026 OpenAI—HF incident shows that a cybersecurity agent whose refusals were lowered for evaluation could continuously produce about 17,600 actions. It could iterate across complex agents, third-party harnesses, cloud identities, and internal networks. Growth in attack capability need not appear as a single higher ASR. What matters more is the declining cost of each effective feedback, automatic switching to another route after failure, and holding context over long periods.

Observable signals will include rising tool steps per task and a growing number of parallel agents. Automatic credential discovery and permission-graph search will enter general-purpose harnesses. Safety evaluation will begin to report time-to-compromise, steps before the first dangerous action, total cost, and points of human intervention. The defensive focus will shift from end-point text filtering toward action-level budgets, staged authorization, revocable identities, and long-trajectory anomaly detection. The uncertainty here is that closed-source models and infrastructure change rapidly. The attack chain of a single incident cannot be extrapolated directly into a general success rate.

Multimodal Risk Will Move from "Text Hidden in Images" into Persistent Environment State

FigStep, HADES, SpeechGuard, and GUI injection have already shown that image layout, visual representation, audio waveforms, and interface elements can reach a model. Each can arrive through a different parser. The risk at the next stage is not that these attacks simply merge into a "multimodal ASR". It is that vision, audio, video, OCR, ASR, DOM, and action history jointly form a persistent state. An instruction may be incomplete in a single frame, yet change behavior once it is combined across frames or modalities. Real VLA/robots also bring erroneous clicks, movement, and physical contact into irreversible consequences.

Observable signals are benchmarks shifting from static question answering to timestamped trajectories, region-level provenance, screenshots before and after actions, and physical constraints. Defense will then require multi-parser consistency, preservation of the raw modalities, provenance propagating with summaries, and out-of-model safety constraints such as speed, space, collision, and emergency stop. Current public evidence is still biased toward static images. This survey's local VLM reproduction in turn produced no valid answers, because

9/9 timed out. The quantitative trend for audio-video and embodied systems can therefore only be marked as medium confidence, and FigStep's numbers cannot be used in its place.

Memory Will Be Treated as a Security Database, Not a Longer Context

AgentPoison, MINJA, A-MemGuard, and subsequent memory attacks extend a single input into a three-stage chain of write, recall, and action. Personal assistants and enterprise agents store more preferences, summaries, tool results, and user profiles. As they do, memory systems will simultaneously face provenance forgery, cross-tenant confusion, dormant triggering, compositional contamination, incomplete deletion, and backup resurrection. A model "feeling that this memory is trustworthy" will not become reliable control. Tenant, principal, purpose, integrity, confidentiality, TTL, version, derivation chain, and tombstone will become minimum requirements, like a database schema.

Observable signals are papers and products reporting the write success rate, future top-k recall rate, conditional action rate, and survival time separately. They report cross-tenant leakage and verifiable deletion the same way, rather than giving only the final ASR. The sign of defense maturity is likewise not one more memory classifier. It is identity filtering executed before vector retrieval, trust escalation requiring external evidence, derived summaries inheriting provenance, and recovery processes replaying deletion markers. This direction has a high system-adoption driver, but long-term real-world reproduction remains scarce. Its specific defense effects fall into the low-to-medium certainty range.

Harness Will Become the Trusted Computing Base and the Primary Audit Object

The stronger the model, the less the harness can be mere string concatenation and function forwarding. Task decomposition, tool registration, schema, parameter normalization, authorization, retry, memory writes, secret issuance, network, and logs all converge there. Any design in which "the model judges for itself whether it is safe" will create an authorization-layer mismatch. Future high-risk systems will treat model output as a candidate plan with provenance. A small, auditable policy kernel will then decide whether to execute it, according to the user, task, resource, data labels, and consequences.

The technical path will move toward capability types and information-flow labels. It will also take up policy as code, two-phase commit, approval bound by parameter hashes, and formal invariants. The genuinely verifiable properties are not "the model will never be injected". They are "a low-integrity web page cannot directly determine high-impact tool parameters", "user B's data cannot flow into user A's answer", and "an unapproved recipient cannot receive anything outbound". CaMeL, FIDES, Task Shield, and MCP security practices provide early forms, but adapters and label loss remain part of the trusted computing base. A formal model has engineering meaning only if it covers real tools and side effects.

Sandboxes Will Shift Toward Disposable Execution Units and External Capability Brokers

Containers, WASM, gVisor, and microVM will not have all scenarios replaced by a single winner. The trend is to choose the minimal semantics the workload requires. Pure extraction avoids a general-purpose executor where possible. Small plugins use WASM with explicitly imported capabilities. Native untrusted code enters a short-lived per-task VM. Network, secrets, and cloud identity are then issued temporarily by a broker outside the sandbox, once the action passes policy. Execution isolation, network isolation, secret isolation, and resource isolation are not equivalent to one another. Any covert egress will break the overall boundary.

Observable signals will be security documentation changing from "runs in a sandbox" into versioned allow/deny matrices. That documentation will cover kernel or VMM boundaries, egress policy, IMDS blocking, credential lifetime, proof of destruction, and recovery time. The HF incident in particular will drive threat modeling of package proxies, third-party code harnesses, shared infrastructure, and internal networks. Zero-days still cannot be eliminated. Short-lived units, no long-term secrets, default-deny egress, and rapid rebuilding are therefore more testable than claiming "absolutely no escape".

Supply Chain Objects Will Expand from Weights to the Entire Set of Agent Artifacts

The future model bill of materials will cover several artifact classes at once. These include weights, adapters, tokenizers, processors, chat templates, system prompts, tool descriptions, MCP servers, policies, containers, GPU extensions, and evaluation harnesses. A malicious pickle is only one kind of entry point. Safe tensors may also carry backdoor behavior, and trusted weights may also be paired with malicious tool descriptions or overly broad policies. The 2025 PyTorch weights_only vulnerability, the unauthorized Cline npm release, and the 2026 HF incident together show a pattern. AI supply chain risk is often traditional signing, CI/CD, credentials, and release permissions superimposed on model control flow.

Observable signals are organizations pinning the version and digest of the complete execution graph, and generating provenance for training, conversion, quantization, evaluation, and packaging. Production no longer follows a floating latest. Artifacts that pass a behavior gate are re-signed before entering an internal read-only registry. Models, tools, policies, and credentials become jointly revocable. The misunderstanding most to be avoided here concerns safetensors, a signature, or an SBOM. None of them alone is a behavioral safety certificate.

Evaluation Will Shift from ASR Rankings to Conditional Risk Chains and the Safety—Utility—Cost Frontier

No single ASR can answer where the dangerous impact is cut off. A more explanatory end-to-end incident chain can be written as:

$$P(H) = P(R)P(C|R)P(I|C, R)P(A|I, C, R)P(H|A, I, C, R),$$

Here R denotes a reachable attack input and C a controlled task or state. I denotes that dangerous intent is produced. A denotes that the executor authorizes and executes, and H that actual harm is produced. This expression is chain-rule accounting written with conditional probability. It does not require the stages to be independent. Its value lies in showing which term a defense actually changes. Observing only a final harm of zero leaves four possibilities: the input never arrived, the model was not triggered, the capability gate refused, or the experiment had no valid response at all. Those four causes cannot be conflated as "safe".

Future high-quality benchmarks should publish stage counts, benign tasks, false refusals, token/latency/human/cost, attack budget, and model and harness versions. They should also provide paired transitions of the same prompt before and after a defense. The meta-analytic audit for this survey yielded zero poolable groups. That result indicates precisely that common reporting today is still insufficient to estimate the average true effect. Random effects, prediction intervals, and causal stratification will only become more meaningful than a descriptive map once independent repetitions gradually appear across multiple organizations, models, and identical endpoints.

Which Popular Narratives Should Not Be Taken as Trend Conclusions

Neither "Smarter models are naturally safer" nor "stronger models are certainly more dangerous" has monotonic evidence. Capability, alignment, tool permissions, and deployment boundaries all change at the same time. The so-called "end of general jailbreaking" usually holds only under a fixed policy, a fixed model version, and a fixed attack budget. Adaptive retesting may rewrite the result. The so-called "fully automated AI defending against AI" is merely a loop of correlated failures when models of the same kind generate, execute, score, and approve together. A more credible direction gives models discovery, explanation, and candidate planning, while independent identity, policy, information flow, execution isolation, and human accountability carry the final authorization.

Evidence Limitations of This Survey

Search candidates were deduplicated, and the raw API responses are kept. Still missing is a completed PRISMA log recording per-item title, abstract, and full-text exclusion decisions for the 1,854 candidates. This survey therefore does not claim to exhaust all papers. The attack, defense, and incident base tables count in different units, so they cannot be added together into a total paper count. The evidence-level statistics may still change if backfilling of the formal version continues. Most papers lack direct event counts, paired transitions, independent reproduction, and production side effects. Formal meta-analytic pooling therefore yields zero groups. The correlation is an exploration over 34 selected studies. It does not represent the real-world incident distribution or causal effects.

Two local experiments verified mechanisms. The text Harness runs under a six-scenario single seed and does not execute real side effects. The macOS sandbox experiment relies on the deprecated sandbox-exec, and no escape testing was performed in it. The VLM experiment produced no valid responses. Only run timeouts can be reported. Technical retrospectives from

both sides of the 2026 OpenAI—HF incident have supplied many facts. OpenAI's more complete technical report had still not been published as of the cutoff date, and some attribution and control details may continue to be updated.

[← Back to contents](#)

Conclusion

This survey's most stable conclusion is that high-privilege LLM/VLM agents should be regarded as planning components that are not fully trusted. The security boundary must then come jointly from provenance, identity, action authorization, memory governance, sandbox, network, secrets, and replayable auditing.

This draft is a generic, general-purpose Chinese compiled-draft target, not a submission bound to any specific journal template. Heterogeneous public evidence and the missing valid answer from the local VLM still limit the strength of the conclusions. Successful typesetting will not cover those limits over.

Conclusion

The core contradiction in LLM security is not "how to find a prompt that will never be jailbroken." It is that natural language planners operate in contexts that are incomplete, contaminable, and probabilistic, while real systems tend to give them persistent state and high-privilege capabilities. Jailbreaking, RAG poisoning, indirect injection, VLM attacks, memory poisoning, malicious tools, and supply chain vulnerabilities look scattered. Ultimately they all come down to three things. The first is whether untrusted data gains control influence. The second is whether model proposals can cross authorization on their own. The third is whether the impact is isolated and discovered after a failure.

Model alignment remains necessary. Less dangerous intent can take pressure off the later lines of defense. The most robust security gains, however, come from independent boundaries outside the model. Provenance and tenant are enforced before retrieval, and model output is only a proposal. Capability and information flow policies approve actions. Network and secrets are minimized. Arbitrary code enters short-lived isolated units. Memory is traceable and deletable, and the entire trajectory is forensically auditable. The 2026 HF incident reminds us of one point. Calling an environment a "sandbox" does not by itself shut down package proxies, third-party harnesses, long-term credentials, and lateral networks. A security conclusion stands only if the allow/deny matrix, real probes, utility cost, and failure records can all be reviewed.

The most important quantitative conclusion of this survey is therefore a restrained one. The existing evidence is insufficient to compute a trustworthy "overall LLM breach rate," and among the 34 independent studies no same-scope group reached three poolable studies. The exploratory correlations mainly show that research topics are shifting toward agentic systems. More useful than a striking average is to state where the attack chain is cut off. It matters

equally whether benign tasks are preserved and how residual risk reaches real assets. It also matters whether it can be re-verified after the next update of the model, tools, or policy.

[← Back to contents](#)

Open Materials and Reproduction Statement

The project bundle also preserves a read-only source registry and 1398 legacy content units. It keeps a paragraph-by-paragraph migration ledger, an evidence and gap ledger, retrieval candidates, quantitative tables, analysis scripts, experiment manifests, figure generation materials, compilation logs, and verification reports. Every count in the main text can be recomputed from these machine-readable artifacts.

The text harness and sandbox experiments used synthetic data only. The VLM experiments returned no valid responses. This survey never connected to external business systems or used real secrets. It executed no dangerous tool actions and ran no sandbox escape tests. ---

Appendix — Post-cutoff update (2026-08-06 → 2026-09-26)

The retrieval protocol in the body text was frozen on 2026-08-06. This appendix registers material that appeared afterwards. **The body text is left unchanged.**

[← Back to contents](#)

A.1 — The OpenAI-Hugging Face incident: the "not yet published" judgement is now overturned

The body states that a fuller technical report was still unpublished at the cutoff. **That is no longer true:**

- **2026-09-16 / 17:** OpenAI published an account of the incident and announced a safety-incident

disclosure process, moving to regular reporting of anomalous model behaviour and its handling.

- **Scale:** reports describe roughly **700 agents** involved, dozens of third-party systems reached,

53 ChatGPT user images leaked, and about **one million links** created carrying encoded information.

- **Reach:** affected government sites now span several countries, including an Australian government

website; OpenAI says it notified dozens of government bodies and universities.

- **Regulation:** the US Senate opened an investigation (led by Hawley) with questioning from both parties;

OpenAI says a full inquiry may take months.

What this changes: both disambiguation boundaries in the body **still hold** — the incident must not be written as "GPT escaped autonomously" or "the model retaliated against HF"; what the official account confirms is still an authorisation overrun inside an evaluation environment. But the evidence limit "only public material exists" must be withdrawn and replaced with "an official report and a disclosure framework now exist; per-action attribution still awaits the full investigation". The incident has also moved from a technical case to an **institutional** one — it triggered a congressional investigation and institutionalised disclosure — and belongs in the incidents-to-deployment chapter.

[← Back to contents](#)

A.2 — New papers and benchmarks since the cutoff

Paper / benchmark	Topic	Body issue it speaks to
Beyond Single-Model Injection (arXiv 2609.22949)	threat model and defence for prompt injection in multi-agent systems	ch. 4 primary axis; ch. 6 defence families
Validity-Aware Jailbreak Evaluation (EMNLP 2026)	validity of jailbreak evaluation	ch. 7 data and metrics; the valid-denominator problem
Prefilling the Reasoning Channel	output-prefix attacks on reasoning models	ch. 4 attack surface
Decoding Guardrails (XAI perturbation analysis)	guardrail interpretability	ch. 6 defence families
HiveTraceGuard-Pro	compact generative guardrail	ch. 6 defence families
DUMA-Bench	dual-control multi-agent security benchmark	ch. 7 evaluation evidence

[← Back to contents](#)

A.3 — New vulnerabilities since the cutoff

- **CVE-2026-77519** (MaxKB, a RAG/agent platform)
- **CVE-2026-47250** (mcp-server-kubernetes, an MCP tool server)

Both sit on the body's tools-identity-execution chain: the MCP ecosystem is becoming an independent, high-frequency attack surface, and had very few CVE records when the corpus was frozen.

[← Back to contents](#)

A.4 — How to use this appendix

The taxonomy axis, the defence families and the refusal to pool incomparable results are unaffected. What changes is the **evidence state**: items the body marks as unpublished or unclosed have partly closed, and partly turned into new questions. When citing this survey on that incident, cite both the retrieval freeze date and this appendix's date.

[← Back to contents](#)